RK0

The Real-Time Kernel '0'

System: v0.13.0 | Docbook: 0130.1.0.0 | www.kernel0.org

Table of Contents

1. THE KERNEL AT A GLANCE . . . 1
  1.1. The design approach . . . 1
    1.1.1. Architecture . . . 1
    1.1.2. Programming with RK0 . . . 1
    1.1.3. Suitable Applications . . . 2
  1.2. Kernel Services . . . 2
2. **Task Scheduler** . . . 3
  2.1. **Task switching dynamics** . . . 3
    2.1.1. Scheduler Design Internals . . . 4
      2.1.1.1. Task Control Block . . . 4
      2.1.1.2. Task Queues . . . 5
      2.1.1.3. Ready Queue Table . . . 6
      2.1.1.4. Waiting Queues . . . 6
      2.1.1.5. The *choose-next* algorithm . . . 6
3. Timers and Delays . . . 9
  3.1. Busy delay . . . 9
  3.2. Sleep Timers . . . 10
    3.2.1. Sleep Delay . . . 10
    3.2.2. Compensated Sleep Delays . . . 11
      3.2.2.1. Periodic Release Sleep . . . 11
      3.2.2.2. Sleep Until (local anchor base) . . . 13
  3.3. Blocking Time-out . . . 15
  3.4. Callout Timers (Application Timers) . . . 15
  3.5. System Tick . . . 15
4. System Tasks . . . 17
5. Memory Allocator . . . 18
  5.1. How it works . . . 18
6. ***Inter-Task Communication*: Signals** . . . 20
  6.1. Task Events (Event Flags) . . . 20
    6.1.1. Usage Example: Supervisor Task . . . 21
  6.2. Semaphores . . . 22
    6.2.1. Counting Semaphore and Binary Semaphores . . . 23
    6.2.2. Semaphores in RK0 . . . 23
    6.2.3. Flushing or Broadcast . . . 23
      6.2.3.1. Bounded Buffer with Semaphores . . . 23
    6.2.4. Mutex Semaphore (Locks) . . . 26
      6.2.4.1. Priority Inversion and PIP . . . 26
    6.2.5. Mutexes vs Binary Semaphores . . . 30

6.3. Scheduler Lock . . . 31
6.4. Sleep Queue . . . 31
6.5. Conditional Critical Regions . . . 32
6.5.1. Monitor Invariants . . . 32
6.5.2. Signalling Discipline . . . 33
6.5.3. Condition Variable Model . . . 34
6.5.3.1. Usage Example: Synchronisation Barrier . . . 34
6.5.3.2. Usage Example: Readers Writers Lock . . . 37
6.6. Deferred Signals . . . 40
6.6.1. Use cases . . . 40
6.6.2. Functional overview . . . 40
6.6.3. Usage Pattern . . . 41
6.6.4. Minimal kDSignal example . . . 42
6.6.4.1. Separation from Message Passing . . . 43
6.6.4.2. Separation from Task Events . . . 43
6.6.4.3. Bounded and Minimal by Design . . . 44
7. **Inter-Task Communication: Messages** . . . 45
7.1. Task Mailbox . . . 45
7.1.1. Description . . . 45
7.1.2. Suitable usage patterns . . . 45
7.1.2.1. Latest-snapshot handoff . . . 45
7.1.2.2. Command/response using Mailbox and Events (synchronous) . . . 46
7.1.2.3. Double-buffering . . . 49
7.1.2.4. Coalesced updates . . . 50
7.1.2.5. Producer allocates - consumer frees . . . 51
7.1.3. Message Queue . . . 52
7.1.3.1. Size of a Message . . . 53
7.1.3.2. Public Mailbox . . . 53
7.1.3.3. Ports . . . 53
7.1.3.4. Blocking and non-blocking behaviour . . . 53
7.1.3.5. Ownership and Priority Inversion . . . 54
7.1.3.6. Notify callback . . . 54
7.1.3.7. Usage Examples . . . 54
7.1.3.7.1. Mail Queue pattern . . . 54
7.1.3.7.2. Averaging Sensor Values . . . 56
7.1.3.7.3. Queue Select using Notify Callback . . . 60
7.1.3.7.4. Synchronous Client-Server Procedure Call . . . 63
7.1.4. Most-Recent Message Protocol (MRM) . . . 66
7.1.4.1. Functional Description . . . 66
7.1.4.2. MRM Control Block Configuration . . . 67
7.1.4.3. Usage Example . . . 67

8. ***Error Handling*** 72
8.1. Fail fast 72
8.2. Stack Overflow 72
8.3. Deadlocks 73

Chapter 1. THE KERNEL AT A GLANCE

1.1. The design approach

RK0 Blog: About Processes, Tasks and Threads

1.1.1. Architecture

The layered architecture can be split — roughly — into two: a top and a bottom layer. On the top, the *Executive* manages the resources needed by the application.

On the bottom, the *Low-level Scheduler* works as a software extension of the CPU.

Together, they implement the *Task* abstraction. This primitive is the *Concurrency Unit* and follows the *Thread* model. A *Task* is a *thread*.

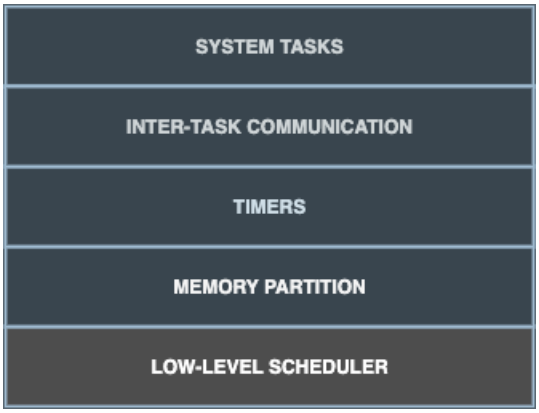

In systems design jargon, the Executive enforces policy (what should happen). The Low-level Scheduler provides the mechanism (how it gets done). The services are the primitives that gradually translate policy decisions into concrete actions executed by the Scheduler.

RK0's goal is determinism on low-end devices. Its multitasking engine does not split user space from kernel space. Tasks execute in privileged mode and use a dedicated process stack pointer, distinct from the system stack. The rationale:

- Application tasks are not unknown entities at run time.
- Implementing system calls as traps either forces a non-preemptive kernel, or increases complexity in critical paths, degrading determinism.
- Relying on the ARMv6/7-M MPU decreases memory usage efficiency and introduces latency on control paths. It does not fit RK0's deterministic execution model.

1.1.2. Programming with RK0

This document discusses the design internals, design choices and give usage examples. If not on actual usage examples, the text does not present the

formal API, but simplifications. Code snippets might present outdated API.

The current API is found here: RK0 API

1.1.3. Suitable Applications

Given the architecture, *RK0* targets applications with the following characteristics:

1. They are designed to handle particular devices in which real-time responsiveness is imperative.
2. Applications and middleware may be implemented alongside appropriate drivers.
3. Drivers may even include the application itself.
4. *Untested programs are not loaded*: After the software has been tested, it can be assumed reliable.

1.2. Kernel Services

RK0 has *Core Services* (always enabled) and optional services (enabled by configuration).

Core Services:

- Scheduler
- Partition Memory Allocator
- Timer Delays
- Task's Event Register
- Deferred Signals
- Task Exchange (per-task 1-slot inbox)

Optional Services

- Timers:
  - Application Timer (Callouts)
  - Sleep Queues
  - Counting/Binary Semaphores
  - Mutex Semaphores
  - Message Queues / Mailboxes
  - Ports
  - Most-Recent Message Protocol (MRM)

When compiled with only *Core Services*, one gets a functional Executive with less than ~3 KB ROM.

Chapter 2. Task Scheduler

Aside from design details that will be discussed later, *RK0* scheduler is a priority-based, preemptive scheduler, supporting 32 different priorities. Quite standard. The major difference from others kernels is tha it deliberately *has no built-in time-slice*.

2.1. Task switching dynamics

In practice, tasks are either *RUNNING* or '*waiting*' for their turn to run.

We need to define *WAITING* and *READY* clearly:

- A *READY* task will be dispatched; therefore, switch to *RUNNING* whenever it is the highest priority *READY* task.
- A WAITING task depends on a condition—generalised as an event—to transition to the READY state.
- Logically, the *WAITING* state will assume different pseudo-states related to the kind of event that will switch a task to *READY* (a message, a combination of event flags, a signal on semaphore, etc.)

(You can see these states on `kcommondefs.h`.)

Knowing the that the 'ready' tasks are within a table, and each row is a FIFO Queue related to priority:

> *The highest priority ready task is at the head of highest priority non-empty ready queue.*

We can outline the scheduling:

- A task must switch to the *READY* state before being eligible for scheduling.
- After a task is dispatched it will keep RUNNING until preempted by a higher priority task, blocked or yielding.
- A yield will only switch tasks if there is another task with equal or higher priority that is ready.
- A task will switch from *RUNNING* to *READY* if yielding or if being preempted by a higher priority task. Otherwise it can only go to a *WAITING* state, and eventually switch back to *READY*.
- When a task is preempted by a higher priority task, it switches from *RUNNING* to *READY* and is placed back on the *head* position of its Ready Queue. This means that it will be resumed as soon as it is the highest priority ready task again.
- On the contrary, if a task *yields*, it tells the scheduler it has completed its cycle. Then, it will be enqueued on the ready queue tail - the last queue position.
- When a task *waits* it is suspended until a condition is satisfied.
- When the condition is satisfied, it switches from *WAITING* to *READY*, and is enqueued on the tail.

- So, tasks with the same priority cooperate by either *yielding* or *waiting*.
- If a task never yields or waits, other tasks with the same or lower priority *will starve*.
- Finally, Tasks with the same priority are *initially* placed on the *Ready Queue* associated with that priority in the order they are *created*.

> *RK0* can handle context-switching with an extended frame when a float-point co-processor is available. This must be informed when compiling by defining the symbol `__FPU_PRESENT=1`.

2.1.1. Scheduler Design Internals

A notable scheduler characteristic is constant-time complexity (*O(1)*) with low latency. This was achieved by carefully composing the data structures and an efficient *'choose-next'* algorithm, detailed below.

2.1.1.1. Task Control Block

Every primitive is associated to a data structure we refer to as its *Control Block*. A Task Control Block is a record for stack, resources, and time management. The table below partially represents a Task Control Block (as this document is live, this might not reflect the exact fields of the current version).

Task Control Block
Task name
Task ID
Status
Assigned Priority
Effective Priority
Saved Stack Pointer
Stack Address
Stack Size
Last wake-time
Next wake-time
Time-out Flag
Preemption Flag
Owned Resources List
Waiting Resources List
Event Register Control Block
Mailbox
Timeout List Node

Task Control Block
TCB List Node

Tasks are static — they are not created (or destroyed) on runtime. There is no fork or join.

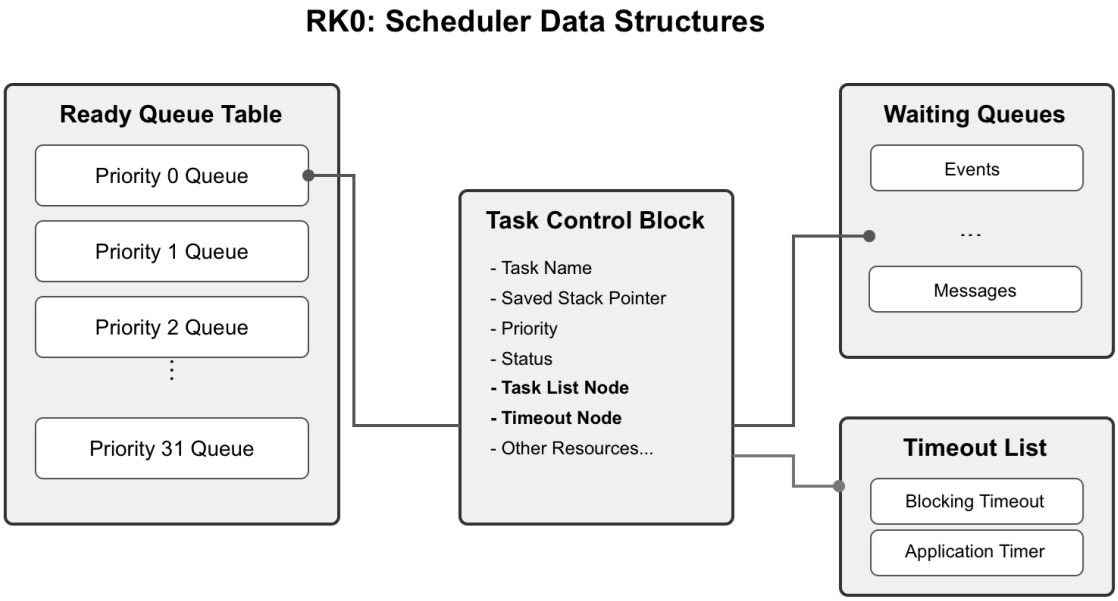

The scheduler rules, not the heap.

RK0 tasks are static.

It's a design decision rooted in real-time correctness.

Besides an application-specific system software does not need to treat tasks as 'unknown' objects.

The wins:

- A memory layout the systems programmer knows.
- No alignment traps.
- Link-time visibility:
  - Each task's stack is a named symbol in the linker map.
  - You can inspect and verify the memory layout before flashing.
  - A simple `objdump` reveals all stack allocations — that's peace of mind.

2.1.1.2. Task Queues

The backbone of the queues where tasks will wait for their turn to run is a circular doubly linked list: removing any item from a double list takes O(1) (provided we don't need to search the item). As the kernel knows each task's address, adding and removing is always O(1). Singly linked lists can't achieve O(1) for removal.

2.1.1.3. Ready Queue Table

Another design choice to achieve O(1) is the global ready queue, which is a table of FIFO queues—each queue dedicated to a priority—and not a single ordered queue. So, enqueuing a ready task is always O(1). Given the sorting needed, the time complexity would be O(n) if tasks were placed on a single ready queue.

2.1.1.4. Waiting Queues

The scheduler does not have a unique waiting queue. Every kernel object that can block a task has an associated waiting queue. Because these queues are a scheduler component, *they follow a priority discipline*: the highest priority task is dequeued first, *always*.

When an event capable of switching tasks from *WAITING* to *READY* happens, one or more tasks (depending on the mechanism) are then placed on the ready list, unique to their priority. Now, they are waiting to be picked by the scheduler—that is the definition of *READY*.

2.1.1.5. The *choose-next* algorithm

As the ready queue table is indexed by priority - the index 0 points to the queue of ready tasks with priority 0, and so forth, and there are 32 possible priorities - a 32-bit integer can represent the state of the ready queue table. It is a BITMAP:

```
The BITMAP computation: ((1a) OR (1b)) AND (2), s.t.:

(1a) Every Time a task is readied, update: BITMAP |= (1U << task->priority );
(1b) Every Time an empty READY QUEUE becomes non-empty, update: BITMAP |= (1U <<
queueIndex)
(2): Every Time READY QUEUE becomes empty, update: BITMAP &= ~(1U << queueIndex);
EXAMPLE:

  Ready Queue Index :     (6)5 4 3 2 1 0
          Not empty :      1 1 1 0 0 1 0
                           ------------->
                 (LOW)  Effective Priority  (HIGH)
In this case, the scenario is a system with 7 priority task levels. Queues with
priorities 6, 5, 4, and 1 are not empty.
```

Having the Ready Queue Table bitmap, we find the highest priority non-empty task list as follows:

(1) Isolate the **rightmost** '1':

```
RBITMAP = BITMAP & -BITMAP. (- is the bitwise operator for two's complement: ~BITMAP +
1) `
```

In this case:

```
                          [31]       [0]  :  Bit Position
```

```
                              0...1110010   :  BITMAP
                              1...0001110   : -BITMAP
                             =============
                              0...0000010   :  RBITMAP
                                      [1]
```

The rationale here is that, for a number N, its 2's complement -N, flips all bits - except the rightmost '1' (by adding '1') . Then, N & -N results in a word with all 0-bits except for the less significant '1'.

(2) Extract the **rightmost '1' *position***:

- For ARMv7M, we benefit from the `CLZ` instruction to count the *leading zeroes*. As they are the number of zeroes on the left of the rightmost bit, '1', this value is subtracted from 31 to find the Ready Queue index.

```
RK_FORCE_INLINE static inline
unsigned __getReadyPrio(unsigned readyQBitmap)
{
    unsigned ret;
    __ASM volatile (
        "clz    %0, %1     \n"
        "neg    %0, %0     \n"
        "add    %0, %0, #31\n"
        : "=&r" (ret)
        : "r" (readyQBitmap)
        :
    );
    return (ret);
}
```

This instruction would return #30, and #31 - #30 = #01 in the example above.

- For ARMv6M there is no suitable hardware instruction. The algorithm is written in C and counts the *trailing zeroes*, thus, the index number. Although it might vary depending on your compiler settings, it takes ~11 cycles (*note it is still O(1)*):

```
/*
  De Brujin's multiply+LUT
  (Hacker's Delight book)
*/

/* table is on a ram section  for efficiency */
 const static unsigned readyPrioTbl[32] =
{
 0, 1, 28, 2, 29, 14, 24, 3, 30, 22, 20, 15, 25, 17, 4, 8,
 31, 27, 13, 23, 21, 19, 16, 7, 26, 12, 18, 6, 11, 5, 10, 9
};

RK_FORCE_INLINE static inline
```

```
unsigned __getReadyPrio(unsigned readyQBitmap)
{
    unsigned mult = readyQBitmap * 0x077CB531U;

    /* Shift right the top 5 bits
     */
    unsigned idx = (mult >> 27);

    /* LUT */
    unsigned ret = (unsigned)readyPrioTbl[idx];
    return (ret);
}
``

For the example above, `mult = 0x2 * 0x077CB531 = 0x0EF96A62`. The 5 leftmost bits
(the index) are `00001` -> `table[1] = 1`.

During a context switch, the procedures to find the highest priority non-empty ready
queue table index are as follows:
``
static inline RK_PRIO kCalcNextTaskPrio_(VOID)
{
    if (readyQBitMask == 0U)
    {
        return (idleTaskPrio);
    }
    readyQRightMask = readyQBitMask & -readyQBitMask;
    RK_PRIO prioVal = (RK_PRIO) (__getReadyPrio(readyQRightMask));
    return (prioVal);
}

VOID kSwtch(VOID)
{
    /* O(1) complexity */
    nextTaskPrio = kCalcNextTaskPrio_();

    RK_TCB* nextRunPtr = NULL;

    /* O(1) complexity */
    kTCBQDeq(&readyQueue[nextTaskPrio], &nextRunPtr);

    runPtr = nextRunPtr;

}
```

Chapter 3. Timers and Delays

3.1. Busy delay

A busy-wait delay `kBusyDelay(t)` keeps a task spinning for `t` ticks. That is, the task does nothing but does not suspend or yield (but can be preempted). This service finds its use when simulating workloads.

> Context switching is probably the most significant overhead on a kernel. The time spent on the System Tick handler contributes to much of this overhead.
>
> Design Choice:
>
> - Timers are kept on a single list; only the head element needs to be updated using a delta-queue approach.
> - Application Timers that trigger callbacks are run on a deferred, non-preemptible system task.
>
> Benefits:
>
> - Keep the overhead of updating timers as minimal as possible with the delta queue;
> - Deferring the Application Timer to a high-priority, non-preemptible system task meet the requested callback period while keeping the ability to track system ticks.

Timeout Node
Timeout Type
Absolute Interval (Ticks)
Relative Interval (Ticks)
Waiting Queue Address
Next Timeout Node
Previous Timeout Node

Every task is prone to events triggered by timers described in this section. Every Task Control Block has a node to *a timeout list*. This list is doubly linked treated a delta-sequence.

A set T_{set} `= {(T1,8), (T2,6), (T3,10)}` will be started at a relative time `0` as a sequence T_{seq} `= <(T2,6), (T1,2), (T3,2)>`.

Thus, for every system tick, only the head element on the list needs to be decreased — yielding O(1) on *decreasing*, that happens on the Hardware interrupt for the *System Tick*.

The ordering for the delta-queue is not O(1), it is O(n). A decrease happens on every SysTick interrupt; and the ordering happens only when adding a new node to the list.

3.2. Sleep Timers

There are three `sleep` primitives in RK0 — they behave differently.

3.2.1. Sleep Delay

The `sleepdelay()` (aliased as `sleep()`) puts a task to sleep for the exact number of `t` ticks on every call — no matter when the last call has happened.

Example:

```
VOID Task1(VOID* args)
{
    RK_UNUSEARGS
    UINT count = 0;
    while (1)
    {

        logPost("Task1: sleep");
        kSleep(300);
        /* wake here */
        count += 1U;
        if (count >= 5)
        {
            kDelay(25); /* spin */
            count=0;
            /* every 5 activations there will be a drift */
        }
    }
}
```

Output:

```
0 ms :: Task1: sleep
300 ms :: Task1: sleep  <-- +300
600 ms :: Task1: sleep  <-- +300
900 ms :: Task1: sleep  <-- +300
1200 ms :: Task1: sleep <-- +300
1525 ms :: Task1: sleep <-- +325
1825 ms :: Task1: sleep <-- +300
2125 ms :: Task1: sleep <-- +300
2425 ms :: Task1: sleep
2725 ms :: Task1: sleep
3050 ms :: Task1: sleep
3350 ms :: Task1: sleep
3650 ms :: Task1: sleep
3950 ms :: Task1: sleep
4250 ms :: Task1: sleep
```

```
4575 ms :: Task1: sleep
```

3.2.2. Compensated Sleep Delays

These are suspensions that recompute the time considering the drift between calls. They are typically used to create periodic tasks with explicit periods. The general pattern is:

```
VOID Task(VOID *args)
{
    <initialisation>

    while(1)
    {
        <periodic code>; /* this has an execution time */
        scheduleNext(PERIOD); /* cycle is finished: compute delay to keep PERIOD */
    }
}
```

The RMS algorithm considers all tasks have a common phase grid, that is when they are made READY, or are eligible to be scheduled. This means that from the very first activation, the time elapsed is already take into account to when the next release should happen.

Nevertheless, it is common for kernels to provide sleep primitives that take into account a local anchor time normally set on the `<initialisation>` code — and will work as `waitUntil(&anchor, PERIOD)`. RK0 provides both, and they suit different cases, as exposed below.

3.2.2.1. Periodic Release Sleep

`sleeprelease(P)` is used to delay a task so it is released on periodic rate. The sleep time is recalculated by the kernel on every call.

If the task wakes late by `N` ticks with $0 < N < P$, the kernel compensates by scheduling the next wake earlier (shortening the next sleep) so that over two periods the phase is preserved:

Say a task is expected to return from its k_{eth} sleep at $T_{k+1} = T_k + P$ [ticks]. If the task is resumed at $T_{k+1} = T_k + P + N$, upon detecting this drift, the kernel sets: $(T_{k+2} = T_{k+1} + P - N)$ for $N < P$.

This can be rewritten as:

$(T_{k+2} = T_k + P + N + P - N) \leftrightarrow (T_{k+2} - T_k = 2P)$

Example:

```
VOID Task1(VOID* args)
{
    RK_UNUSEARGS
    UINT count = 0;
    while (1)
    {
```

```
        logPost("Task1 released.");

         count += 1U;
         if (count >= 5)
         {
             kDelay(25); /* spin */
             count=0;
         }

         kSleepRelease(300); /*P=300 ticks; tick=1ms*/


    }
}
```

Output:

```
.
.

/* R is release time */

1200 ms :: Task1: sleep periodic (R==4P)        (n)
1525 ms :: Task1: sleep periodic (5P<R<6P)       |
1800 ms :: Task1: sleep periodic (6P)            |
2100 ms :: Task1: sleep periodic (7P)            |
2400 ms :: Task1: sleep periodic (8P)            |
2700 ms :: Task1: sleep periodic (9P)            |
3025 ms :: Task1: sleep periodic (10P<R<11P)     |
3300 ms :: Task1: sleep periodic (11P)           |
3600 ms :: Task1: sleep periodic (12P)           |
3900 ms :: Task1: sleep periodic (13P)           |
4200 ms :: Task1: sleep periodic (14P)           |
4525 ms :: Task1: sleep periodic (15P<R<16P)     |
4800 ms :: Task1: sleep periodic (16P)          (m)  m-n=12
.                                              -----
.                                           Phase=3600=12xP
.
```

This mechanism is phase-locked. When the lateness is greater or equal to `P` it *skips* one or more releases to stay locked to the phase grid. In some sense, the `period` value can be seen as a *deadline* — if not met, the scheduler rejects to run on that activation.

`sleeprelease()` makes easier to perform worst-response time analysis on periodic tasks.

> A set of periodic tasks must have priorities assigned properly (highest request rate, highest priority — the lower the period, the higher the priority). For `sleeprelease()` this is mandatory given the common phase grid.

3.2.2.2. Sleep Until (local anchor base)

`sleepuntil(anchor, period)` is somehow similar to `sleeprelease()`, but differs in two important aspects:

- The reference used to calculate how long to suspend to keep its rate is local to each task. It means the time before the first time the task is dispatched is dismissed.
- A late release longer than 1 period will return and run immediately. It prioritises execution count within a time-window - not the phase across releases.

The snippet belows clearly demonstrates how each mechanism handles lateness that are longer than 1 period:

```
/* Every 3rd call both tasks will add up a delay longer than the task's period */

VOID HTask(VOID* args) /* higher priority: Period is 300 ticks */
{
    RK_UNUSEARGS
    UINT count = 0;
    while (1)
    {

        logPost("Higher: begin\r\n");
        /* wake here */
        count += 1U;
        kDelay(5);
        if (count >= 3)
        {
            kSleep(400); /* suspend */
            count=0;
        }
        logPost("Higher: end\r\n");
        kSleepRelease(300);

    }
}


VOID LTask(VOID *args) /* lower priority: Period is 400 ticks */
{
    RK_UNUSEARGS
    RK_TICK anchor = kTickGet();
    UINT count=0;
    while (1)
    {

        logPost("Lower: begin\r\n");
        /* wake here */
        count += 1U;
```

```
        kDelay(5);
        if (count >= 3)
        {
            kSleep(500);
            count=0;
        }
        logPost("Lower: end\r\n");
        kSleepUntil(&anchor, 400);
    }
}
```

Output:

```
       0 ms :: Higher: begin
       5 ms :: Higher: end
       5 ms :: Lower: begin
      10 ms :: Lower: end
     300 ms :: Higher: begin
     305 ms :: Higher: end
     405 ms :: Lower: begin
     410 ms :: Lower: end

    /* H 3rd run, expected next at 900ms */
     600 ms :: Higher: begin


   /* L 3rd run, expected next at 1205 ms */
     805 ms :: Lower: begin

    /* H Drift: 1005ms - 600ms = 405 ms > 300ms */
    1005 ms :: Higher end

    /* H released again @ next multiple of 300. */
    1200 ms :: Higher: begin
    1205 ms :: Higher: end

   /* L Drift: 1310ms - 805ms = 505 ms > 400ms */
    1310 ms :: Lower: end
    /* it runs again  immediately */
    1310 ms :: Lower: begin
```

One normally does not write a code with periodic tasks expecting they will not keep their rate. But on the field a transient overload might cause it to happen. If it does, you choose the policy that is best-fit for your task: preserve phase (skip) or preserve execution count.

3.3. Blocking Time-out

These are internal timers associated with kernel calls that are blocking. Thus, establishing an upper-bound waiting time might benefit them. When the time for unblocking is up, the kernel call returns, indicating a timeout. This value is passed as a number of ticks.

When blocking is associated with a kernel object (other than the Task Control Block), the timeout node will store the object waiting for queue's address, so it can be removed if time expires.

A kernel call is made non-blocking, that is *try semantics*, by assigning the value `RK_NO_WAIT`, the function returns immediately if unsuccessful. The value `RK_WAIT_FOREVER` suspends a task indefinitely until the condition is satisfied. Timeout values above `RK_MAX_PERIOD` are invalid.

In practice, we often block either using `RK_WAIT_FOREVER` or do not block (*try semantics*, `RK_NO_WAIT`).

Use a bounded timeout only when you expect occasional misses and you know how to handle them.

> Importantly, *an ISR shall **never** block*. Any blocking call from an ISR is invalid and triggers fault handling when error checking is enabled.

3.4. Callout Timers (Application Timers)

Timer Control Block
Option: Reload/One-Shot
Phase (Initial Delay)
Callout Function Pointer
Callout Argument
Timeout Node

These are Application Timers that will issue a callback when expiring. In addition to a callout function, an Application Timer receives an initial phase delay and a period and can choose to run once (one-shot) or auto-reload itself.

The callback runs within a System Task with priority 0 and is non-preemptible, which makes the scheduler prioritise it over other tasks. Callouts must be short and unblocking, as they can cause high CPU contention.

Application Timers (with autoreload) will keep track of delays in between activations.

3.5. System Tick

A dedicated peripheral that generates an interrupt after a defined period provides the kernel time reference. For ARMv6/7M, this peripheral is the built-in SysTick, a 24-bit counter timer. The handler performs some housekeeping on every tick and assesses the need to call a context switch.

The 'housekeeping' accounts for global timer tracking and any tick-dependent condition that might

change a task status. When a timer expires, it may switch a task from `WAITING` to `READY` or schedule a callback. For callbacks, execution is deferred to the `PostProcSysTask` system task, where callbacks run and timer state is updated.

Note that tasks might switch from `WAITING` to `READY` for reasons other than tick-related. In these cases, context switching might be triggered immediately if the readied task can preempt the running task.

Chapter 4. System Tasks

System Tasks perform housekeeping and other kernel maintenance outside interrupt handlers.

Currently there are the *Idle Task*, the *PostProcSysTask*, and the *SigHndl* task that runs the deferred signal handler path.

The *PostProcSysTask* executes *Application Timer* callbacks and any work deferred from ISRs, such as broadcast-style flushes on Sleep Queues or Semaphores. Because this task runs at priority 0 and is not preemptible by user tasks, the deferred post-processing loop applies a time budget to flush/wake processing: it drains queued flush/wake requests until either the queue is empty or the budget (a fraction of a tick, configurable to match the platform timing constraints) is consumed. Any remaining requests stay queued and are finished on the next dispatch. This amortises large fan-out wakes while keeping ISR paths short and preventing the non-preemptible PostProcSysTask from starving ready user tasks.

> Keep application timer callbacks minimal and non-blocking. Avoid flushing from ISR; enqueue the flush/wake so the PostProcSysTask can process it within its budget.

The *Idle Task* runs whenever there is no other ready task to be dispatched. The CPU enters low-power mode. If `RK_CONF_MIN_PRIO` is set to 31—and it can be even if not all priorities are occupied (with a small memory overhead)—the IdleTask is dispatched when the Ready bitmap is `0x00000000`.

Chapter 5. Memory Allocator

Memory Allocator Control Block
Associated Block Pool
Number of Blocks
Block Size
Number of Free Blocks
Free Block List

The standard C library `malloc()` leads to fragmentation and (also, because of that) is highly indeterministic. Unless we use it once - to allocate memory before starting up, it doesn't fit. But often, we need to 'multiplex' memory amongst tasks over time, that is, to dynamically allocate and deallocate.

To avoid fragmentation, we use fixed-size memory blocks. A simple approach would be a static table marking each block as free or taken. With this pattern, you will need to 'search' for the next available block, if any - the time for searching changes - bounding this search to a maximum number of blocks, or *O(n)*. To optimise, an approach is to keep track of what is free using a dynamic table—a linked list of addresses. Now we have *O(1)*.

We use "meta-data" to initialise the linked list. Every address holds the "next" address value. All addresses are within the range of a pool of fixed-size blocks. This approach limits the minimal size of a block to the size of a memory address—32 bits for our supported architecture.

Yet, this is the cheapest way to store meta-data. If not stored on the empty address itself, an extra 32-bit variable would be needed for each block, so it could have a size of less than 32 bits.

Allocating memory at runtime is a major source of latency (1), indeterministic (2) behaviour, and footprint overhead (3).

Design choice: the allocator's design achieves low-cost, deterministic, fragmentation-free memory management by using fixed-size word-aligned block sizes (1)(2) and embedding metadata within the memory blocks themselves (3).

Benefits: Run-time memory allocation benefits have no real-time drawbacks.

The kernel will always round up the block size to the next multiple of 4. Say the user creates a memory pool, assigning blocks to be 6-byte wide; they will turn into 8-byte blocks.

5.1. How it works

When a routine calls `alloc()`, the address to be returned is the one a "free list" is pointing to, say `addr1`. Before returning `addr1` to the caller, we update the free list to point to the value stored within `addr1` - say `addr8` at that moment.

When a routine calls `free(addr1)`, we overwrite whatever has been written in addr1 with the value-free list point to (if no more `alloc()` were issued, it would still be `addr8`), and `addr1` becomes the free list head again.

Allocating and deallocating fixed-size blocks using this structure and storing meta-data this way is as deterministic (*O(1)*) and economical as we can get for dynamic memory allocation.

Chapter 6. *Inter-Task Communication*: Signals

RK0 Blog:

- About Inter-Task Communication - Part 1
- About Inter-Task Communication - Part 2

Inter-Task Communication (ITC) refers to the mechanisms that enable tasks to coordinate/cooperate/synchronise by means of sending or receiving information that falls into two logical categories: *Signals* or *Messages*.

- ***Signals***: A *Signal is either present or absent*. Its meaning is implicit.
- ***Messages***: A *Message* is a means of coordinating and exchanging information altogether. Different from a *Signal*, a Message needs to be parsed.

Semantics for inter-task communication is a design concern when targeting predicatibility. For this reason, some mechanisms that on the surface can be seen as 'too high-level' for a lean kernel, as PORTS or the MRM Protocol, are offered as primitives because through semantics they enforce policy that meet real-time demands.

In this text when genericallly describing an event-based mechanism, we use the terms `wait` or `sleep` to talk about the operation of checking for an event and blocking if it is not present (for a bounded waiting time or indefenitely); `try` to check for an event without blocking - returning `SUCCESSFUL/UNSUCESSFUL` immediately.

6.1. Task Events (Event Flags)

Within Task Control Block
Event Register (32 flags)
Required Events
Satisfy Condition (ALL/ANY)

Each Task Control Block stores a 32-bit event register (one event per bit). The API is bit-mask oriented:

- `set(taskHandle, mask)`: `mask` is OR'ed to the current event register value of the target task. If the task is waiting for a combination of events, and the result of this operation satisfies, the target task switches to ready.
- `get(required, ANY|ALL, saveAddr, timeout)`: checks if `ANY` or `ALL required` bits are set on the

caller's event register. If true, saves the current event register state on `saveAddr` (if it is non-NULL), before clearing all required flags. Return SUCCESS. If conditions are not met, wait for the specified timeout or immmediately return unsuccesful.

- `query(taskHandle)` reads current flags. No side effects.
- `clear(mask)` clears selected flags.

`0x00` is invalid as a `set()` mask and as `required` for `get()`.

6.1.1. Usage Example: Supervisor Task

One common pattern is to start each loop by checking pending events. In a supervisor task, this becomes an event hub.

```
/* events */
#define PRESSURE_EVENT           RK_EVENT_01
#define TEMP_EVENT               RK_EVENT_02
#define AIRFLOW_EVENT             RK_EVENT_02



/* task handles */
#define TEMPERATURE_TASK         tempCtlHandle
#define PRESSURE_TASK            pressCtl.Handle
#define AIRFLOW_TASK             airCtlHandle


/* highest priority user task */
VOID SupervisorTask(VOID *args)
{
    RK_UNUSEARGS;

    RK_TICK anchor = kGetTick();
    while(1)

    {
        ULONG gotFlags = 0UL;

        RK_ERR err = kTaskEventGet(0xFFFF, /* event range */
                                   RK_FLAGS_ANY, /* condition: any event signalled */
                                   &gotFlags, /*store on gotFlags */
                                   RK_NO_WAIT);

        if (err == RK_ERR_SUCCESS && gotFlags != 0)
        {

            if (gotFlags & PRESSURE_EVENT)
            {
                /* signal pressure handler tasks about this event */
            }
```

```
            if (gotFlags &  TEMP_EVENT)
            {
                /* signal pressure handler tasks about this event */
            }
            /* and so forth */

        }
        /* sleep so tasks signalled will handle events */
        kSleepUntil(&anchor, RK_SUPERVISOR_PERIOD);
    }
}

/* actuator tasks: start pending on event */

VOID ActuatorTask(VOID* args)
{
    RK_UNUSEARGS

    while(1)
    {

      kEventGet(/*wait forever any related events */);

      /* handle each event */

      }
}
```

6.2. Semaphores

Semaphore Control Block
Counter (Unsigned Integer)
Maximum Value
Waiting Queue

Semaphores are *public* kernel objects for signalling and waiting on countable events. Differently from task events, any task can `wait` or `signal` a semaphore.

> A semaphore S is a nonnegative integer variable, apart from the operations it is subjected to. S is initialised to a nonnegative value. The two operations, called P and V, are defined as follows:
>
> `P(S): if S > 0 then S := S-1, else the process is suspended until S > 0.`
>
> `V(S): if there are processes waiting, then one of them is resumed; else S := S+1.`
>
> (Dijkstra, 1968)

`V()` in RK0 semaphores maps to `post()` and `P()` to `pend()`.

6.2.1. Counting Semaphore and Binary Semaphores

The typical use case for semaphores is as a "credit tracker": use `pend()` to consume a credit and `post()` to return a credit (for example, free slots in a queue). These are *Counting Semaphores*.

A *Binary Semaphore* is a counting semaphore with maximum value `1`: the state is either available or unavailable.

Binary semaphores are often used for task-to-task or ISR-to-task synchronisation, and sometimes for mutual exclusion (with caveats discussed later).

6.2.2. Semaphores in RK0

To initialise a semaphore in RK0, provide two values: initial count and maximum count. When the counter is at maximum, `post()` does not increment it and returns `RK_ERR_SEMA_FULL`.

A *Binary Semaphore* is therefore created by setting maximum count to `1`. A counting semaphore that is intended never to saturate can use `UINT32_MAX`.

Besides `init()`, `post()`, and `pend()`, `query()` inspects current state: non-negative means current count; negative means number of tasks waiting.

6.2.3. Flushing or Broadcast

A `flush()` is a *broadcast* signal: it releases all tasks currently pending on the semaphore. If successful, the semaphore counter remains `0`.

> Flushing/broadcasting must be used with care. Releasing many tasks at once can increase scheduling pressure and hurt response time.

Flushing is allowed from ISRs but you should have a good reason. Nevertheless, they are deferred to the *PostProcSysTask* - so they will happen asynchronously.

6.2.3.1. Bounded Buffer with Semaphores

Items are buffered within a memory region whose capacity is `K` items.

Thus: `0 ≤ (Number of Inserted) − (Number of Extracted) ≤ K`.

Using semaphores the pattern is as follows:

1. A semaphore with `K` tokens to track free slots, preventing producers from proceeding when there are no free slots.
2. Another semaphore, with `K` tokens, for the number of items, not allowing consumers to proceed if there are no items.
3. A 1-token semaphore so only one task manipulates the buffer at a time.

Bounded buffer: semaphore roles and critical section

```
/* a ring buffer of items */
#define BUFSIZ (K)
static ITEM_t buf[BUFSIZ]={0};
static UINT getIdx = 0U;
static UINT putIdx = 0U;
/* this indexes==0 could either mean FULL or EMPTY for a regular
circular buffer with wrap-around.
With semaphores the state is well defined.
*/

RK_SEMAPHORE  itemSema;
RK_SEMAPHORE  slotSema;
RK_SEMAPHORE  acquireSema;


VOID kApplicationInit(VOID)
{

  /*buffer is initialised empty */
    kSemaphoreInit
    (   &itemSema,
        0,   /* no item  */
        K    /*max items */
    );

    kSemaphoreInit
    (   &slotSema,
        K, /* K free slots */
        K  /* max slots */
    );

    /* and free */
    kSemaphoreInit
    (   &acquireSema,
        1, /* free to access */
        1  /* 1 max task allowed */
    );


VOID PutItem(ITEM_t* insertItemPtr)
{
    /* wait for room */
    kSemaphorePend(&slotSema, RK_WAIT_FOREVER);

    /* wait for availability */

    kSemaphorePend(&acquireSema,  RK_WAIT_FOREVER);
    buf[putIdx] = *insertItemPtr;
    putIdx += 1U; putIdx %= BUFSIZ;
```

```
    /* signal availability */
    kSemaphorePost(&acquireSema);

    /* signal item */
    kSemaphorePost(&itemSema);
}


 VOID GetItem(ITEM_t* extractItemPtr )
{

    /* wait for an item */
    kSemaphorePend(&item, RK_WAIT_FOREVER);

    /* wait for availability */
    kSemaphorePend(&acquireSema,  RK_WAIT_FOREVER);

    *extractItemPtr = buf[getIdx];
    getIdx+=1U; getIdx %= BUFSIZ;

    /* signal availability */
    kSemaphorePost(&acquireSema);

    /* signal room */
    kSemaphorePost(&slotSema);
}
```

The solution above has Put() and Get() as blocking methods.

If the producer and the consumer run at different rates, eventually, they will synchronise to the lowest rate.

The numbers below are from a run with a buffer of 32 items (integers being incremented are the produced data).

The producer is twice faster than the consumer. Initially at every 2 insertions there is a single remove. Eventually, the buffer is filled up, and tasks run at lockstep (put, get, put, get...), at the consumer pace:

```
Put 59 <-
Put 60 <-
------
Got 30 ->
------
Put 61 <-
Put 62 <-
------
Got 31 ->
------
```

```
Put 63 <-
Put 64 <-
--------
Got 32  | ->
Put 65  . <-
       <x>[Full Queue, Producer blocks]
Got 33  | ->  [Consumer unblocks producer...]
Put 66  . <-
       <x>[Full Queue]
Got 34  | ->  [Consumer unblocks producer...]
Put 67  . <-
       <x>[Full Queue]
```

6.2.4. Mutex Semaphore (Locks)

Mutex Control Block
Locked State (Boolean)
Owner
Protocol Flag (`RK_NO_INHERIT` / `RK_INHERIT`)
Waiting Queue
Mutex Node (list node within the owner TCB)

Some regions are critical and must not be executed by more than one task at once. Acquiring (`lock()`) a mutex before entering and releasing (`unlock()`) after leaving makes the region mutually exclusive.

A *Mutex* is a binary semaphore with ownership: once a task locks a mutex only that task can unlock it.

If a task tries to acquire a locked mutex, it switches to `BLOCKED` until the owner unlocks it. When released, the highest-priority waiter is dequeued first. Unlike semaphores, unlocking by a non-owner is invalid and rejected.

Mutexes are only for mutual exclusion; they are not signalling primitives.

PS: RK0 mutexes are non-recursive. Re-entrant locking of the same mutex returns `RK_ERR_MUTEX_REC_LOCK` *and is considered a fault.*

6.2.4.1. Priority Inversion and PIP

Let TH, TM, and TL be three tasks with priority high (H), medium (M) and low (L), respectively. Say TH is dispatched and blocks on a mutex that 'TL' has acquired (i.e.: *"TL is blocking TH*").

If 'TM' does not need the resource, it will run and preempt 'TL'. And, by transition, 'TH'.

From now on, 'TH' has an *unbounded waiting time* because any task with priority higher than 'L' that does not need the resource indirectly prevents it from being unblocked — *awful.*

The Priority Inheritance (PI) Protocol avoids this unbounded waiting. It is characterised by an invariant, simply put:

> PIP Invariant: *At any instant a Task assumes the highest priority among the tasks it is blocking.*

If employed on the situation described above, task TM cannot preempt TL, whose effective priority would have been raised to 'H'.

It is straightforward to reason about this when you consider the scenario of a single mutex.

When locks nest, the protocol also needs to be:

- Transitive: if T1 blocks T2 and T2 blocks T3, the highest priority (T3) must propagate back to T1 through T2.

This is the hard part of a correct implementation: updates must preserve the invariant across changing wait chains and multiple mutexes.

> This blog shows an even more intricate case of priority inversion handling.
>
> RK0 Blog: Transitive Priority Inheritance on RK0 mutexes

Below, a PIP-only case in which locks nest:

```
/* Task1 has the Highest nominal priority */
/* Task2 has the Medium nominal priority */
/* Task3 has Lowest nominal priority */

/* Note Task3 starts as 1 and 2 are delayed */

RK_DECLARE_TASK(task1Handle, Task1, stack1, STACKSIZE)
RK_DECLARE_TASK(task2Handle, Task2, stack2, STACKSIZE)
RK_DECLARE_TASK(task3Handle, Task3, stack3, STACKSIZE)


RK_MUTEX mutexA;
RK_MUTEX mutexB;

VOID kApplicationInit(VOID)
{
    K_ASSERT(!kCreateTask(&task1Handle, Task1, RK_NO_ARGS, "Task1", stack1, \
        STACKSIZE, 1, RK_PREEMPT));
    K_ASSERT(!kCreateTask(&task2Handle, Task2, RK_NO_ARGS, "Task2", stack2, \
        STACKSIZE, 2, RK_PREEMPT));
    K_ASSERT(!kCreateTask(&task3Handle, Task3, RK_NO_ARGS, "Task3", stack3, \
        STACKSIZE, 3, RK_PREEMPT));
```

```
/* mutexes initialised with priority inheritance enabled */
    kMutexInit(&mutexA, RK_INHERIT);
    kMutexInit(&mutexB, RK_INHERIT);
}



VOID Task3(VOID *args)
{
    RK_UNUSEARGS
    while (1)
    {
        printf("@ %lums: [TL] Attempting to LOCK 'A' | Eff: %d | Nom: %d\r\n",
kTickGet(),
               runPtr->priority, runPtr->prioNominal);

        kMutexLock(&mutexA, RK_WAIT_FOREVER);

        printf("@ %lums: [TL] LOCKED 'A' (in CS) | Eff: %d | Nom: %d\r\n", kTickGet(),
               runPtr->priority, runPtr->prioNominal);

        kDelay(60); /* <-- important */

        printf("@%lums: [TL] About to UNLOCK 'A' | Eff: %d | Nom: %d\r\n", kTickGet(),
               runPtr->priority, runPtr->prioNominal);

        kMutexUnlock(&mutexA);

        printf("--->");
        printf("@%lums: [TL] Exit CS | Eff: %d | Nom: %d\r\n", kTickGet(),
               runPtr->priority, runPtr->prioNominal);

        kSleep(4);
    }
}

VOID Task2(VOID *args)
{
    RK_UNUSEARGS
    while (1)
    {
        kSleep(5);

        printf("@%lums: [TM] Attempting to LOCK 'B' | Eff: %d | Nom: %d\r\n",
kTickGet(),
               runPtr->priority, runPtr->prioNominal);
        kMutexLock(&mutexB, RK_WAIT_FOREVER);

        printf("@%lums: [TM] LOCKED 'B', now trying to LOCK 'A' | Eff: %d | Nom:
%d\r\n",
```

```
                kTickGet(), runPtr->priority, runPtr->prioNominal);
        kMutexLock(&mutexA, RK_WAIT_FOREVER);

        printf("@%lums: [TM] LOCKED 'A' (in CS) | Eff: %d | Nom: %d\r\n", kTickGet(),
                runPtr->priority, runPtr->prioNominal);
        kMutexUnlock(&mutexA);

        printf("@%lums: [TM] UNLOCKING 'B' | Eff: %d | Nom: %d\r\n", kTickGet(),
                runPtr->priority, runPtr->prioNominal);

        kMutexUnlock(&mutexB);

        printf("--->");

        printf("@%lums: [TM] Exit CS | Eff: %d | Nom: %d\r\n", kTickGet(),
                runPtr->priority, runPtr->prioNominal);
    }
}

VOID Task1(VOID *args)
{
    RK_UNUSEARGS
    while (1)
    {
        kSleep(2);

        printf("@%lums: [TH] Attempting to LOCK 'B'| Eff: %d | Nom: %d\r\n",
kTickGet(),
                runPtr->priority, runPtr->prioNominal);

        kMutexLock(&mutexB, RK_WAIT_FOREVER);

        printf("@%lums: [TH] LOCKED 'B' (in CS)  | Eff: %d | Nom: %d\r\n", kTickGet(),
                runPtr->priority, runPtr->prioNominal);

        kMutexUnlock(&mutexB);

        printf("--->");

        printf("@%lums: [TH] Exit CS | Eff: %d | Nom: %d\r\n", kTickGet(),
                runPtr->priority, runPtr->prioNominal);
    }
}
```

Result and comments:

```
>>>> TL locks 'A'. Higher priority tasks are sleeping. <<<<

@ 14720ms: [TL] Attempting to LOCK 'A' | Eff: 3 | Nom: 3
@ 14720ms: [TL] LOCKED 'A' (in CS) | Eff: 3 | Nom: 3
```

```
@14721ms: [TM] Attempting to LOCK 'B' | Eff: 2 | Nom: 2

>>>> TM acquires 'B' and is blocked by TL on 'A'. TL inherits TM's  priority. <<<<

@14721ms: [TM] LOCKED 'B', now trying to LOCK 'A' | Eff: 2 | Nom: 2

>>>> TH will blocked by TM on 'B': <<<<

@14722ms: [TH] Attempting to LOCK 'B'| Eff: 1 | Nom: 1

>>>> TM inherits TH's priority. TL inherits TH's priority via TM. <<<<

@14780ms: [TL] About to UNLOCK 'A' | Eff: 1 | Nom: 3

>>>> Upon unlocking 'A', TL is preempted by TM. It means TL's priority has been
restored, as it is no longer blocking a higher priority task. <<<<


>>>> Now TM acquires 'A' <<<<

@14780ms: [TM] LOCKED 'A' (in CS) | Eff: 1 | Nom: 2

>>>> After releasing 'A', but before releasing 'B', TM's priority is still '1', as it
is blocking TH while holding 'B'. <<<<

@14780ms: [TM] UNLOCKING 'B' | Eff: 1 | Nom: 2

>>>> Upon unlocking 'B' TM is preempted by TH. (TM's priority has been restored.) <<<<

@14780ms: [TH] LOCKED 'B' (in CS)  | Eff: 1 | Nom: 1

>>> RESULT: even though priority inversion was enforced, tasks leave the nested lock
ordered by their nominal priority. <<<

--->@14780ms: [TH] Exit CS | Eff: 1 | Nom: 1
--->@14780ms: [TM] Exit CS | Eff: 2 | Nom: 2
--->@14780ms: [TL] Exit CS | Eff: 3 | Nom: 3
```

Importantly, the worst-case time is bounded by the time the lowest priority task holds a lock (60 ms in the example: 14720ms → 14780ms).

As for each priority update we check each waiting queue for each mutex a task owns, t he time-complexity is linear `O(owner*mutex)`. But, typically no task ever holds more than a few mutexes. Yet, one should not be encouraged to nest locks if not needed.

6.2.5. Mutexes vs Binary Semaphores

There is (or used to be) a lot of fuss about whether binary semaphores are appropriate to use as locks. As a practical guideline, if all tasks sharing the resource have the same priority, using a

binary semaphore *can be appropriate* — because a binary semaphore is considerably faster. It all depends on the case.

The drawback is the lack of ownership: any task can accidentally release the resource. On a large codebase, this can become a real problem. Nonetheless, this is a problem for semaphores in general.

For tasks with different priorities, binary semaphores should never be considered for mutual exclusion unless priority inversion is not a problem (how?).

Counting semaphores initialised as 1 is too risky. Besides the priority inversion, if the count ever increases above 1, mutual exclusion is lost, and multiple tasks can enter the critical section at once.

6.3. Scheduler Lock

Often, we need a task to perform operations without being preempted. A mutex serialises access to a code region but does not prevent a task from being preempted while operating on data. Depending on the case, this can lead to inconsistent data state.

An aggressive way is to disable interrupts globally. For kernel services often it is the only way to keep data integrity. On the higher level it is feasible for very short operations and/or when you need to protect data from interrupts altogether.

A less aggressive approach is to make the task non-preemptible with `kSchLock()` before entering the critical region and `kSchUnlock()` when leaving. This way, interrupts are still being sensed, and even higher-priority tasks might switch to a ready state, but the running thread will not be preempted.

The priority inversion it potentially causes is bounded. If a higher-priority task is readied while the scheduler is locked, the context switch happens immediately after unlocking.

Note that for locking/unlocking the scheduler the global interrupts will be disabled for the time to increment/decrement a counter, therefore, if your atomic operation is as short as that (3 to 4 cycles), disabling/enabling global interrupts is a better alternative.

To add to the discussion, when two threads need to access the same data to 'read-modify-write', a lock-free mechanism is the LDREX/STREX operations of ARMv7M (or more generally C11 atomics). They do not avoid preemptions, and particularly in ARMv7m, if the data is touched by an ISR before the store-exclusive concludes, the ownership is lost. Typically used for multi-core spin-locking.

6.4. Sleep Queue

Sleep Queue Control Block
Task Waiting Queue

`RK_SLEEP_QUEUE` is a queue of tasks sleeping until explicitly woken. This the crudest primitive for

coordinating tasks in *RK0*.

An *RK_SLEEP_QUEUE* is stateless: it does not record whether an event has happened. Therefore a `wait()` always suspends a task, switching it to `RK_SLEEPING`, unless `RK_NO_WAIT` is used — making it a no-operation.

A `signal()` wakes the task with the higher priority that has been enqueued first. ` A `wake(n)` wakes at most `n` tasks if any. This provides some control over the the always questionable broadcast ignal. If `n==0` it is indeed `flush()`.

If broadcasting from an ISR, the operation is deferred for the *PostProcessingTask*.

A `query()` returns the number of waiting tasks. A `suspend()` moves a `READY` task to a sleep queue, and its state switches `SLEEPING_SUSPENDED`. This is done to prevent a task from being scheduled, so Tasks in other states are not allowed as target. A `ready()` does the opposite, moving a specific task from that sleep queue to `READY`.

This primitive was formerly named `RK_EVENT`. Sleep Queue names usually reflect the *condition* they need or the *action* a sensed wake/signal event triggers. (for example, `notFull`, `notEmpty`, `goWriters`, `goReaders`)

Because Sleep Queues are stateless, they are prone to lost wake-ups when used alone. They are best used as building blocks for Conditional Critical Regions (monitor-like constructs).

6.5. Conditional Critical Regions

Task Events and semaphores work by atomically updating state and testing predicates that control execution flow (for example, `pend` on a semaphore with count `0` blocks the caller).

A critical region guarded by a lock is either free or taken. What if we need to wait on a richer condition?

We can combine a Mutex and a Sleep Queue, but there is an important detail: operations on both must be coordinated atomically.

When that combination is wrapped as data plus methods, it is a *Monitor*.

6.5.1. Monitor Invariants

1. a single task can be active within a monitor;
2. only the active task within a monitor can check or change its state.

Since tasks signal or sleep *within* a monitor based on its *internal* state, and only one task can be active, we need to establish how to keep both invariants true.

6.5.2. Signalling Discipline

At any moment a single task can be active within a monitor. When the sleeping task is signalled, there are 3 common disciplines to follow: signal-and-leave (Hansen), signal-and-wait (Hoare) or signal-and-continue (Mesa).

Arguably, the most common is *signal-and-continue* - rather than leaving or suspending itself the active task might continue within the monitor. That is possible if the active task holds a lock the waking task needs to acquire to enter. Upon leaving, the active task must release the lock.

The major implication is that by the time the woken task enters the monitor, the condition it was waiting for might no longer be true. It sounds odd because a Monitor is about encapsulating a conditional critical region so no outsiders change its state. But, either a flush, a bad design — or a preemption anomaly — can violate that somehow.

Mesa Monitor has a typical test-loop pattern:

```
--- snippet ---
while (condition.isFalse())
{
   /*unlock-wait sequence:*/
    DISABLE_PREEMPTION
    unlock(mutex);
    sleep(condition)
    ENABLE_PREEMPTION
    lock(mutex);
    /* when waking, the while clause is tested again */
}
--- snippet ---
```

The figure below shows a producer-consumer problem if implemented on a Monitor-idiom.

```
                        Entry Queue [][][]▯>
        ..............................................
        -------------------
        | Internal Data   |
        -------------------  |[Active Task]|
        | Sleep Queues    |
        |Sleeping for Data|
        |Sleeping for Room|
        -------------------
                             -> [][][] Exit Queue
       ..............................................
```

The `UNLOCK-SLEEP` sequence within the testing loop has preemption disabled because after releasing the lock, the task cannot be allowed to resume within the monitor again, for any reason that is not the monitor predicate being satisfied.

6.5.3. Condition Variable Model

The *Condition Variable Model* allows a task to wait within a monitor-construct and if active, operate using `signal()`, `wait()` and `broadcast()` respecting the monitor invariant.

Sleep Queues are actually the canonical *Condition Variable* introduced by Hoare:

> *Note that a condition "variable" is neither true nor false; indeed, it does not have any stored value accessible to the program. In practice, a condition variable will be represented by an (initially empty) queue of processes which are currently waiting on the condition; but this queue is invisible both to waiters and signallers. This design of the condition variable has been deliberately kept as primitive and rudimentary as possible (...)*
>
> (*Monitors: An Operating System Structuring Concept, Hoare, 1974*)

RK0 follows *pthreads condition variable* semantics, aligned with Mesa monitors.

The difference is that there is no standalone *CondVar* primitive. The programmer combines Sleep Queues and Mutexes and uses these helpers:

- `kCondVarWait(&sleepq, &mutex, timeout)`
- `kCondVarSignal(&sleepq)`
- `kCondVarBroadcast(&sleepq)`

The `condWait` is the real helper. When using it, a Mesa testing-loop reduces to:

```
while(!condition)
{
  kCondVarWait(&condQueue, &monitorLock, timeout);
}
```

If you need a monitor policy different from Mesa, you can build it from the same primitives.

> `kCondVarWait`, `kCondVarSignal`, and `kCondVarBroadcast` are task-context APIs and cannot be called from ISRs.

6.5.3.1. Usage Example: Synchronisation Barrier

A given number of tasks must reach a point in the program before *all can proceed*, so every task calls a `barrWait(&barrier)` to catch up with the set of tasks it must synchronise.

The last task entering the barrier will broadcast a signal to all tasks waiting for the wake condition.

At any moment within a Monitor a single task is RUNNING (what is an invariant of the kernel), all other tasks within the monitor are either SLEEPING (for some condition) or BLOCKED (on a mutex).

```
/* Synchronisation Barrier */
```

```
typedef struct
{
    RK_MUTEX lock;
    RK_SLEEP_QUEUE allSynch;
    UINT count; /* number of tasks in the barrier */
    UINT round; /* increased every time all tasks synch */
    UINT nRequired; /* number of tasks required */
} Barrier_t;

VOID BarrierInit(Barrier_t *const barPtr, UINT nRequired)
{
    kMutexInit(&barPtr->lock, RK_INHERIT);
    kSleepQueueInit(&barPtr->allSynch);
    barPtr->count = 0;
    barPtr->round = 0;
    barPtr->nRequired = nRequired;

}

VOID BarrierWait(Barrier_t *const barPtr)
{
    UINT myRound = 0;
    kMutexLock(&barPtr->lock, RK_WAIT_FOREVER);

    /* save round number */
    myRound = barPtr->round;
    /* increase count on this round */
    barPtr->count++;

    if (barPtr->count == barPtr->nRequired)
    {
        /* reset counter, inc round, broadcast to sleeping tasks */
        barPtr->round++;
        barPtr->count = 0;
        kCondVarBroadcast(&barPtr->allSynch);
    }
    else
    {
        /* sequence: a proper wake signal might happen after inc round */
        while ((UINT)(barPtr->round - myRound) == 0U)
        {
            RK_ERR err = kCondVarWait(&barPtr->allSynch, &barPtr->lock,
RK_WAIT_FOREVER);
            K_ASSERT(err==RK_ERR_SUCCESS);
        }
    }
    kMutexUnlock(&barPtr->lock);
}
```

```
#define N_REQUIRED 3

Barrier_t syncBarrier;

VOID kApplicationInit(VOID)
{

    K_ASSERT(!kCreateTask(&task1Handle, Task1, RK_NO_ARGS, "Task1", stack1, STACKSIZE,
2, RK_PREEMPT));
    K_ASSERT(!kCreateTask(&task2Handle, Task2, RK_NO_ARGS, "Task2", stack2, STACKSIZE,
3, RK_PREEMPT));
    K_ASSERT(!kCreateTask(&task3Handle, Task3, RK_NO_ARGS, "Task3", stack3, STACKSIZE,
1, RK_PREEMPT));
    BarrierInit(&syncBarrier, N_REQUIRED);
}
VOID Task1(VOID* args)
{
    RK_UNUSEARGS
    while (1)
    {
        kPuts("Task 1 is waiting at the barrier...\r\n");
        BarrierWait(&syncBarrier);
        kPuts("Task 1 passed the barrier!\r\n");
        kSleep(8);
    }
}

VOID Task2(VOID* args)
{
    RK_UNUSEARGS
    while (1)
    {
        kPuts("Task 2 is waiting at the barrier...\r\n");
        BarrierWait(&syncBarrier);
        kPuts("Task 2 passed the barrier!\r\n");
        kSleep(5);
    }
}

VOID Task3(VOID* args)
{
    RK_UNUSEARGS
    while (1)
    {
        kPuts("Task 3 is waiting at the barrier...\r\n");
        BarrierWait(&syncBarrier);
        kPuts("Task 3 passed the barrier!\r\n");
        kSleep(3);
    }
}
```

```
Task 3 is waiting at the barrier...
Task 2 is waiting at the barrier...
Task 1 is waiting at the barrier...
Task 3 passed the barrier!
Task 1 passed the barrier!
Task 2 passed the barrier!
Task 3 is waiting at the barrier...
Task 2 is waiting at the barrier...
Task 1 is waiting at the barrier...
Task 3 passed the barrier!
Task 1 passed the barrier!
Task 2 passed the barrier!
Task 3 is waiting at the barrier...
Task 2 is waiting at the barrier...
Task 1 is waiting at the barrier...
Task 3 passed the barrier!
Task 1 passed the barrier!
Task 2 passed the barrier!
Task 3 is waiting at the barrier...
Task 2 is waiting at the barrier...
Task 1 is waiting at the barrier...
Task 3 passed the barrier!
Task 1 passed the barrier!
Task 2 passed the barrier!
```

Monitor/CondVars are not to be signalled by ISRs. Because as they do not accumulate and an ISR is not a normal task, it can't be within a monitor and a signal might be lost.

6.5.3.2. Usage Example: Readers Writers Lock

Several readers and writers share a piece of memory. Readers can concurrently access the memory to read; a single writer is allowed (otherwise, data would be corrupted).

When a writer finishes, it checks for any readers waiting. If there is, the writer flushes the readers waiting queue. If not, it wakes a single writer, if any. When the last reader finishes, it signals a writer.

Every read or write operation begins with an acquire and finishes with a release.

*PS: This RWLock implementation has a **reader-preference** policy, as when a writer finishes, it flushes sleeping readers. When the last reader finishes, it will signal writer waiting queue.*

```
/* RW-Lock */

/* a single writer is allowed if there are no readers */
/* several readers are allowed if there is no writer*/
typedef struct
{
    RK_MUTEX     lock;
    RK_SLEEP_QUEUE   writersGo;
    RK_SLEEP_QUEUE   readersGo;
    INT          rwCount; /* number of active readers if > 0 */
                          /* active writer if -1             */
```

```
}RwLock_t;

VOID RwLockInit(RwLock_t *const rwLockPtr)
{

    kMutexInit(&rwLockPtr->lock, RK_INHERIT);
    kSleepQueueInit(&rwLockPtr->writersGo);
    kSleepQueueInit(&rwLockPtr->readersGo);
    rwLockPtr->rwCount = 0;
}

/* A writer can acquire if  rwCount = 0 */
/* An active writer is indicated by rwCount = -1; */
VOID RwLockAcquireWrite(RwLock_t *const rwLockPtr)
{
    kMutexLock(&rwLockPtr->lock, RK_WAIT_FOREVER);
    /* if different than 0, there are either writers or readers */
    /* sleep to be signalled */
    while (rwLockPtr->rwCount != 0)
    {
        kCondVarWait(&rwLockPtr->writersGo, &rwLockPtr->lock, RK_WAIT_FOREVER);
        /* mutex is locked when waking up*/
    }
    /* woke here, set an active writer */
    rwLockPtr->rwCount = -1;
    kMutexUnlock(&rwLockPtr->lock);
}

/* a writer releases, waking up all waiting readers, if any */
/* if there are no readers, a writer can get in */
VOID RwLockReleaseWrite(RwLock_t *const rwLockPtr)
{
    kMutexLock(&rwLockPtr->lock, RK_WAIT_FOREVER);

    rwLockPtr->rwCount = 0; /* indicate no writers*/

    /* if there are waiting readers, flush */
    ULONG nWaitingReaders=0;
    kSleepQueueQuery(&rwLockPtr->readersGo, &nWaitingReaders);
    if (nWaitingReaders > 0)
    {
        /* condVarBroadcast is just an alias for an event flush */
        kCondVarBroadcast(&rwLockPtr->readersGo);
    }
    else
    {
        /* wake up a single writer if any */
        kCondVarSignal(&rwLockPtr->writersGo);
    }
    kMutexUnlock(&rwLockPtr->lock);
}
```

```
/* a reader can acquire if there are no writers */
VOID RwLockAcquireRead(RwLock_t *const rwLockPtr)
{
    kMutexLock(&rwLockPtr->lock, RK_WAIT_FOREVER);
    /* if there is an active writer, sleep */
    while (rwLockPtr->rwCount < 0)
    {
        kCondVarWait(&rwLockPtr->readersGo, &rwLockPtr->lock, RK_WAIT_FOREVER);
        /* mutex is locked when waking up*/
    }
    /* increase rwCount, so its > 0, indicating readers */
    rwLockPtr->rwCount ++;
    kMutexUnlock(&rwLockPtr->lock);
}

/* a reader releases and wakes a single writer */
/* if it is the last reader */
VOID RwLockReleaseRead(RwLock_t *const rwLockPtr)
{
    kMutexLock(&rwLockPtr->lock, RK_WAIT_FOREVER);
    rwLockPtr->rwCount --;
    if (rwLockPtr->rwCount == 0)
    {
        kCondVarSignal(&rwLockPtr->writersGo);
    }
    kMutexUnlock(&rwLockPtr->lock);
}
```

In the image below, 4 tasks — a fast writer (Task 1), a slow writer (Task 4) and two readers (Task3 is faster than Task2) — reading from and writing to a shared UINT variable:

```
15340ms : Task 3 read 73
15390ms : Task 4 wrote 75
15500ms : Task 3 read 75
15550ms : Task 2 read 75
15620ms : Task 1 wrote 76
15750ms : Task 3 read 76
15910ms : Task 2 read 76
16000ms : Task 1 wrote 77
16100ms : Task 3 read 77
16200ms : Task 4 wrote 79
Task4 is flushing readers
16300ms : Task 3 read 79
16350ms : Task 2 read 79
16420ms : Task 1 wrote 80
16520ms : Task 3 read 80
16570ms : Task 4 wrote 82
Task4 is flushing readers
16670ms : Task 2 read 82
16750ms : Task 3 read 82
16800ms : Task 1 wrote 83
```

6.6. Deferred Signals

Per-task Deferred Signal Record
Array of handler callbacks indexed by signal ID.
FIFO queue of `(id, sigInfo)` records

Deferred Task Signals (DSignals) are queued, task-directed notifications carrying:

- `id` — numeric signal identifier (1..`RK_DSIGNAL_USER_ID_MAX` for user handlers). Kernel uses reserved IDs above `RK_DSIGNAL_USER_ID_MAX`, such as `RK_DSIGNAL_SYS_EXCHG_CANCEL`, for internal purposes.
- `sigInfo` — optional payload:

Despite the terminology similarity with BSD/UNIX signals, a DSignal does not ressemble a software interrupt.

This is particularly useful to information a task needs to know whenever it is resumed — like a side-effect from the operation of another task (e.g., TaskA has changed parameter that needs TaskB to execute some procedure again).

This provides:

- A single, explicit delivery point.
- No asynchronous execution paths inside a running task.
- Bounded and analysable pre-run overhead.

The delivery model is therefore predictable and compatible with real-time analysis.

6.6.1. Use cases

DSignals are appropriate when:

- A task must notify another task of a side effect.
- The receiver must observe the notification at its next scheduling boundary.
- No ownership transfer or protocol interaction is required.

6.6.2. Functional overview

- DSignals do not pre-empt RUNNING tasks.
- Pending signals are dispatched when the task transitions from READY to RUNNING.
- Dispatch occurs on the system signal-handler task path at the receiver's priority.
- Queue depth is fixed at compile time, FIFO policy
- If the queue is full, `kDSignal()` fails and returns `RK_ERR_SIGNALQ_FULL`.

- Handlers must be short, non-blocking, and must not yield.
- DSignals are not a general message-passing mechanism — use Ports or Message Queues for ordered payload transfer or protocol semantics.

6.6.3. Usage Pattern

```
/* Declare record + handler table  */
RK_DECLARE_TASK_DSIGNAL(taskbDSig, 32);

static volatile RK_BOOL configDirty = RK_FALSE;

static volatile ULONG configTime = DEFAULT; /*defined elsewhere */

#define RK_SIG_CONFIG_CHANGED   1U

/* this handler will hook when TaskB is selected to be dispatched and DSignalID is
pending */
static VOID onConfigChanged(RK_DSIGNAL_ID id,
                            RK_DSIGNAL_INFO info)
{
    /* Minimal side effect only */
    configDirty = RK_TRUE;
    configTime = info.val;
}

VOID ApplicationInit(VOID)
{
    /* Initialise record (clears queue and installs default NOP handlers) */
    kDSignalInit(&taskbDSig,
                 taskbDSigHandlers,
                 RK_CONF_DSIGNAL_QUEUE_SIZE);

    /* Attach record to TaskB */
    kDSignalAttachTask(taskBhandle, &taskbDSig);

    /* Register handler for specific signal ID */
    kDSignalRegisterHandler(&taskbDSig,
                        RK_SIG_CONFIG_CHANGED,
                            onConfigChanged);
}

/* Sender */
TaskA:
--snip--

    UINT ret = updateState();
    if (ret == ERROR_NEED_RECONFIG_THING)
    {
        /* signal and continue */
        kDSignal(taskBHandle,
```

```
                  RK_SIG_CONFIG_CHANGED,
RK_DSIGNAL_VAL_FROM_ULONG(THING_RECONFIG_TIME)); /*time defined elsewhere */
    }
    moreWork():
--snip--

/* Receiver */
TaskB:

    ---snip---

        if (configDirty)
        {
            configDirty = RK_FALSE;
            reconfig();
            kSleep(THING_RECONFIG_TIME);
);
        }
    ---snip---
```

6.6.4. Minimal kDSignal example

The smallest end-to-end flow is: attach a per-task signal record, register a handler, send a signal, let the scheduler deliver it, and read the payload inside the handler.

```
/* 1) Declare record + handler table for one task (queue depth = 4). */
RK_DECLARE_TASK_DSIGNAL(workerDSig, 4);

/* 2) Handler toggles a flag and reads the payload. */
static volatile RK_BOOL gTickFlag = RK_FALSE;
static ULONG gLastTick = 0;

static VOID onTick(RK_DSIGNAL_ID id, RK_DSIGNAL_INFO info)
{
    (void)id;
    gTickFlag = RK_TRUE;
    gLastTick = info.val; /* 32-bit payload */
}

/* 3) Init/attach/register once at startup. */
VOID InitSignals(VOID)
{
    kDSignalInit(&workerDSig, workerDSigHandlers, 4);
    kDSignalAttachTask(workerHandle, &workerDSig);
    kDSignalRegisterHandler(&workerDSig, 1U, onTick);
}

/* 4) Sender posts a deferred signal. */
VOID NotifierTask(VOID *args)
```

```
{
    RK_UNUSEARGS;
    while (1)
    {
        kDSignal(workerHandle, 1U, RK_DSIGNAL_VAL_FROM_ULONG(kGetTick()));
        kSleep(10U);
    }
}

/* 5) Receiver observes the handler effect when it next runs. */
VOID WorkerTask(VOID *args)
{
    RK_UNUSEARGS;
    while(1)
    {
        if (gTickFlag)
        {
            gTickFlag = RK_FALSE;
            consumeTick(gLastTick);
        }
        doWork();
    }
}
```

6.6.4.1. Separation from Message Passing

Ports and Message Queues are intended for structured communication, ownership transfer, and protocol sequencing.

DSignals intentionally do not provide:

- Ordered payload transport beyond a small bounded FIFO.
- Request–reply semantics.
- Ownership transfer guarantees.
- Flow control.

Their payload is limited to a small union (`ULONG` or `void*`) to prevent misuse as a transport mechanism.

Applications requiring structured interaction must use Ports or Message Queues.

6.6.4.2. Separation from Task Events

Task Events represent state (present/absent) and are typically used with explicit waiting primitives.

DSignals differ in that:

- They preserve multiplicity (multiple signals can be queued).
- They carry a small payload.

- They do not require the receiver to pre-arm a wait operation.
- They are guaranteed to be observed at the next scheduling boundary.

This makes DSignals appropriate for lightweight, deferred side-effect notification.

6.6.4.3. Bounded and Minimal by Design

DSignal queue depth is fixed at compile time. Dispatch order is strictly FIFO. Handler execution is constrained to short, non-blocking operations.

These constraints ensure:

- Bounded memory usage.
- Bounded dispatch latency.
- No hidden scheduling side effects.

Chapter 7. Inter-Task Communication: Messages

So far we have been discussing shared-state inter-task communications. *Events* are on the form of *Signals* that, defined by being 'present/absent'.

A *Message* is defined by its presence and the information it conveys. As taskA can only receive a message *after* a taskB has sent it, message-passing is a means of coordinating/synchronising tasks.

In *RK0* we distinguih *Mails* from *Messages* willing for clarity. A *Mail* is a 32-bit message. It is often a pointer but not always.

7.1. Task Mailbox

Within a Task Control Block
Pointer to a message (`VOID*`)

7.1.1. Description

In RK0 a Task Mailbox is always present: a single pointer slot (`VOID *`) inside each TCB. The slot is either EMPTY (`NULL`) or FULL (holds one pointer). Writers overwrite any unread pointer—*last message wins*. Only the owning task may receive; any task may send.

How to use it:

- Sender: drop a pointer into the target's mailbox; the call never blocks and simply replaces stale mail.
- Receiver: reads own mailbox. If empty, block with a suitable timeout or returns unsuccesful. On success the pointer is returned and the slot is cleared.

7.1.2. Suitable usage patterns

A Task Mailbox is a single-slot inbox, not a queue. It works best when the receiver only cares about the newest pointer, or when a higher-level protocol guarantees that only one message can be outstanding at a time.

A mailbox is not suitable when every message must be preserved. In those cases the Message Queues are to be used, or another RYO buffering

Below several usage examples are demonstrated to show how the simplicity of the core mechanisms combined is enough to cover many real use-cases.

7.1.2.1. Latest-snapshot handoff

The producer updates a shared snapshot and mails its address to the consumer. If the consumer falls behind, stale snapshots are naturally dropped and only the latest one remains.

```
typedef struct
{
    FLOAT x;
    FLOAT y;
    FLOAT z;
} Position_t;

static volatile Position_t g_lastPos;

/* Producer posts newest sample; never blocks and overwrites if reader is late */
VOID SensorTask(VOID *args)
{
    RK_UNUSEARGS;

    while (1)
    {
        HAL_getPosition((Position_t *)&g_lastPos);
        kMailSend(consumerHandle, (VOID *)&g_lastPos);
        /* wait for next sample */
    }
}

/* Consumer blocks for the newest snapshot */
VOID Consumer(VOID *args)
{
    RK_UNUSEARGS;
    Position_t *recvPtr = NULL;
    Position_t  sample;

    while (1)
    {
        if (kMailRecv((VOID **)&recvPtr, RK_WAIT_FOREVER) == RK_ERR_SUCCESS)
        {
            sample = *recvPtr;   /* copy immediately */
            processSample(&sample);
        }
    }
}
```

7.1.2.2. Command/response using Mailbox and Events (synchronous)

This pattern is useful for a lightweight RPC-style exchange between one client and one server. The client allocates a command block, sends its pointer to the server, then waits for a response event. The server processes the command and returns the same pointer through the client mailbox.

```
#define EV_RESP RK_EVENT_1

typedef ULONG (*CmdOp_t)(ULONG arg1, ULONG arg2);
```

```
typedef struct
{
    CmdOp_t opFunc;
    ULONG   arg1;
    ULONG   arg2;
    ULONG   result;
} CMD_t;

CMD_t *fillCmd(ULONG arg1, ULONG arg2, CmdOp_t func)
{
    CMD_t *cmd = kMemAlloc(cmdMemPool);

    if (cmd != NULL)
    {
        cmd->opFunc = func;
        cmd->arg1   = arg1;
        cmd->arg2   = arg2;
        cmd->result = 0;
    }

    return cmd;
}

/* Client task */
VOID Client(VOID *args)
{
    RK_UNUSEARGS;
    CMD_t *cmd     = NULL;
    CMD_t *respPtr = NULL;

    while (1)
    {
        cmd = fillCmd(10u, 20u, addOp);
        if (cmd == NULL)
        {
            assert(0);
        }

        kMailSend(serverHandle, cmd);

        /* Wait for reply notification, then pull reply from own mailbox */
        kEventGet(EV_RESP, RK_EVENT_FLAGS_ANY, NULL, RK_WAIT_FOREVER);

        if (kMailRecv((VOID **)&respPtr, RK_WAIT_FOREVER) == RK_ERR_SUCCESS)
        {
            useResult(respPtr->result);
            kMemFree(cmdMemPool, respPtr);
        }
    }
}
```

```
/* Server task */
VOID Server(VOID *args)
{
    RK_UNUSEARGS;
    CMD_t *recvPtr = NULL;

    while (1)
    {
        if ((kMailRecv((VOID **)&recvPtr, RK_WAIT_FOREVER) == RK_ERR_SUCCESS) &&
            (recvPtr != NULL))
        {
            recvPtr->result = recvPtr->opFunc(recvPtr->arg1, recvPtr->arg2);

            kMailSend(clientHandle, recvPtr);
            kEventSet(clientHandle, EV_RESP);
        }
    }
}
''''
Use this only when at most one request can be pending in the server mailbox.
If several clients may post concurrently, their requests can overwrite each
other unless access is serialized externally.

==== Latest-snapshot handoff (one-way)

This pattern is ideal when only the newest sample matters. The producer updates
a shared snapshot and mails its address to the consumer. If the consumer falls
behind, stale snapshots are naturally dropped and only the latest one remains.

'''c
typedef struct
{
    FLOAT x;
    FLOAT y;
    FLOAT z;
} Position_t;

static volatile Position_t g_lastPos;

/* Producer posts newest sample; never blocks and overwrites if reader is late */
VOID SensorTask(VOID *args)
{
    RK_UNUSEARGS;

    while (1)
    {
        /* blocks until an ISR  make the position avaiable */
        kEventGet(RK_EVENT_POSITION_AVAIABLE, RL_WAIT_FOREVER);
        kMailSend(consumerHandle, (VOID *)&g_lastPos);
    }
}
```

```
/* Consumer blocks for the newest snapshot */
VOID Consumer(VOID *args)
{
    RK_UNUSEARGS;
    Position_t *recvPtr = NULL;
    Position_t  sample;

    while (1)
    {
        if (kMailRecv((VOID **)&recvPtr, RK_WAIT_FOREVER) == RK_ERR_SUCCESS)
        {
            sample = *recvPtr;   /* copy immediately */
            processSample(&sample);
        }
    }
}
```

7.1.2.3. Double-buffering

In last example, producer still owns the storage. The consumer should copy the data immediately if the producer may update it again. If the writer can modify the object while the reader is copying, use double-buffering. We keep a `WRITE-HERE][READ-HERE]` and swap:

```
typedef struct
{
    FLOAT x;
    FLOAT y;
    FLOAT z;
} Position_t;

static Position_t    g_lastPos[2];
static volatile UINT g_pubPos = 1U;   /* most recently completed buffer */
static volatile UINT g_wrtPos = 0U;   /* buffer ISR writes next */

VOID ISR(VOID)
{

    UINT w = g_wrtPos;                 /* capture once */

    HAL_getPosition(&g_lastPos[w]);    /* fill complete snapshot */

    RK_DMB                          /* data visible before publish */
    g_pubPos = w;                      /* publish just-written buffer */
    g_wrtPos = w ^ 1U;                 /* swap */

    kEventSet(sensorTask, RK_EVENT_POSITION_AVAIABLE);
}

/* Producer posts newest sample; overwrites old pointer if reader is late */
```

```
VOID SensorTask(VOID *args)
{
    RK_UNUSEARGS;

    while (1)
    {
        UINT bufpos;

        kEventGet(RK_EVENT_POSITION_AVAIABLE, RK_WAIT_FOREVER);

        bufpos = g_pubPos;

        kMailSend(consumerHandle, (VOID *)&g_lastPos[bufpos]);
    }
}

/* Consumer blocks for the newest snapshot */
VOID Consumer(VOID *args)
{
    RK_UNUSEARGS;
    Position_t *recvPtr = NULL;
    Position_t  sample;

    while (1)
    {
        if (kMailRecv((VOID **)&recvPtr, RK_WAIT_FOREVER) == RK_ERR_SUCCESS)
        {
            sample = *recvPtr;   /* copy immediately */
            processSample(&sample);
        }
    }
}
```

7.1.2.4. Coalesced updates

Small integers can be passed as pointer casts when both sides agree on the encoding.

```
typedef enum
{
    MODE_IDLE = 0,
    MODE_RUN  = 1,
    MODE_CAL  = 2
} Mode_t;

static VOID postMode(Mode_t mode)
{
    kMailSend(controlTaskHandle, (VOID *)(ULONG)mode);
}

VOID ControlTask(VOID *args)
```

```
{
    RK_UNUSEARGS;
    VOID  *raw  = NULL;
    Mode_t mode = MODE_IDLE;

    while (1)
    {
        if (kMailRecv((VOID **)&raw, RK_WAIT_FOREVER) == RK_ERR_SUCCESS)
        {
            mode = (Mode_t)(ULONG)raw;
            applyMode(mode);
        }
    }
}
```

It is not suitable for edge-triggered commands where every transition must be observed.

7.1.2.5. Producer allocates - consumer frees

This pattern transfers ownership of a partition-allocated object from one task to another. *With a Mailbox is suitable when only one item may be in flight at a time.*

```
#define EV_BUF_DONE RK_EVENT_2

struct databuf
{
    UCHAR data[128];
    ULONG len;
} K_ALIGN(4);

typedef struct databuf Buffer_t;

/* Producer gives ownership to consumer */
VOID Producer(VOID *args)
{
    RK_UNUSEARGS;
    Buffer_t *buf = NULL;

    while (1)
    {
        buf = kMemAlloc(bufPart);
        if (buf == NULL)
        {
            continue;
        }

        fillBuffer(buf);

        kMailSend(consumerHandle, buf);
```

```
        /* Do not send another until the previous buffer is consumed */
        kEventGet(EV_BUF_DONE, RK_EVENT_FLAGS_ANY, NULL, RK_WAIT_FOREVER);
    }
}

/* Consumer becomes owner after receive */
VOID Consumer(VOID *args)
{
    RK_UNUSEARGS;
    Buffer_t *buf = NULL;

    while (1)
    {
        if (kMailRecv((VOID **)&buf, RK_WAIT_FOREVER) == RK_ERR_SUCCESS)
        {
            consumeBuffer(buf);
            kMemFree(bufPart, buf);
            kEventSet(producerHandle, EV_BUF_DONE);
        }
    }
}
```

Do not use a Task Mailbox when:

- every message must be preserved
- several producers may post independent work items to one task
- the receiver must count notifications rather than observe the latest state
- multiple outstanding requests must be handled concurrently

Message Queues are the straight mechanism. You can consider creating your own ring buffer along with semaphores as well.

7.1.3. Message Queue

Message Queue Control Block
Buffer Address
Message Size
Number of Mesages
Write Position
Read Position
Server Status
Owner Task
Notify callback

Message Queue Control Block
Waiting queue

Message Queues (`RK_MESG_QUEUE`) are public message-passing kernel objects. Each messsage queue has a back storage holding *N* messages of a fixed-size *S*. We say that `C = N x S [Words]` is the queue capacity.

Messages Queues transmit *by copy*. Sending and Receiving are either blocking or non-blocking. When *N=S=1* we have a public `RK_MAILBOX`.

7.1.3.1. Size of a Message

Each declared queue has a *fixed message-size* at initialisation, and can assume, 1, 2, 4 or 8 WORDs (4, 8, 16, 32 BYTEs). This constraint is intentional. Word-aligned copies are faster, predictable and safer for type casting.

(A word-aligned *single copy* will take ~5 cycles in Cortex-M3/4/7, and ~6 cycles on Cortex-M0/M0+.)

7.1.3.2. Public Mailbox

When *N=1* and *S=1* we have a public *Mailbox*. (`RK_MAILBOX`). Note that a `RK_MAILBOX` pass messages by copy, and behave as a message queue of capacity 1.

Public Mailboxes can block on both `post()` and `pend()` operations, differently from Task Mailboxes that only block when sending.

Like *Task Mailboxes* they will also accept last-message semantics by the operation `postovw()`. Since they are public one often wants that the counterpart that is receiving does not clear the mailbox, so a a `peek()` operation is used.

7.1.3.3. Ports

Ports (`RK_PORTS) are for fully-synchronous *unbuffered* message-passing.

have messages whose types are labeled as:

- `RK_PORT_MESG_<0,2,4>WORD` — the number indicates the *payload*. In case of '0' the message contains only meta-data/

7.1.3.4. Blocking and non-blocking behaviour

A producer task can optionally block on a full queue, switching its state to `SENDING`. A consumer (optionally) blocks on an empty queue switching its state to `RECEIVING`.

Note that when using non-blocking calls (`RK_NO_WAIT`) they immediately return if unsuccessful.

One can use a `peek()` to read from the head of a queue without extracting the message. A `jam()` method is used to put a message on the queue head, but it does not overwrite. `jam()` is meaningless for Mailboxes, since they a have a single slot.

A `postovw()` overwrites the oldest message and is allowed solely for 1-message queues. If N-message

queues were allowed to be overwritten, all unread messages would leak as read and write pointers are adjusted on the operation. In case one needs overwriting full queues continuously a classic ring buffer will do.

`peek()` and `postovw()` are normally used on Mailboxes for last-message semantics — the Mailbox never goes empty; a new message is placed by overwriting the current one.

7.1.3.5. Ownership and Priority Inversion

Message Queues (and by extension, Mailboxes) *can* have owners, and thus, only the owners can receive at these channels. One might think it makes it a *Port*, but not yet. The fact is, if a receiver is blocked on full queue, the sender can boost its priority.

Ports are for _ synchronous procedure calls_. The server needs to run at the client's priority, which means either boosting or demoting priority.

7.1.3.6. Notify callback

A callback can be registered for when a queue sends a message successfuly, as a means of event notification.

7.1.3.7. Usage Examples

Below there is a rich set of usage examples. They demonstrate how to use message-passing in RK0, its API and helper macros to declare storage and messages at the appropriate size.

7.1.3.7.1. Mail Queue pattern

A mail queue is the strongest pattern for message-passing.

It is done by combining a Memory Partition and 1-word-size message-queue, with a `N>1`.

Sender allocates, receiver frees the partition. This keep integrity and low overhead — data is 'moved' between tasks.

```
/* Application Logger pattern */

/* standard log structure */
struct log
{
    RK_TICK t; /* timestamp */
    CHAR s[LOGLEN]; /*formatted string */
    UINT    level; /* level 0=message, 1=fault */
} K_ALIGN(4);

typedef struct log Log_t;

/* logger mem allocator + mem pool */
static RK_MEM_PARTITION qMem;
static Log_t logBufPool[LOGPOOLSIZ] K_ALIGN(4);
```

```
/* backing buffer for the logger queue */
/* (messages are 1-word-size, number equals the pool) */
RK_DECLARE_MESG_QUEUE_BUF(logQBuf, VOID *, LOGPOOLSIZ)

/* logger mail queue */
static RK_MESG_QUEUE logQ;

/* a sender will allocate a buffer write the log
message and enqueue on the mail queue, not-blocking
if the queue is full it returns the buffer immediately
if the memory pool is empty it drops the operation */

```c
/* excerpt of logPost(...) */
VOID logPost(/*formatted string */)
    ---snip---
    Log_t *logPtr = (Log_t*)kMemPartitionAlloc(&qMem);
    RK_BARRIER
    if (logPtr) /* available buffer */
    {

       < fill the buffer >

        /* enqueue address */
        if (kMesgQueueSend(&logQ, &p, RK_NO_WAIT) != RK_ERR_SUCCESS)
        {
            /* queue is full, deallocate buf */
            RK_ERR err = kMemPartitionFree(&qMem, &p);
            K_ASSERT(err==RK_ERR_SUCCESS);
        }

    }
    ---snip---

/* excerpt of the logger task */

static VOID LoggerTask(VOID *args)
{
    RK_UNUSEARGS
    while (1)
    {

        VOID *recvPtr = NULL;

        /* drain the queue: keep receiving while successful */
        while (kMesgQueueRecv(&logQ, &recvPtr, RK_WAIT_FOREVER) == RK_ERR_SUCCESS)
        {
            < print buffer contents >

            /* deallocate */
            RK_ERR err = kMemPartitionFree(&qMem, recvPtr);
```
```

```
            K_ASSERT(err == RK_ERR_SUCCESS);

        }
    }
}
```

This is one of the many ways of using the Mail Queue pattern.

The entire implementation can be seen at `app\logger.c`.

7.1.3.7.2. Averaging Sensor Values

A task receives measured sensor values from an ISR on a periodic rate. (A Soft Timer emulates the ISR).

Then it enqueues this data to a consumer - that will process the average value for each of 4 sensors.

Message queues (`RK_MESG_QUEUE`) move data by copy between tasks. Here is the pattern we use in this example:

1. A periodic interrupt (emulated by a timer) posts raw samples into a mailbox.
2. A high-priority producer task pends on that mailbox, then enqueues the sample into a message queue consumed by a lower-priority task.
3. If the queue fills, the producer drops the latest sample and signals the consumer.
4. The consumer drains the queue, computes averages, and blocks on an event flag when the queue is empty.

```
/* Task2 blocks on a Task Event flag (shared state) when the queue is empty. */

/* sensor type */
typedef enum
{
    TEMPERATURE=1, HUMIDITY, CO2, FLOW
}SensorType_t;

struct sensorMsg
{
    SensorType_t sensorType;
    ULONG sensorValue;

};
typedef struct sensorMsg Mesg_t;

#define N_SENSOR    4
#define AVG_WINDOW_SIZE   10 /* 10 samples */
#define BATCH_READY RK_EVENT_1

/* the queue */
RK_MESG_QUEUE sensorStream;
```

```
/* convenience macro to declare the queue storage */
#define N_MESSAGE 8
RK_DECLARE_MESGQ_BUF(mesgBuf, Mesg_t, N_MESSAGE)

/* timer to mimic isr */
RK_TIMER timerT1;

/* the mailbox the sensor task pends */
RK_MAILBOX sensorBox;

/* callback fwded declaration */
VOID callBackISR(VOID* args);

VOID kApplicationInit( VOID)
{
    K_ASSERT(!kMesgQueueInit(&sensorStream,
                            mesgBuf,
                             RK_MESGQUEUE_MESG_SIZE(Mesg_t), /* set mesg size */
                             N_MESSAGE));

    /* timer @ every 10 ms */
    K_ASSERT(!kTimerInit(&timerT1, 0, 10, callBackISR, NULL, RK_TIMER_RELOAD));
    K_ASSERT(!kMailboxInit(&sensorBox));
}


VOID callBackISR(VOID *args)
{

    static Mesg_t sample = {0};
    static Mesg_t *samplePtr = &sample;
    static UINT sampleErr;

    /* fill sample */
    sample.sensorType = (rand() % 4) + 1;
    switch (sample.sensorType)
    {
        case TEMPERATURE:
            sample.sensorValue = ( ULONG) rand() % 50;
            break;
        case HUMIDITY:
            sample.sensorValue = ( ULONG) rand() % 100;
            break;
        case CO2:
            sample.sensorValue = ( ULONG) rand() % 1000;
            break;
        case FLOW:
            sample.sensorValue = ( ULONG) rand() % 10;
            break;
        default:
```

```
            break;
    }
    /* send address */
    RK_ERR err = kMailboxPost(&sensorBox, &samplePtr, RK_NO_WAIT);
    if (err != RK_ERR_SUCCESS)
        sampleErr ++;

}

/* Producer - higher priority, blocks on mailbox */
VOID Task1(VOID *args)
{
    RK_UNUSEARGS
    Mesg_t *recvSample = NULL;
    while (1)
    {
        /* receive a pointer */
        RK_ERR errmbox = kMailboxPend(&sensorBox, &recvSample,RK_WAIT_FOREVER);
        K_ASSERT( errmbox==RK_ERR_SUCCESS);

        /* enqueue by copy into the stream */
        RK_ERR err = kMesgQueueSend(&sensorStream, recvSample, RK_NO_WAIT);

        K_ASSERT(err >= 0); /* either succesful or unsuccesful */
        if (err == RK_ERR_SUCCESS)
        {
            CHAR const *sensorTypeStr = NULL;
            if (recvSample->sensorType == 1)
                sensorTypeStr = "TEMP";
            if (recvSample->sensorType == 2)
                sensorTypeStr = "HUM";
            if (recvSample->sensorType == 3)
                sensorTypeStr = "CO2";
            if (recvSample->sensorType == 4)
                    sensorTypeStr = "FLOW";
            printf( "ENQ: [@%lums, %s, %lu] \r\n", kTickGet(), sensorTypeStr,
                        recvSample->sensorValue);
        }
        /* full, drop this sample and signal task2 */
        else if (err == RK_ERR_MESGQ_FULL)
        {
            kTaskEventSet(task2Handle, BATCH_READY);
        }
    }
}

/* for each sensor:
     . a ring buffer of AVG_WINDOW_SIZE values
 . sum of values
 . an index table (=enum - 1 eg., HUMIDITY IDX=2-1=1)
 */
```

```
static ULONG ringBuf[N_SENSOR][AVG_WINDOW_SIZE];
static ULONG ringSum[N_SENSOR] = {0};
static UINT ringIndex[N_SENSOR] = {0};

void Task2( void *args)
{

    RK_UNUSEARGS
    Mesg_t readSample;
    while (1)
    {
        RK_ERR err = kMesgQueueRecv(&sensorStream, &readSample, RK_NO_WAIT);
        if (err == RK_ERR_SUCCESS)
        {
            UINT sensorIdx = readSample.sensorType - 1;

/* remove oldest sample */
            ULONG oldest = ringBuf[sensorIdx][ringIndex[sensorIdx]];
            ringSum[sensorIdx] -= oldest;

/* push new sample */
            ringBuf[sensorIdx][ringIndex[sensorIdx]] = readSample.sensorValue;
            ringSum[sensorIdx] += readSample.sensorValue;

/* index incr-wrap */
            ringIndex[sensorIdx] ++;
            ringIndex[sensorIdx] %= AVG_WINDOW_SIZE;

/* simple average */
            ULONG avg = ringSum[sensorIdx] / AVG_WINDOW_SIZE;


            CHAR const *sensorTypeStr = NULL;
            if (readSample.sensorType == 1)
                sensorTypeStr = "TEMP";
            if (readSample.sensorType == 2)
                sensorTypeStr = "HUM";
            if (readSample.sensorType == 3)
                sensorTypeStr = "CO2";
            if (readSample.sensorType == 4)
                sensorTypeStr = "FLOW";

            printf( "DEQ: [@%lums, %s, %lu] | AVG: %lu \r\n", kTickGet(),
                    sensorTypeStr, readSample.sensorValue, avg);

        }
        else
        {
            ULONG gotFlags = 0UL;
            kTaskEventGet(BATCH_READY, RK_EVENT_ANY, &gotFlags, RK_WAIT_FOREVER);
        }
```

```
    }
}
```

```
OUTPUT:

ENQ: [@550ms, CO2, 571]
ENQ: [@560ms, FLOW, 4]
ENQ: [@570ms, FLOW, 4]
ENQ: [@580ms, HUM, 25]
ENQ: [@590ms, CO2, 931]
ENQ: [@600ms, CO2, 487]
ENQ: [@610ms, FLOW, 7]
ENQ: [@620ms, HUM, 79]

>>> Queue is full. Now offload and process. Note the order remains <<<

DEQ: [@630ms, CO2, 571] | AVG: 460
DEQ: [@631ms, FLOW, 4] | AVG: 5
DEQ: [@632ms, FLOW, 4] | AVG: 5
DEQ: [@633ms, HUM, 25] | AVG: 52
DEQ: [@634ms, CO2, 931] | AVG: 553
DEQ: [@635ms, CO2, 487] | AVG: 549
DEQ: [@636ms, FLOW, 7] | AVG: 5
DEQ: [@637ms, HUM, 79] | AVG: 55

>>> Consumer is preempted <<<
ENQ: [@640ms, CO2, 913]
ENQ: [@650ms, CO2, 134]
ENQ: [@660ms, HUM, 47]
ENQ: [@670ms, HUM, 30]
ENQ: [@680ms, TEMP, 7]
ENQ: [@690ms, CO2, 726]
ENQ: [@700ms, FLOW, 7]
ENQ: [@710ms, TEMP, 43]

DEQ: [@720ms, CO2, 913] | AVG: 578
DEQ: [@721ms, CO2, 134] | AVG: 543
DEQ: [@722ms, HUM, 47] | AVG: 51
DEQ: [@723ms, HUM, 30] | AVG: 44
DEQ: [@724ms, TEMP, 7] | AVG: 20
DEQ: [@725ms, CO2, 726] | AVG: 592
DEQ: [@726ms, FLOW, 7] | AVG: 5
DEQ: [@727ms, TEMP, 43] | AVG: 23
```

7.1.3.7.3. Queue Select using Notify Callback

A task is receiving from many queues and need to know which one has been able to complete.

The `notifyCbk(queue*)` is executed every time a send is successful. In this case it is using an Event Signal to a task. The Signals Flag indicate the queue number - as a contract - of which queue has completed sends. Note that as sends may coalesce, while a flag caps at 1, the consumer will drain each queue until it is empty, or it is preempted.

```
/* Many-to-1 queue channels */

/* Consumer Select queue based on its event
flags, that a succesfull send triggers  */

#define LOG_PRIORITY 4 /* keep logger as lowest-priority user task */
#define STACKSIZE 256

#define NQUEUES 3 /* number of queues */
#define QSIZ 8 /* depth of each queue */

#define Q0_FLAG   RK_EVENT_1 // (1<<0)
#define Q1_FLAG   RK_EVENT_2 // (1<<1)
#define Q2_FLAG   RK_EVENT_3 // (1<<2)
#define QFLAGS   (ULONG)(Q0_FLAG | Q1_FLAG | Q2_FLAG)

typedef struct
{
    RK_TASK_HANDLE producer;
    UINT payload;
} MESG_t;

/* Succesful Send callbacks */
/* each callback follows this pattern */
static inline
VOID sendNotify0(RK_MESG_QUEUE *qPtr)
{
    (VOID)qPtr;
    kTaskEventSet(consumerHandle, Q0_FLAG);
    /* Q1 flag for queue1 and so forth */
}

/* each callback in installed using kMesgQueueInstallSendCbk
on kApplicationInit() */

/* helper to send */
static inline
VOID enqueueSample(RK_MESG_QUEUE *qPtr UINT payload)
{
    MESG_t mesg = {
        .payload = payload,
        .producer = RK_RUNNING_HANDLE,
    };
    RK_ERR err = kMesgQueueSend(qPtr, &mesg, RK_WAIT_FOREVER);
    K_ASSERT(err == RK_ERR_SUCCESS);
```

```
}


VOID Prod0Task(VOID *args)
{
    RK_UNUSEARGS
    UINT seq = 0U;

    while (1)
    {
        enqueueSample(&queues[0], seq++);
        kSleepRelease(25); /* fast producer */
    }
}

VOID Prod1Task(VOID *args)
{
    UINT seq = 0U;
    RK_UNUSEARGS

    while (1)
    {
        enqueueSample(&queues[1], seq++);

        /* every fourth sample, also tickle the third queue */
        if ((seq & 0x3U) == 0U)
        {
            enqueueSample(&queues[2], seq);
        }

        kSleepRelease(60);
    }
}

/* Consumer listens on all queues, selecting those on its signal flags. */
VOID ConsumerTask(VOID *args)
{
    RK_UNUSEARGS

    MESG_t recv = {0};
    ULONG flags = 0UL;

    while (1)
    {
        flags = 0UL;
        kTaskEventGet(QFLAGS, RK_EVENT_FLAGS_ANY, &flags,
                            RK_WAIT_FOREVER);

        for (UINT i = 0; i < NQUEUES; ++i)
        {
            if (flags & (1UL << i))
```

```
            {
                while (kMesgQueueRecv(&queues[i], (VOID*)&recv, RK_NO_WAIT) ==
                       RK_ERR_SUCCESS)
                {
                    logPost("Q%u <- sender=%s payload=%u", i,
                            RK_TASK_NAME(recv.producer), recv.payload);
                }
            }
        }
    }
}
```

```
       0 ms :: Q1 <- sender=Prod1 payload=0
     250 ms :: Q0 <- sender=Prod0 payload=1
     500 ms :: Q0 <- sender=Prod0 payload=2
     600 ms :: Q1 <- sender=Prod1 payload=1
     750 ms :: Q0 <- sender=Prod0 payload=3
    1000 ms :: Q0 <- sender=Prod0 payload=4
    1200 ms :: Q1 <- sender=Prod1 payload=2
    1250 ms :: Q0 <- sender=Prod0 payload=5
    1500 ms :: Q0 <- sender=Prod0 payload=6
    1750 ms :: Q0 <- sender=Prod0 payload=7
    1800 ms :: Q1 <- sender=Prod1 payload=3
    1800 ms :: Q2 <- sender=Prod1 payload=4
    2000 ms :: Q0 <- sender=Prod0 payload=8
    2250 ms :: Q0 <- sender=Prod0 payload=9
    2400 ms :: Q1 <- sender=Prod1 payload=4
    2500 ms :: Q0 <- sender=Prod0 payload=10
    2750 ms :: Q0 <- sender=Prod0 payload=11
    3000 ms :: Q0 <- sender=Prod0 payload=12
    3000 ms :: Q1 <- sender=Prod1 payload=5
    3250 ms :: Q0 <- sender=Prod0 payload=13
    3500 ms :: Q0 <- sender=Prod0 payload=14
    3600 ms :: Q1 <- sender=Prod1 payload=6
    3750 ms :: Q0 <- sender=Prod0 payload=15
    4000 ms :: Q0 <- sender=Prod0 payload=16
    4200 ms :: Q1 <- sender=Prod1 payload=7
    4200 ms :: Q2 <- sender=Prod1 payload=8
    4250 ms :: Q0 <- sender=Prod0 payload=17
    4500 ms :: Q0 <- sender=Prod0 payload=18
    4750 ms :: Q0 <- sender=Prod0 payload=19
    4800 ms :: Q1 <- sender=Prod1 payload=8
    5000 ms :: Q0 <- sender=Prod0 payload=20
    5250 ms :: Q0 <- sender=Prod0 payload=21
```

7.1.3.7.4. Synchronous Client-Server Procedure Call

The example below computes a CRC on the server for the client's payload and returns it as the reply code. Logically, this is unbuffered: the client blocks until it receives the reply.

For kPortSendRecv(), each client task must pass the address of a Mailbox to store the client response (4-byte).

Note that the server may temporarily adopt the client priority while serving the request. After kPortReplyDone(), it returns to nominal priority.

```
#include <application.h>
#include <logger.h>
#define STACKSIZE 256
#define PORT_MSG_WORDS 4U   /* 2 words meta + 2 words payload */
#define PORT_CAPACITY 16

/* tasks */
RK_DECLARE_TASK(serverHandle, ServerTask,    stack1, STACKSIZE)
RK_DECLARE_TASK(clientHandle, ClientTask,    stack2, STACKSIZE)

/* port */
static RK_PORT serverPort;
RK_DECLARE_PORT_BUF(portBuf, PORT_MSG_WORDS, PORT_CAPACITY)

/* 4-word message format; first two words are metadata */
typedef RK_PORT_MESG_4WORD RpcMsg;

static inline UINT crc32(const VOID *data, ULONG size);
static inline BYTE xorshift8(void);

VOID kApplicationInit(void)
{
    K_ASSERT(!kCreateTask(&serverHandle, ServerTask, RK_NO_ARGS, "Server", stack1,
STACKSIZE, 1, RK_PREEMPT));
    K_ASSERT(!kCreateTask(&clientHandle, ClientTask, RK_NO_ARGS, "Client", stack2,
STACKSIZE, 2, RK_PREEMPT));

    /* init port */
    K_ASSERT(!kPortInit(&serverPort, portBuf, PORT_MSG_WORDS, PORT_CAPACITY,
serverHandle));

    logInit();
}

VOID ServerTask(VOID *args)
{
    RK_UNUSEARGS
    RpcMsg msg;
    while(1)
    {
        /* receive next request; server may adopt client priority here */
        K_ASSERT(!kPortRecv(&serverPort, &msg, RK_WAIT_FOREVER));

        BYTE  *vector = (BYTE*) msg.payload[0];
```

```
        ULONG   size  =          msg.payload[1];
        UINT    crc   = crc32(vector, size);

        logPost("[SERVER] Will Reply CRC=0x%04X | Eff Prio=%d | Nom Prio=%d",
                crc, runPtr->priority, runPtr->prioNominal);

        /* must end with kPortReplyDone */
        K_ASSERT(!kPortReplyDone(&serverPort, (ULONG const*)&msg, crc));

        logPost("[SERVER] Finished. | Eff Prio: %d | Nom Prio: %d", runPtr->priority,
runPtr->prioNominal);
    }
}

static RK_MAILBOX replyBox;

VOID ClientTask(VOID *args)
{
    RK_UNUSEARGS
    static BYTE vec[8];
    for (UINT i = 0; i < 8; ++i)
        vec[i] = xorshift8();

    RpcMsg msg = {0};
    msg.payload[0] = (ULONG) vec;  /* pointer as one word */
    msg.payload[1] = 8;            /* number of bytes */

    UINT reply = 0;
    while(1)
    {
        /* Send-Receive: a call */
        UINT want = crc32(vec, 8);
        K_ASSERT(!kPortSendRecv(&serverPort, (ULONG*)&msg, &replyBox  &reply,
RK_WAIT_FOREVER));
        logPost("[CLIENT] Need=0x%04X | Recvd=0x%04X", want, reply);
        /* if reply is correct, generate a new payload */
        if (want == reply)
            for (UINT i = 0; i < 8; ++i) vec[i] = xorshift8();
        kSleepRelease(1000);
    }
}
```

```
   28000 ms :: [CLIENT] Need=0xC2A6C337 | Recvd=0xC2A6C337
   29000 ms :: [SERVER] Will Reply CRC=0x93F4110A | Eff Prio=2 | Nom Prio=1
   29000 ms :: [SERVER] Finished. | Eff Prio: 1 | Nom Prio: 1
   29000 ms :: [CLIENT] Need=0x93F4110A | Recvd=0x93F4110A
   30000 ms :: [SERVER] Will Reply CRC=0x7A8FA006 | Eff Prio=2 | Nom Prio=1
   30000 ms :: [SERVER] Finished. | Eff Prio: 1 | Nom Prio: 1
   30000 ms :: [CLIENT] Need=0x7A8FA006 | Recvd=0x7A8FA006
   31000 ms :: [SERVER] Will Reply CRC=0x9051C8B1 | Eff Prio=2 | Nom Prio=1
```

```
31000 ms :: [SERVER] Finished. | Eff Prio: 1 | Nom Prio: 1
31000 ms :: [CLIENT] Need=0x9051C8B1 | Recvd=0x9051C8B1
32000 ms :: [SERVER] Will Reply CRC=0x11F29117 | Eff Prio=2 | Nom Prio=1
32000 ms :: [SERVER] Finished. | Eff Prio: 1 | Nom Prio: 1
```

7.1.4. Most-Recent Message Protocol (MRM)

MRM Control Block
MRM Buffer Allocator
Data Buffer Allocator
Current MRM Buffer Address
Data Size (Message Size)

MRM Buffer
Data Buffer Address
Readers Count

Data Buffer
Application-dependent

> There is little practical difference between a message that does not arrive and one with no valid (stale) data. But when wrong (or stale) data is processed - e.g., to define a set point on a loop - a system can fail badly.
>
> Design Choice: provide a broadcast asynchronous message-passing scheme that guarantees data freshness and integrity for all readers.
>
> Benefits: The system has a mechanism to meet strict deadlines that cannot be predicted on design time.

Control loops reacting to unpredictable time events—like a robot scanning an environment or a drive-by-wire system—require a different message-passing approach. Readers cannot "look at the past" and cannot block. The most recent data must be delivered lock-free and have guaranteed integrity.

As *Ports*, the *MRM* is a high-level mechanism. It was chosen to be provided as a kernel service, given its distinctive nature and suitability for *RK0* applications.

7.1.4.1. Functional Description

An *MRM* works as a *1-to-many asynchronous Mailbox* - with a lock-free specialisation that enables several readers to get the most recent deposited message with no integrity issues. Whenever a reader reads an MRM buffer, it will find the most recent data transmitted. It can also be seen as an extension of the Double Buffer pattern for a 1:N communication.

The core idea of the MRM protocol is that readers can only access the buffer that is classified as the '*most recent buffer*'. After a writer *publish()* a message, that will be the only message readers can *get()* — any former message being processed by a reader was grabbed *before* a new *publish()* - and, from now on, can only be *unget()*, eventually returning to the pool.

To clarify further, the communication steps are listed:

1. A producer first reserves an MRM Buffer - the reserved MRM Buffer is not available for reading until it is published.
2. A message buffer is allocated and filled, and its address is within an MRM Buffer. The producer *publishes* the message. From now on, it is *the most recent message*. Any former published buffer is no longer visible to new readers
3. A reader starts by *getting* an MRM Buffer. A `get()` operation delivers a copy of the message to the reader's scope. Importantly, this operation increases the number of readers associated to that MRM Buffer.

Before ending its cycle, the task releases (`unget()`) the buffer; on releasing, the kernel checks if the caller task is the last reader and if the buffer being released is not the current MRM Buffer.

If the above conditions are met, the `unget()` operation will return the MRM buffer to the pool. If there are more readers, OR if it is the current buffer, it remains available.

When the `reserve` operation detects that the most recent buffer still has readers, a new buffer is allocated to be written and published. If it has no readers, it is reused.

This way, the worst case is a sequence of `publish()` with no `unget()` at all — this would lead to the writer finding no buffer to reserve. This is prevented by making: `N Buffers = N tasks + 1`.

7.1.4.2. MRM Control Block Configuration

What might lead to some confusion when initialising an MRM Control Block is the need for two different pools:

- One pool will be the storage for the MRM Buffers, which is the data structure for the mechanism.
- Another pool is for the actual payload. The messages.

Both pools must have the same number of elements: the number of tasks communicating + 1.

- The size of the data buffers is application-dependent - and is passed as a number of *words*. The minimal message size is 32-bit.
- If using data structures, keep it aligned to 4 to take advantage of the performance of aligned memory.

7.1.4.3. Usage Example

Consider a modern car - speed variations are of interest in many modules. With a somehow "naive" approach, let us consider three modules and how they should react when speed varies:

1. **Cruiser Control:** For the Cruiser Control, a speed increase might signify the driver wants manual control back, and it will likely turn off.
2. **Windshield Wipers:** If they are on, a speed change can affect the electric motor's adjustments to the air resistance.
3. **Radio:** Speed changes reflect the aerodynamic noise - the radio volume might need adjustment.

As the variations are unpredictable, we need a mechanism to deliver the last speed in order of importance for all these modules. From highest to lowest priority, Cruise, Wipers, and Radio are the three modules that range from safety to comfort.

To emulate this scenario, we can write an application with a higher priority task that sleeps and wakes up at pseudo-random times to produce random values that represent the (unpredictable) speed changes.

The snippet below has 4 periodic tasks. Tasks are periodic using the `kSleepRelease()` primitive.

There is a sequential counter. When a task wakes and see it was not incremented it just runs.

The producer publishes new data at a random interval, preempting whatever task is running at the moment. Given tha random updates,

```
typedef struct
{
    UINT speed;
    ULONG timeStamp;
} Mesg_t;

#define STACKSIZE 256
#define N_MRM (5)                          /* Number of MRMs N Tasks + 1 */
#define MRM_MESG_SIZE (sizeof(Mesg_t) / 4) /* In WORDS */
RK_MRM MRMCtl;                             /* MRM control block */
RK_MRM_BUF buf[N_MRM];                     /* MRM pool */
Mesg_t data[N_MRM];                        /* message data pool */

RK_DECLARE_TASK(speedSensorHandle, SpeedSensorTask, stack1, STACKSIZE)
RK_DECLARE_TASK(cruiserHandle, CruiserTask, stack2, STACKSIZE)
RK_DECLARE_TASK(wiperHandle, WiperTask, stack3, STACKSIZE)
RK_DECLARE_TASK(radioHandle, RadioTask, stack4, STACKSIZE)
volatile UINT seq = 0;
VOID kApplicationInit(VOID)
{

    kCreateTask(&speedSensorHandle, SpeedSensorTask, RK_NO_ARGS, "SpeedTsk",
                stack1, STACKSIZE, 1, RK_PREEMPT);

    kCreateTask(&cruiserHandle, CruiserTask, RK_NO_ARGS, "CruiserTsk", stack2,
                STACKSIZE, 2, RK_PREEMPT);

    kCreateTask(&wiperHandle, WiperTask, RK_NO_ARGS, "WiperTsk", stack3,
```

```
                STACKSIZE, 3, RK_PREEMPT);

    kCreateTask(&radioHandle, RadioTask, RK_NO_ARGS, "RadioTsk", stack4,
                STACKSIZE, 4, RK_PREEMPT);

    kMRMInit(&MRMCtl, buf, data, N_MRM, MRM_MESG_SIZE);

    logInit(5);
}

VOID SpeedSensorTask(VOID *args)
{
    RK_UNUSEARGS

    Mesg_t sendMesg = {0};
    while (1)
    {
        RK_TICK sleepTicks = ((RK_TICK)rand() % 18) + 1;
        kSleep(sleepTicks);
        RK_TICK currTick = kTickGetMs();
        UINT speedValue = (UINT)(rand() % 170) + 1;
        sendMesg.speed = speedValue;
        sendMesg.timeStamp = currTick;
        /* grab a buffer */
        RK_MRM_BUF *bufPtr = kMRMReserve(&MRMCtl);
        if (bufPtr != NULL)
        {
            K_ASSERT(!kMRMPublish(&MRMCtl, bufPtr, &sendMesg));
            printf("!!!!! @%lums SPEED UPDATE: %u mph\r\n", kTickGetMs(),
speedValue);
            seq += 1;
        }
        else
        { /* cannot fail */
            logError("MRM protocol could not find a free buffer\r\n");
        }
        /* publish  */
    }
}

VOID CruiserTask(VOID *args)
{
    RK_UNUSEARGS
    Mesg_t recvMesg = {0};
    while (1)
    {
        RK_MRM_BUF *readBufPtr = kMRMGet(&MRMCtl, &recvMesg);
        if (readBufPtr)
        {
            logPost("CRUISER: (%u mph, %lu ms)", recvMesg.speed,
recvMesg.timeStamp);
```

```
            kMRMUnget(&MRMCtl, readBufPtr);
        }
        kSleepRelease(4);
    }
}

VOID WiperTask(VOID *args)
{
    RK_UNUSEARGS
    Mesg_t recvMesg = {0};

    while (1)
    {

        RK_MRM_BUF *readBufPtr = kMRMGet(&MRMCtl, &recvMesg);
        if (readBufPtr)
        {
            logPost("WIPERS: (%u mph, %lu ms)", recvMesg.speed,
recvMesg.timeStamp);

            kMRMUnget(&MRMCtl, readBufPtr);
        }
        kSleepRelease(7);
    }
}
VOID RadioTask(VOID *args)
{
    RK_UNUSEARGS
    Mesg_t recvMesg = {0};
    while (1)

    {

        RK_MRM_BUF *readBufPtr = kMRMGet(&MRMCtl, &recvMesg);

        if (readBufPtr)
        {
            logPost("RADIO: (%u mph, %lu ms) ", recvMesg.speed,
recvMesg.timeStamp);
            kMRMUnget(&MRMCtl, readBufPtr);
        }
        kSleepRelease(11);
    }
}
```

Thus, different situations can happen:

- All tasks read the updated pair (speed, time)

- Not all tasks receive the updated pair because another update happens in between.
- No tasks receive an update - because another happens too soon.
- No update happens between in the period of a given task. It receives the same value. No problems.

All these cases are on the log:

```
Logs show: (last speed record, record time)

  !!!@120ms SPEED UPDATE: 164 mph
     120 ms :: CRUISER: (164 mph, 120 ms)
     140 ms :: WIPERS: (164 mph, 120 ms)
 !!! @150ms SPEED UPDATE: 80 mph
     160 ms :: CRUISER: (80 mph, 150 ms)
     200 ms :: CRUISER: (80 mph, 150 ms)
     210 ms :: WIPERS: (80 mph, 150 ms)
     220 ms :: RADIO: (80 mph, 150 ms)
     240 ms :: CRUISER: (80 mph, 150 ms)
  !!!@280ms SPEED UPDATE: 49 mph
     280 ms :: CRUISER: (49 mph, 280 ms)
     280 ms :: WIPERS: (49 mph, 280 ms)
     320 ms :: CRUISER: (49 mph, 280 ms)
     330 ms :: RADIO: (49 mph, 280 ms)
     350 ms :: WIPERS: (49 mph, 280 ms)
     360 ms :: CRUISER: (49 mph, 280 ms)
     400 ms :: CRUISER: (49 mph, 280 ms)
     420 ms :: WIPERS: (49 mph, 280 ms)
     440 ms :: CRUISER: (49 mph, 280 ms)
     440 ms :: RADIO: (49 mph, 280 ms)
  !!!@450ms SPEED UPDATE: 87 mph
     480 ms :: CRUISER: (87 mph, 450 ms)
     490 ms :: WIPERS: (87 mph, 450 ms)
     520 ms :: CRUISER: (87 mph, 450 ms)
 !!! @540ms SPEED UPDATE: 110 mph
     550 ms :: RADIO: (110 mph, 540 ms)
     560 ms :: CRUISER: (110 mph, 540 ms)
     560 ms :: WIPERS: (110 mph, 540 ms)
     600 ms :: CRUISER: (110 mph, 540 ms)
     630 ms :: WIPERS: (110 mph, 540 ms)
 !!! @640ms SPEED UPDATE: 22 mph
     640 ms :: CRUISER: (22 mph, 640 ms)
     660 ms :: RADIO: (22 mph, 640 ms)
```

The highlight is that controllers can keep their pace, while receiving fresh data - you can see it on the timestamp on the image.

Again, they might receive the same data more than once or miss samples; what is important is that *they are not lagging and consuming stale data.*

Chapter 8. *Error Handling*

8.1. Fail fast

While tracing and error handling are yet to be largely improved (and that is when the 1.0.0 version will be released), currently *RK0* employs a policy of *failing fast* in **debug mode**.

When Error Checking is enabled, every kernel call will be 'defensive', checking for correctness of parameters and invariants, null dereferences, etc.

In these cases is more useful to allow the first error to halt the execution by calling an Error Handler function to observe the program state.

A trace structure records the address of the running TCB, its current stack pointer, the link register (that is, the PC at kErrHandler was called), and a time stamp.

This record is on a `.noinit` RAM section, so it is visible if CPU resets. A fault code is stored in a global `faultID` and on the trace structure. Developers can hook in custom behaviour.

If the kernel is configured to not halt on a fault, but Error Checking is enabled, functions will return negative values in case of an error.

On the other hand, when Error Checking is disabled or `NDEBUG` is defined nothing is checked, reducing code size and improving performance.

(*Some deeper internal calls have assertion. For those, only* `NDEBUG` *defined ensures they are disabled.*)

8.2. Stack Overflow

Stack overflow is detected (not prevented) using a "stack painting" with a sentinel word. Stack Overflow detection is enabled by defining the assembler preprocessor `__KDEF_STACKOVFLW` when compiling.

As a matter of fact, sizing your stack is something you must do diligently when programming a system. I would say a mechanism for stack overflow detection is on the bottom of the list of 'must-have' features.

One can take advantage of the static task model - *it is possible to predict offline* the deepest call within any task. The compiler flag `-fstack-usage` creates `.su` files indicating the depth of every function within a module. This is an example of compilation-unit output:

```
core/src/ksema.c:34:8:kSemaphoreInit    88  static
core/src/ksema.c:74:8:kSemaphorePend    96  static
core/src/ksema.c:189:8:kSemaphorePost   88  static
core/src/ksema.c:265:8:kSemaphoreFlush  72  static
core/src/ksema.c:306:5:kSemaphoreQuery  40  static
core/src/kmutex.c:128:8:kMutexInit  16  static
core/src/kmutex.c:165:8:kMutexLock  120 static
```

```
core/src/kmutex.c:325:8:kMutexUnlock    96  static
core/src/kmutex.c:425:6:kMutexQuery 56  static
```

These are the worst cases. Now, you identify the depth of the longest *chain of calls* for a task using these services and add a generous safety margin — 30%. The cap depends on your budget.

Importantly, you also have to size the System Stack. This initial size is defined in `linker.ld` by the symbol `Min_Stack_Size`. In this case, account for the depth of `main()`, `kApplicationInit()`, and all interrupt handlers; again, inspect the longest call chain depth. Assume interrupts always add to the worst static depth, and account for *nested interrupts*.

8.3. Deadlocks

There are deadlock-recovery techniques in the literature, but they are generally unfeasible here. The kernel provides bounded waiting, enforces priority-ordered waiting queues, applies mutex priority inheritance, and offers lock-free primitives and period-drift compensation. None of these techniques prevents deadlocks by itself (and with bounded blocking plus lock-free primitives, one can still get *livelocks*).

- Ordered Locking:

For those programming the application, despite following the RMS rule of higher priority for higher request rate tasks, the golden rule for locking is acquiring resources in an unidirectional order *throughout the entire application*:

```
acquire(A);
acquire(B);
acquire(C);
   .
   .
   .
release(C);
release(B);
release(A);
```

This breaks circular waiting.

For instance:

```
TaskA:
   wait(R1);
   wait(R2);
    /* critical section */
   signal(R2);
   signal(R1);

TaskB:
   wait(R1);
```

```
    wait(R2);
     /* critical section */
    signal(R2);
    signal(R1);
```

But, if:

```
TaskA:
    wait(R1);
    wait(R2);
    .
    .

TaskB:
    wait(R2);
    wait(R1);
    .
    .
```

There are some possible outcomes:

1. Deadlock:
    - TaskA runs: acquires R1
    - TaskB runs: acquires R2
    - TaskA runs: tries to acquire R2 — blocked
    - TaskB runs: tries to acquire R1 — blocked
2. No deadlock:
    - TaskA runs: acquires R1
    - TaskA runs: acquires R2 (nobody is holding R2)
    - TaskA releases both; TaskB runs and acquires both (in either order)

Overall, there is no deadlock if tasks do not overlap in critical sections. That is why systems run for years without deadlocks and eventually: *ploft*.

- Use a 'master-lock' with 'try' semantics

Another technique that can be employed is if one needs to acquire multiple locks—acquire them all or none using a try-lock (`RK_NO_WAIT`). If any of the tries fail, the task gives up on acquiring the resources and backs off, releasing all successful locks to retry later (most simply, using a sleep queue). That is easier said than done, though, and, as mentioned, if not well done, instead of deadlocks, one gets livelocks.

(*Livelocks* are when a couple of tasks keep running, but the program does not advance.)

```
/pseudo code acquire M1 and M2 or nothing
```

```
    while (attempts-- > 0u)
    {
        if (lock(&m1, RK_NO_WAIT) != OK)
        {
            /* suspend to try again */
            sleep(interval);
            continue;
        }

        if (lock(&m2, RK_NO_WAIT) == OK)
        {
            /*both acquired*/

            return (OK);
        }
        /* failed m2 */

        /* back off m1 */
        unlock(m1);

        /* suspend to try again */
        sleep(interval);
    }
    return (ERROR_TIMEOUT);
```

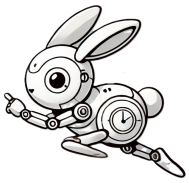

© 2026 Antonio Giacomelli | All Rights Reserved | www.kernel0.org