

RK0

The Real-Time Kernel '0'

System: v0.9.14-dev | Docbook: v0914.1.0.0 | www.kernel0.org

Table of Contents

1. THE KERNEL AT A GLANCE	1
1.1. The design approach	1
1.1.1. Architecture	1
1.1.2. Programming with RK0	1
1.1.3. Suitable Applications	2
1.2. Kernel Services	2
2. Task Scheduler	3
2.1. Scheduler Data Structures	3
2.1.1. Task Control Block	3
2.1.2. Task Queues	5
2.1.3. Ready Queue Table	5
2.1.4. Waiting Queues	5
2.1.5. The scheduling algorithm	6
2.2. Handling the scheduler	8
3. Timers and Delays	10
3.1. Busy-wait delay	10
3.2. Sleep Timers	11
3.2.1. Sleep Delay	11
3.2.2. Compensated Sleep Delays	12
3.2.2.1. Periodic Release Sleep	12
3.2.2.2. Sleep Until (local anchor base)	14
3.3. Blocking Time-out	16
3.4. Callout Timers (Application Timers)	16
3.5. System Tick	16
3.6. System Tasks	17
4. Memory Allocator	18
4.1. How it works	18
5. <i>Inter-Task Communication</i>	20
5.1. Task Events (Task Event Flags)	20
5.1.1. Usage Example: Supervisor Task	20
5.2. Semaphore	22
5.2.1. Counting Semaphore and Binary Semaphores	22
5.2.2. Semaphores in RK0	22
5.2.3. Flushing or Broadcast	23
5.2.4. Mutex Semaphore (Locks)	25
5.2.4.1. Priority Inversion and PIP	26
5.2.5. Mutexes vs Binary Semaphores	30
5.3. Scheduler Lock	31

5.4. Sleep Queue	31
5.5. Conditional Critical Regions.....	32
5.5.1. Monitor Invariants	32
5.5.2. Signalling Discipline.....	32
5.5.3. Condition Variable Model for (Mesa) Monitors	33
5.5.3.1. Usage Example: Synchronisation Barrier	34
5.5.3.2. Usage Example: Readers Writers Lock	37
5.6. Message Queue	39
5.6.1. Size of a Message.....	40
5.6.2. Blocking and non-blocking behaviour.....	40
5.6.3. Ownership and Priority Inversion	40
5.6.4. Notify callback.....	41
5.6.5. Usage Examples.....	41
5.6.5.1. Averaging Sensor Values	41
5.6.5.2. Queue Select using Notify Callback.....	46
5.6.5.3. Synchronous Client-Server Procedure Call.....	49
5.7. Most-Recent Message Protocol (MRM).....	51
5.7.1. Functional Description	52
5.7.2. MRM Control Block Configuration	52
5.7.3. Usage Example	53
6. Error Handling.....	58
6.1. Fail fast	58
6.2. Stack Overflow	58
6.3. Deadlocks	59
7. RK0 Services API	61
7.1. Convention	61
7.2. Return Values	61

Chapter 1. THE KERNEL AT A GLANCE

1.1. The design approach

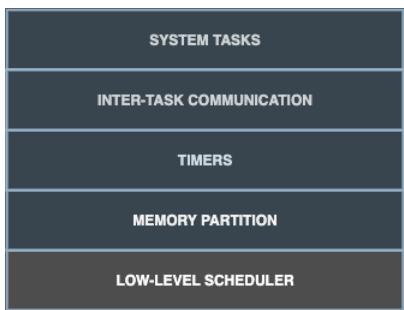
RK0 Blog: [About Processes, Tasks and Threads](#)

1.1.1. Architecture

The layered architecture can be split—roughly—into two: a top and a bottom layer. On the top, the *Executive* manages the resources needed by the application.

On the bottom, the *Low-level Scheduler* works as a software extension of the CPU.

Together, they implement the *Task* abstraction. This primitive is the *Concurrency Unit* and follows the *Thread* model. A *Task* is a *thread*.



In systems design jargon, the Executive enforces policy (what should happen). The Low-level Scheduler provides the mechanism (how it gets done). The services are the primitives that gradually translate policy decisions into concrete actions executed by the Scheduler.

RK0's goal is determinism on low-end devices. Its multitasking engine does not split user space from kernel space. Tasks execute in privileged mode and use a dedicated process stack pointer, distinct from the system stack. The rationale:

- Application tasks are not unknown entities at run time.
- Implementing system calls as traps either forces a non-preemptive kernel, or increases complexity in critical paths, degrading determinism.
- Relying on the ARMv6/7-M MPU decreases memory usage efficiency and introduces latency on control paths. It does not fit RK0's deterministic execution model.

1.1.2. Programming with RK0

RK0 is designed so that it does not get in the programmer's way.

It aims to be transparent, composable, deterministic, and with clear semantics.

Transparent means that the kernel does not hide its primitives behind opaque pointers nor introduce unexpected side effects.

Its components are composable because each feature is self-contained yet designed to work with others.

Clear semantics means mechanisms behave as expected across different usage patterns.

Determinism is pursued by avoiding execution paths whose worst case cannot be bounded or reasoned about in advance.

When possible, operations are $O(1)$, rely on word-aligned memory, and use static allocation.

When dynamic allocation is unavoidable, its worst-case behaviour can still be characterised offline (unless the application itself prevents any meaningful offline bound).

1.1.3. Suitable Applications

Given the architecture, *RK0* targets applications with the following characteristics:

1. They are designed to handle particular devices in which real-time responsiveness is imperative.
2. Applications and middleware may be implemented alongside appropriate drivers.
3. Drivers may even include the application itself.
4. *Untested programs are not loaded*: After the software has been tested, it can be assumed reliable.

1.2. Kernel Services

RK0 has *Core Services* (always enabled) and optional services (enabled by configuration).

Core Services:

- Scheduler
- Partition Memory Allocator
- Timer Delays
- Task Event Flags

Optional Services

- Application Timer (Callouts)
- Semaphore
- Mutex
- Sleep Queue (Condition Variable)
- Message Queue (and its extensions)
- Most-Recent Message Protocol (MRM)

When compiled with only *Core Services*, one gets a functional Executive with less than ~3 KB ROM.

Chapter 2. Task Scheduler

RK0 employs a priority-based preemptive scheduler, aligned with a Rate-Monotonic Assignment. Tasks are typically assigned priorities according to their request rates - i.e., tasks with shorter periods are assigned to higher priorities. The highest priority is represented by the value '0'; the lowest is represented by the value '31'.

A notable scheduler characteristic is constant-time complexity ($O(1)$) with low latency. This was achieved by carefully composing the data structures and an efficient '*choose-next*' algorithm, detailed below.

Time-slice was deprecated on version 0.5.0.

2.1. Scheduler Data Structures

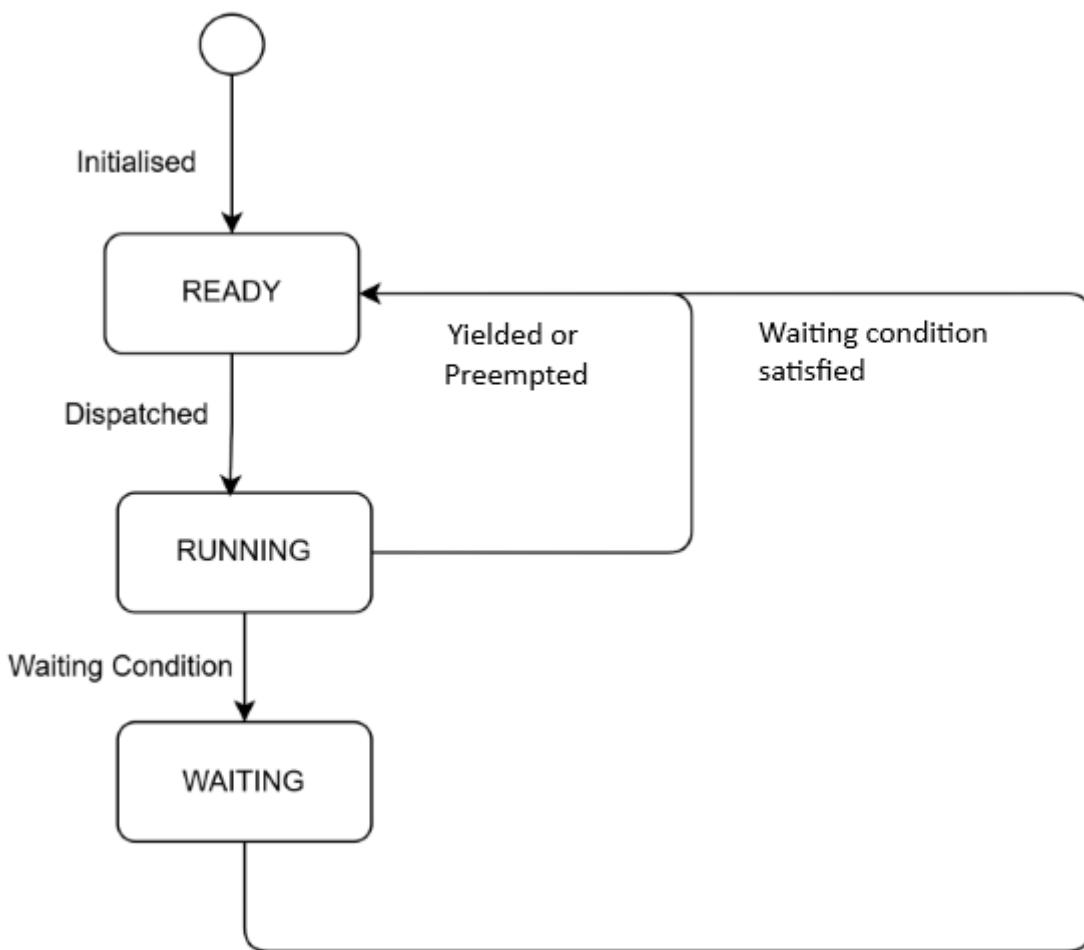
2.1.1. Task Control Block

Every primitive is associated to a data structure we refer to as its *Control Block*. A Task Control Block is a record for stack, resources, and time management. The table below partially represents a Task Control Block (as this document is live, this might not reflect the exact fields of the current version).

Task Control Block
Task name
Task ID
Status
Assigned Priority
Effective Priority
Saved Stack Pointer
Stack Address
Stack Size
Event Register Control Block
Last wake-time
Next wake-time
Time-out Flag
Preemption Flag
Owned Resources List
Waiting Resources List
Timeout List Node
TCB List Node

Tasks are static — they are not created (or destroyed) on runtime. There is no fork or join.

In practice, tasks are either *RUNNING* or 'waiting' for their turn to run.



We need to define *WAITING* and *READY* clearly:

1. A *READY* task will be dispatched; therefore, switch to *RUNNING* whenever it is the highest priority *READY* task.
2. A *WAITING* task depends on a condition, generalised as an *event* to switch to *READY*.
3. Logically, the *WAITING* state will assume different pseudo-states related to the kind of event that will switch a task to *READY*:
 - *SLEEPING*: a task is either sleeping (delayed) for a given time, or sleeping for a wake signal to be delivered to a *_Sleep Queue*.
 - *PENDING*: the task suspended itself, waiting for a combination of Event Flags on its Event Register.
 - *BLOCKED*: A task is blocked on a mutex or semaphore.
 - *SENDING/RECEIVING*: A producer task, when blocking on a Message Passing object, switches its status to *SENDING*, and a consumer to *RECEIVING*.

The scheduler rules, not the heap.

RK0 tasks are static.

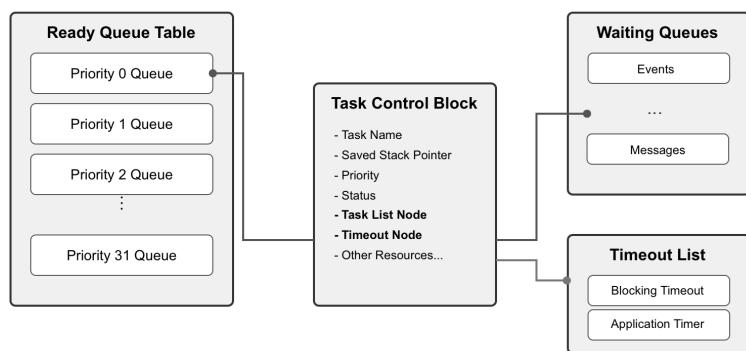
It's a design decision rooted in real-time correctness.

Besides an application-specific system software does not need to treat tasks as 'unknown' objects.

The wins:

- A memory layout the systems programmer knows.
- No alignment traps.
- Link-time visibility:
 - Each task's stack is a named symbol in the linker map.
 - You can inspect and verify the memory layout before flashing.
 - A simple `objdump` reveals all stack allocations — that's peace of mind.

RK0: Scheduler Data Structures



2.1.2. Task Queues

The backbone of the queues where tasks will wait for their turn to run is a circular doubly linked list: removing any item from a double list takes O(1) (provided we don't need to search the item). As the kernel knows each task's address, adding and removing is always O(1). Singly linked lists can't achieve O(1) for removal.

2.1.3. Ready Queue Table

Another design choice to achieve O(1) is the global ready queue, which is a table of FIFO queues—each queue dedicated to a priority—and not a single ordered queue. So, enqueueing a ready task is always O(1). Given the sorting needed, the time complexity would be O(n) if tasks were placed on a single ready queue.

2.1.4. Waiting Queues

The scheduler does not have a unique waiting queue. Every kernel object that can block a task has an associated waiting queue. Because these queues are a scheduler component, *they follow a priority discipline*: the highest priority task is dequeued first, *always*.

When an event capable of switching tasks from *WAITING* to *READY* happens, one or more tasks (depending on the mechanism) are then placed on the ready list, unique to their priority. Now, they are waiting to be picked by the scheduler—that is the definition of *READY*.

2.1.5. The scheduling algorithm

As the ready queue table is indexed by priority - the index 0 points to the queue of ready tasks with priority 0, and so forth, and there are 32 possible priorities - a 32-bit integer can represent the state of the ready queue table. It is a BITMAP:

The BITMAP computation: ((1a) OR (1b)) AND (2), s.t.:

- (1a) Every Time a task is readied, update: $\text{BITMAP} |= (1U << \text{task->priority})$;
 - (1b) Every Time an empty READY QUEUE becomes non-empty, update: $\text{BITMAP} |= (1U << \text{queueIndex})$
 - (2): Every Time READY QUEUE becomes empty, update: $\text{BITMAP} \&= \sim(1U << \text{queueIndex})$;
- EXAMPLE:

Ready Queue Index :	(6)5 4 3 2 1 0
Not empty :	1 1 1 0 0 1 0
----->	
(LOW) Effective Priority (HIGH)	

In this case, the scenario is a system with 7 priority task levels. Queues with priorities 6, 5, 4, and 1 are not empty.

Having the Ready Queue Table bitmap, we find the highest priority non-empty task list as follows:

(1) Isolate the **rightmost '1'**:

$\text{RBITMAP} = \text{BITMAP} \& -\text{BITMAP}$. (- is the bitwise operator for two's complement: $\sim\text{BITMAP} + 1$)

In this case:

[31]	[0]	: Bit Position
0...1110010		: BITMAP
1...0001110		: -BITMAP
=====		
0...0000010		: RBITMAP
		[1]

The rationale here is that, for a number N , its 2's complement $-N$, flips all bits - except the rightmost '1' (by adding '1'). Then, $N \& -N$ results in a word with all 0-bits except for the less significant '1'.

(2) Extract the **rightmost '1' position**:

- For ARMv7M, we benefit from the **CLZ** instruction to count the *leading zeroes*. As they are the

number of zeroes on the left of the rightmost bit, '1', this value is subtracted from 31 to find the Ready Queue index.

```
RK_FORCE_INLINE static inline
unsigned __getReadyPrio(unsigned readyQBitmap)
{
    unsigned ret;
    __ASM volatile (
        "clz    %0, %1      \n"
        "neg    %0, %0      \n"
        "add    %0, %0, #31\n"
        : "=r" (ret)
        : "r" (readyQBitmap)
        :
    );
    return (ret);
}
```

This instruction would return #30, and $\#31 - \#30 = \#01$ in the example above.

- For ARMv6M there is no suitable hardware instruction. The algorithm is written in C and counts the *trailing zeroes*, thus, the index number. Although it might vary depending on your compiler settings, it takes ~11 cycles (*note it is still O(1)*):

```
/*
 De Bruijn's multiply+LUT
 (Hacker's Delight book)
 */

/* table is on a ram section for efficiency */
RK_SECTION(getReadyTable)
const static unsigned readyPrioTbl[32] =
{
    0, 1, 28, 2, 29, 14, 24, 3, 30, 22, 20, 15, 25, 17, 4, 8,
    31, 27, 13, 23, 21, 19, 16, 7, 26, 12, 18, 6, 11, 5, 10, 9
};

RK_FORCE_INLINE static inline
unsigned __getReadyPrio(unsigned readyQBitmap)
{
    unsigned mult = readyQBitmap * 0x077CB531U;

    /* Shift right the top 5 bits
     */
    unsigned idx = (mult >> 27);

    /* LUT */
    unsigned ret = (unsigned)readyPrioTbl[idx];
    return (ret);
}
```

```
}
```

For the example above, `mult = 0x2 * 0x077CB531 = 0xEF96A62`. The 5 leftmost bits (the index) are `00001` → `table[1] = 1`.

During a context switch, the procedures to find the highest priority non-empty ready queue table index are as follows:

```
static inline RK_PRI0 kCalcNextTaskPrio_(VOID)
{
    if (readyQBitMask == 0U)
    {
        return (idleTaskPrio);
    }
    readyQRightMask = readyQBitMask & -readyQBitMask;
    RK_PRI0 prioVal = (RK_PRI0) (_getReadyPrio(readyQRightMask));
    return (prioVal);
}

VOID kSchSwtch(VOID)
{
    /* O(1) complexity */
    nextTaskPrio = kCalcNextTaskPrio_();

    RK_TCB* nextRunPtr = NULL;

    /* O(1) complexity */
    kTCBQDeq(&readyQueue[nextTaskPrio], &nextRunPtr);

    runPtr = nextRunPtr;
}
```

2.2. Handling the scheduler

RKO Blog: [About Real-Time, Responsiveness and Throughput](#)

An essential characteristic of the scheduler is that it is a *preemptive run-to-completion* scheduler. This term, '*run-to-completion*' has slightly different meanings depending on the context. It is often related to strictly cooperative schedulers, in the sense tasks must *yield* the processor. Otherwise, they monopolise the CPU.

In RKO, tasks with the same priority work cooperatively. This is different from schedulers that employ a *time-slice* (quantum) for round-robin: after the slice expires, the task is put at the *tail* of the *Ready Queue*.

The term *run-to-completion* here is to be interpreted as follows:

- The scheduler's behaviour is to choose the highest priority READY task to run. Always.
- The scheduler works on a First-In-First-Out discipline for tasks with the same priority.
- A task must switch to the *READY* state before being eligible for scheduling.
- A task will switch from *RUNNING* to *READY* if yielding or if being preempted by a higher priority task. Otherwise it can only go to a *WAITING* state, and eventually switch back to *READY*.
- When a task is preempted by a higher priority task, it switches from *RUNNING* to *READY* and is placed back on the *head* position of its Ready Queue. This means that it will be resumed as soon as it is the highest priority ready task again.
- On the contrary, if a task *yields*, it tells the scheduler it has completed its cycle. Then, it will be enqueued on the ready queue tail - the last queue position.
- When a task *waits* it is suspended until a condition is satisfied.
- When the condition is satisfied, it switches from *WAITING* to *READY*, and is enqueued on the tail.
- So, tasks with the same priority cooperate by either *yielding* or *waiting*.
- If a task never yields or waits, other tasks with the same or lower priority *will starve*.
- Finally, Tasks with the same priority are *initially* placed on the *Ready Queue* associated with that priority in the order they are *created*.



RK0 can handle context-switching with an extended frame when a float-point co-processor is available. This must be informed when compiling by defining the symbol `__FPU_PRESENT=1`.

Chapter 3. Timers and Delays

3.1. Busy-wait delay

A busy-wait delay `kDelay(t)` keeps a task spinning for t ticks. That is, the task does nothing but does not suspend or yield (but can be preempted). This service finds its use when simulating workloads.

Context switching is probably the most significant overhead on a kernel. The time spent on the System Tick handler contributes to much of this overhead.

Design Choice:

- Timers are kept on a single list; only the head element needs to be updated using a delta-queue approach.
- Application Timers that trigger callbacks are run on a deferred, non-preemptible system task.



Benefits:

- Keep the overhead of updating timers as minimal as possible with the delta queue;
- Deferring the Application Timer to a high-priority, non-preemptible system task meet the requested callback period while keeping the ability to track system ticks.

Timeout Node

Timeout Type

Absolute Interval (Ticks)

Relative Interval (Ticks)

Waiting Queue Address

Next Timeout Node

Previous Timeout Node

Every task is prone to events triggered by timers described in this section. Every Task Control Block has a node to a *timeout list*. This list is doubly linked treated a delta-sequence.

A set $T_{set} = \{(T1, 8), (T2, 6), (T3, 10)\}$ will be started at a relative time 0 as a sequence $T_{seq} = <(T2, 6), (T1, 2), (T3, 2)>$.

Thus, for every system tick, only the head element on the list needs to be decreased — yielding $O(1)$ on *decreasing*, that happens on the Hardware interrupt for the *System Tick*.

The ordering for the delta-queue is not $O(1)$, it is $O(n)$. A decrease happens on every SysTick interrupt; and the ordering happens only when adding a new node to the list.

3.2. Sleep Timers

There are three `sleep` primitives in RK0 — they behave differently.

3.2.1. Sleep Delay

The `sleepdelay()` (aliased as `sleep()`) puts a task to sleep for the exact number of `t` ticks on every call — no matter when the last call has happened.

Example:

```
VOID Task1(VOID* args)
{
    RK_UNUSEARGS
    UINT count = 0;
    while (1)
    {

        logPost("Task1: sleep");
        kSleep(300);
        /* wake here */
        count += 1U;
        if (count >= 5)
        {
            kDelay(25); /* spin */
            count=0;
            /* every 5 activations there will be a drift */
        }
    }
}
```

Output:

```
0 ms :: Task1: sleep
300 ms :: Task1: sleep <-- +300
600 ms :: Task1: sleep <-- +300
900 ms :: Task1: sleep <-- +300
1200 ms :: Task1: sleep <-- +300
1525 ms :: Task1: sleep <-- +325
1825 ms :: Task1: sleep <-- +300
2125 ms :: Task1: sleep <-- +300
2425 ms :: Task1: sleep
2725 ms :: Task1: sleep
3050 ms :: Task1: sleep
3350 ms :: Task1: sleep
3650 ms :: Task1: sleep
3950 ms :: Task1: sleep
4250 ms :: Task1: sleep
```

```
4575 ms :: Task1: sleep
```

3.2.2. Compensated Sleep Delays

These are suspensions that recompute the time considering the drift between calls. They are typically used to create periodic tasks with explicit periods. The general pattern is:

```
VOID Task(VOID *args)
{
    <initialisation>

    while(1)
    {
        <periodic code> /* this has an execution time */
        scheduleNext(PERIOD); /* cycle is finished: compute delay to keep PERIOD */
    }
}
```

The RMS algorithm considers all tasks have a common phase grid, that is when they are made READY, or are eligible to be scheduled. This means that from the very first activation, the time elapsed is already taken into account to when the next release should happen.

Nevertheless, it is common for kernels to provide sleep primitives that take into account a local anchor time normally set on the `<initialisation>` code—and will work as `waitFor(&anchor, PERIOD)`. RK0 provides both, and they suit different cases, as exposed below.

3.2.2.1. Periodic Release Sleep

`sleeprelease(P)` is used to delay a task so it is released on periodic rate. The sleep time is recalculated by the kernel on every call.

If the task wakes late by N ticks with $0 < N < P$, the kernel compensates by scheduling the next wake earlier (shortening the next sleep) so that over two periods the phase is preserved:

Say a task is expected to return from its k_{eth} sleep at $T_{k+1} = T_k + P$ [ticks]. If the task is resumed at $T_{k+1} = T_k + P + N$, upon detecting this drift, the kernel sets: $(T_{k+2} = T_{k+1} + P - N)$ for $N < P$.

This can be rewritten as:

$$(T_{k+2} = T_k + P + N + P - N) \leftrightarrow (T_{k+2} - T_k = 2P)$$

Example:

```
VOID Task1(VOID* args)
{
    RK_UNUSEARGS
    UINT count = 0;
    while (1)
    {
```

```

    logPost("Task1 released.");

    count += 1U;
    if (count >= 5)
    {
        kDelay(25); /* spin */
        count=0;
    }

    kSleepRelease(300); /*P=300 ticks; tick=1ms*/
}

}

```

Output:

```

.
.

/* R is release time */

1200 ms :: Task1: sleep periodic (R==4P)      (n)
1525 ms :: Task1: sleep periodic (5P<R<6P)   |
1800 ms :: Task1: sleep periodic (6P)          |
2100 ms :: Task1: sleep periodic (7P)          |
2400 ms :: Task1: sleep periodic (8P)          |
2700 ms :: Task1: sleep periodic (9P)          |
3025 ms :: Task1: sleep periodic (10P<R<11P) |
3300 ms :: Task1: sleep periodic (11P)          |
3600 ms :: Task1: sleep periodic (12P)          |
3900 ms :: Task1: sleep periodic (13P)          |
4200 ms :: Task1: sleep periodic (14P)          |
4525 ms :: Task1: sleep periodic (15P<R<16P) |
4800 ms :: Task1: sleep periodic (16P)          (m) m-n=12
.
.
.
----- Phase=3600=12xP

```

This mechanism is phase-locked. When the lateness is greater or equal to P it *skips* one or more releases to stay locked to the phase grid. In some sense, the `period` value can be seen as a *deadline*—if not met, the scheduler rejects to run on that activation.

`sleeprelease()` makes easier to perform worst-response time analysis on periodic tasks.



A set of periodic tasks must have priorities assigned properly (highest request rate, highest priority—the lower the period, the higher the priority). For `sleeprelease()` this is mandatory given the common phase grid.

3.2.2.2. Sleep Until (local anchor base)

`sleepuntil(anchor, period)` is somehow similar to `sleeprelease()`, but differs in two important aspects:

- The reference used to calculate how long to suspend to keep its rate is local to each task. It means the time before the first time the task is dispatched is dismissed.
- A late release longer than 1 period will return and run immediately. It prioritises execution count within a time-window - not the phase across releases.

The snippet below clearly demonstrates how each mechanism handles lateness that are longer than 1 period:

```
/* Every 3rd call both tasks will add up a delay longer than the task's period */

VOID HTask(VOID* args) /* higher priority: Period is 300 ticks */
{
    RK_UNUSEARGS
    UINT count = 0;
    while (1)
    {

        logPost("Higher: begin\r\n");
        /* wake here */
        count += 1U;
        kDelay(5);
        if (count >= 3)
        {
            kSleep(400); /* suspend */
            count=0;
        }
        logPost("Higher: end\r\n");
        kSleepRelease(300);

    }
}

VOID LTask(VOID *args) /* lower priority: Period is 400 ticks */
{
    RK_UNUSEARGS
    RK_TICK anchor = kTickGet();
    UINT count=0;
    while (1)
    {

        logPost("Lower: begin\r\n");
        /* wake here */
        count += 1U;
```

```

kDelay(5);
if (count >= 3)
{
    kSleep(500);
    count=0;
}
logPost("Lower: end\r\n");
kSleepUntil(&anchor, 400);
}
}

```

Output:

```

0 ms :: Higher: begin
5 ms :: Higher: end
5 ms :: Lower: begin
10 ms :: Lower: end
300 ms :: Higher: begin
305 ms :: Higher: end
405 ms :: Lower: begin
410 ms :: Lower: end

/* H 3rd run, expected next at 900ms */
600 ms :: Higher: begin

/* L 3rd run, expected next at 1205 ms */
805 ms :: Lower: begin

/* H Drift: 1005ms - 600ms = 405 ms > 300ms */
1005 ms :: Higher end

/* H released again @ next multiple of 300. */
1200 ms :: Higher: begin
1205 ms :: Higher: end

/* L Drift: 1310ms - 805ms = 505 ms > 400ms */
1310 ms :: Lower: end
/* it runs again immediately */
1310 ms :: Lower: begin

```



One normally does not write a code with periodic tasks expecting they will not keep their rate. But on the field a transient overload might cause it to happen. If it does, you choose the policy that is best-fit for your task: preserve phase (skip) or preserve execution count.

3.3. Blocking Time-out

These are internal timers associated with kernel calls that are blocking. Thus, establishing an upper-bound waiting time might benefit them. When the time for unblocking is up, the kernel call returns, indicating a timeout. This value is passed as a number of ticks.

When blocking is associated with a kernel object (other than the Task Control Block), the timeout node will store the object waiting for queue's address, so it can be removed if time expires.

A kernel call is made non-blocking, that is *try semantics*, by assigning the value `RK_NO_WAIT`, the function returns immediately if unsuccessful. The value `RK_WAIT_FOREVER` suspends a task indefinitely until the condition is satisfied. Timeout values above `RK_MAX_PERIOD` are invalid.

In practice, we often block either using `RK_WAIT_FOREVER` or do not block (*try semantics*, `RK_NO_WAIT`).

Use a bounded timeout only when you expect occasional misses and you know how to handle them.



Importantly, *an ISR shall never block*. Any blocking call from an ISR is invalid and triggers fault handling when error checking is enabled.

3.4. Callout Timers (Application Timers)

Timer Control Block

Option: Reload/One-Shot

Phase (Initial Delay)

Callout Function Pointer

Callout Argument

Timeout Node

These are Application Timers that will issue a callback when expiring. In addition to a callout function, an Application Timer receives an initial phase delay and a period and can choose to run once (one-shot) or auto-reload itself.

The callback runs within a System Task with priority 0 and is non-preemptible, which makes the scheduler prioritise it over other tasks. Callouts must be short and unblocking, as they can cause high CPU contention.

Application Timers (with autoreload) will keep track of delays in between activations.

3.5. System Tick

A dedicated peripheral that generates an interrupt after a defined period provides the kernel time reference. For ARMv6/7M, this peripheral is the built-in SysTick, a 24-bit counter timer. The handler performs some housekeeping on every tick and assesses the need to call a context switch.

The 'housekeeping' accounts for global timer tracking and any tick-dependent condition that might

change a task status. When a timer expires, it may switch a task from **WAITING** to **READY** or schedule a callback. For callbacks, execution is deferred to the **TimHandleTask** system task, where callbacks run and timer state is updated.

Note that tasks might switch from **WAITING** to **READY** for reasons other than tick-related. In these cases, context switching might be triggered immediately if the readied task can preempt the running task.

3.6. System Tasks

There are two *System Tasks*: the *Idle Task* and the *TimHandleTask*.

The *TimHandleTask* is currently used for timer callouts. Nominally its priority is **0**, but in practice it can be seen as priority **-1**, because it always takes precedence over user tasks with priority **0**. *It cannot be preempted* by user tasks, so timer callouts must be short. If an application timer callback needs heavier processing than a simple time-mark, use a periodic task with **kSleepPeriodic()** and defer that work.

The *Idle Task* runs whenever there is no other ready task to be dispatched. The CPU enters on *low-power*. The kernel assigns the *Idle Task* priority during initialisation, taking into account all priorities the user has defined. Unless user tasks occupy all 32 priorities, the Idle Task is treated as an ordinary lowest priority and has a position in the ready queue table. Otherwise, it is selected if *Ready Queue Bitmap* is **0x00000000**.

Chapter 4. Memory Allocator

Memory Allocator Control Block

Associated Block Pool

Number of Blocks

Block Size

Number of Free Blocks

Free Block List

The standard C library `malloc()` leads to fragmentation and (also, because of that) is highly indeterministic. Unless we use it once - to allocate memory before starting up, it doesn't fit. But often, we need to 'multiplex' memory amongst tasks over time, that is, to dynamically allocate and deallocate.

To avoid fragmentation, we use fixed-size memory blocks. A simple approach would be a static table marking each block as free or taken. With this pattern, you will need to 'search' for the next available block, if any - the time for searching changes - bounding this search to a maximum number of blocks, or $O(n)$. To optimise, an approach is to keep track of what is free using a dynamic table—a linked list of addresses. Now we have $O(1)$.

We use "meta-data" to initialise the linked list. Every address holds the "next" address value. All addresses are within the range of a pool of fixed-size blocks. This approach limits the minimal size of a block to the size of a memory address—32 bits for our supported architecture.

Yet, this is the cheapest way to store meta-data. If not stored on the empty address itself, an extra 32-bit variable would be needed for each block, so it could have a size of less than 32 bits.

Allocating memory at runtime is a major source of latency (1), indeterministic (2) behaviour, and footprint overhead (3).



Design choice: the allocator's design achieves low-cost, deterministic, fragmentation-free memory management by using fixed-size word-aligned block sizes (1)(2) and embedding metadata within the memory blocks themselves (3).

Benefits: Run-time memory allocation benefits have no real-time drawbacks.



The kernel will always round up the block size to the next multiple of 4. Say the user creates a memory pool, assigning blocks to be 6-byte wide; they will turn into 8-byte blocks.

4.1. How it works

When a routine calls `alloc()`, the address to be returned is the one a "free list" is pointing to, say `addr1`. Before returning `addr1` to the caller, we update the free list to point to the value stored within `addr1` - say `addr8` at that moment.

When a routine calls `free(addr1)`, we overwrite whatever has been written in `addr1` with the value-free list point to (if no more `alloc()` were issued, it would still be `addr8`), and `addr1` becomes the free list head again.

Allocating and deallocating fixed-size blocks using this structure and storing meta-data this way is as deterministic ($O(1)$) and economical as we can get for dynamic memory allocation.

Chapter 5. Inter-Task Communication

RK0 Blog:

- [About Inter-Task Communication - Part 1](#)
- [About Inter-Task Communication - Part 2](#)

Inter-Task Communication (ITC) refers to the mechanisms that enable tasks to coordinate/cooperate/synchronise by means of sending or receiving information that falls into two logical categories: *Signals* or *Messages*.

- **Signals:** A *Signal* is either present or absent. Its meaning is implicit.
- **Messages:** A *Message* is a means of coordinating and exchanging information altogether. Different from a *Signal*, a Message needs to be parsed (interpreted).

5.1. Task Events (Task Event Flags)

Within Task Control Block

Event Register (32 flags)

Required Signal Flags

Options (ALL/ANY)

Each Task Control Block stores a 32-bit event register (one event per bit). A set bit means an event is pending.

The API is bit-mask oriented:

- `kTaskEventSet(task, mask)` posts flags with OR semantics.
- `kTaskEventGet(required, RK_EVENT_ANY|RK_EVENT_ALL, &got, timeout)` waits until the required pattern is met.
- `kTaskEventQuery(task, &flags)` reads current flags.
- `kTaskEventClear(task, mask)` clears selected flags.

If the required pattern is not met, the task can block (logical state: `PENDING`) until timeout or until another task posts matching flags.

On successful `get()`, the matched required positions are cleared from the task event register.

`0x00` is invalid as a `set()` mask and as `required` for `get()`.

5.1.1. Usage Example: Supervisor Task

One common pattern is to start each loop by checking pending events. In a supervisor task, this becomes an event hub.

```

typedef struct
{
    ULONG pendingBit;
    RK_TASK_HANDLE dstTask;
    ULONG dstSignal;
} Route_t;

static const Route_t routes[] =
{
{
    PENDING_AIRFLOW_INCREASE,
    airFlowTaskHandle,
    AIRFLOW_INCREASE_SIGNAL
},
{
    PENDING_TEMP_DECREASE,
    tempTaskHandle,
    TEMP_DECREASE_SIGNAL
},
/* more routes */
};

VOID SupervisorTask(VOID *args)
{
    RK_UNUSEARGS;

    while(1)

    {
        ULONG gotFlags = 0UL;

        RK_ERR err = kTaskEventGet(0xFFFF,
                                   RK_FLAGS_ANY,
                                   &gotFlags,
                                   SUPERVISOR_T_PERIOD);

        if (err == RK_ERR_SUCCESS && gotFlags != 0)
        {
            for (ULONG i = 0; i < ARRAY_LEN(routes); ++i)
            {
                if (gotFlags & routes[i].pendingBit)
                {
                    kTaskEventSet(routes[i].dstTask, routes[i].dstSignal);
                }
            }
        }

        /* if there is anything to do if time out */
    }
}

```

Task Events are the only ITC primitive that cannot be disabled; they are a *Core Mechanism*.

RK0 does not implement BSD/UNIX-like asynchronous signals (POSIX style).



Tasks explicitly wait (`get()`) for bit patterns, and others `set()` bit patterns, making the scheduler aware of all blocking conditions.

Thus, unlike traditional asynchronous signal mechanisms (as BSD/UNIX Signals), no code is ever executed at an arbitrary instruction boundary.

5.2. Semaphore

Semaphore Control Block

Counter (Unsigned Integer)

Maximum Value

Waiting Queue

A semaphore S is a nonnegative integer variable, apart from the operations it is subjected to. S is initialized to a nonnegative value. The two operations, called P and V, are defined as follows:

`P(S): if $S > 0$ then $S := S-1$, else the process is suspended until $S > 0$.`

`V(S): if there are processes waiting, then one of them is resumed; else $S := S+1$.`

(Dijkstra, 1968)

Semaphores are *public* kernel objects for signalling and waiting on countable events.

`V()` in RK0 semaphores maps to `post()` and `P()` to `pend()`.

5.2.1. Counting Semaphore and Binary Semaphores

The typical use case for semaphores is as a "credit tracker": use `pend()` to consume a credit and `post()` to return a credit (for example, free slots in a queue). These are *Counting Semaphores*.

A *Binary Semaphore* is a counting semaphore with maximum value 1: the state is either available or unavailable.

Binary semaphores are often used for task-to-task or ISR-to-task synchronisation, and sometimes for mutual exclusion (with caveats discussed later).

5.2.2. Semaphores in RK0

To initialise a semaphore in RK0, provide two values: initial count and maximum count. When the counter is at maximum, `post()` does not increment it and returns `RK_ERR_SEMA_FULL`.

A *Binary Semaphore* is therefore created by setting maximum count to 1. A counting semaphore that is intended never to saturate can use `UINT32_MAX`.

Besides `init()`, `post()`, and `pend()`, `query()` inspects current state: non-negative means current count; negative means number of tasks waiting.

5.2.3. Flushing or Broadcast

A `flush()` is a *broadcast* signal: it releases all tasks currently pending on the semaphore. If successful, the semaphore counter remains `0`.

In current RK0 behaviour, semaphore flush is synchronous and cannot be called from ISRs.



Flushing/broadcasting must be used with care. Releasing many tasks at once can increase scheduling pressure and hurt response time.

Bounded Buffer with Semaphores

Items are buffered within a memory region whose capacity is `K` items.

Thus: $0 \leq (\text{Number of Inserted}) - (\text{Number of Extracted}) \leq K$.

Using semaphores the pattern is as follows:

1. A semaphore with `K` tokens to track free slots, preventing producers from proceeding when there are no free slots.
2. Another semaphore, with `K` tokens, for the number of items, not allowing consumers to proceed if there are no items.
3. A 1-token semaphore so only one task manipulates the buffer at a time.

```
/* a ring buffer of items */
#define BUFSIZ (K)
static ITEM_t buf[BUFSIZ]={0};
static UINT getIdx = 0U;
static UINT putIdx = 0U;
/* this indexes==0 could either mean FULL or EMPTY for a regular
circular buffer with wrap-around.
```

```
With semaphores the state is well defined.
```

```
*/
```

```
RK_SEMAPHORE itemSema;
RK_SEMAPHORE slotSema;
RK_SEMAPHORE acquireSema;
```

```
VOID kApplicationInit(VOID)
{
```

```
/*buffer is initialised empty */
kSemaphoreInit
( &itemSema,
0, /* no item */
```

```

        K /*max items */
    );

kSemaphoreInit
(   &slotSema,
    K, /* K free slots */
    K /* max slots */
);

/* and free */
kSemaphoreInit
(   &acquireSema,
    1, /* free to access */
    1 /* 1 max task allowed */
);

VOID PutItem(ITEM_t* insertItemPtr)
{
    /* wait for room */
    kSemaphorePend(&slotSema, RK_WAIT_FOREVER);

    /* wait for availability */

    kSemaphorePend(&acquireSema, RK_WAIT_FOREVER);
    buf[putIdx] = *insertItemPtr;
    putIdx += 1U; putIdx %= BUFSIZ;
    /* signal availability */
    kSemaphorePost(&acquireSema);

    /* signal item */
    kSemaphorePost(&itemSema);
}

VOID GetItem(ITEM_t* extractItemPtr )
{
    /* wait for an item */
    kSemaphorePend(&item, RK_WAIT_FOREVER);

    /* wait for availability */
    kSemaphorePend(&acquireSema, RK_WAIT_FOREVER);

    *extractItemPtr = buf[getIdx];
    getIdx+=1U; getIdx %= BUFSIZ;

    /* signal availability */
    kSemaphorePost(&acquireSema);

    /* signal room */
}

```

```
    kSemaphorePost(&slotSema);  
}
```

The solution above has *Put()* and *Get()* as blocking methods.

If the producer and the consumer run at different rates, eventually, they will synchronise to the lowest rate.

The numbers below are from a run with a buffer of 32 items (integers being incremented are the produced data).

The producer is twice faster than the consumer. Initially at every 2 insertions there is a single remove. Eventually, the buffer is filled up, and tasks run at lockstep (*put, get, put, get...*), at the consumer pace:

```
Put 59 <-  
Put 60 <-  
-----  
Got 30 ->  
-----  
Put 61 <-  
Put 62 <-  
-----  
Got 31 ->  
-----  
Put 63 <-  
Put 64 <-  
-----  
Got 32 | ->  
Put 65 . <-  
    <x>[Full Queue, Producer blocks]  
Got 33 | -> [Consumer unblocks producer...]  
Put 66 . <-  
    <x>[Full Queue]  
Got 34 | -> [Consumer unblocks producer...]  
Put 67 . <-  
    <x>[Full Queue]
```

5.2.4. Mutex Semaphore (Locks)

Mutex Control Block

Locked State (Boolean)

Owner

Protocol Flag (RK_NO_INHERIT / RK_INHERIT)

Mutex Control Block

Waiting Queue

Mutex Node (list node within the owner TCB)

Some regions are critical and must not be executed by more than one task at once. Acquiring (`lock()`) a mutex before entering and releasing (`unlock()`) after leaving makes the region mutually exclusive.

A *Mutex* is a binary lock with ownership: once a task acquires it, that task is the owner, and only the owner can release it.

If a task tries to acquire a locked mutex, it switches to `BLOCKED` until the owner unlocks it. When released, the highest-priority waiter is dequeued first.

Unlike semaphores, unlocking by a non-owner is invalid and rejected.

Mutexes are only for mutual exclusion; they are not signalling primitives.

PS: RK0 mutexes are non-recursive. Re-entrant locking of the same mutex returns `RK_ERR_MUTEX_REC_LOCK`.

Mutex protocol modes in RK0:

- `RK_NO_INHERIT`: no priority protocol.
- `RK_INHERIT`: Priority Inheritance Protocol (PIP).

5.2.4.1. Priority Inversion and PIP

Let TH, TM, and TL be three tasks with priority high (H), medium (M) and low (L), respectively. Say TH is dispatched and blocks on a mutex that 'TL' has acquired (i.e.: "*TL is blocking TH*").

If 'TM' does not need the resource, it will run and preempt 'TL'. And, by transition, 'TH'.

From now on, 'TH' has an *unbounded waiting time* because any task with priority higher than 'L' that does not need the resource indirectly prevents it from being unblocked — *awful*.

The Priority Inheritance (PI) Protocol avoids this unbounded waiting. It is characterised by an invariant, simply put:

PIP Invariant: *At any instant a Task assumes the highest priority among the tasks it is blocking.*

If employed on the situation described above, task TM cannot preempt TL, whose effective priority would have been raised to 'H'.

It is straightforward to reason about this when you consider the scenario of a single mutex.

When locks nest, the protocol also needs to be:

- Transitive: if T1 blocks T2 and T2 blocks T3, the highest priority (T3) must propagate back to T1 through T2.

This is the hard part of a correct implementation: updates must preserve the invariant across changing wait chains and multiple mutexes.



This blog shows an even more intricate case of priority inversion handling.

[RK0 Blog: Transitive Priority Inheritance on RK0 mutexes](#)

Below, a PIP-only case in which locks nest:

```

/* Task1 has the Highest nominal priority */
/* Task2 has the Medium nominal priority */
/* Task3 has Lowest nominal priority */

/* Note Task3 starts as 1 and 2 are delayed */

RK_DECLARE_TASK(task1Handle, Task1, stack1, STACKSIZE)
RK_DECLARE_TASK(task2Handle, Task2, stack2, STACKSIZE)
RK_DECLARE_TASK(task3Handle, Task3, stack3, STACKSIZE)

RK_MUTEX mutexA;
RK_MUTEX mutexB;

VOID kApplicationInit(VOID)
{
    K_ASSERT(!kCreateTask(&task1Handle, Task1, RK_NO_ARGS, "Task1", stack1, \
        STACKSIZE, 1, RK_PREEMPT));
    K_ASSERT(!kCreateTask(&task2Handle, Task2, RK_NO_ARGS, "Task2", stack2, \
        STACKSIZE, 2, RK_PREEMPT));
    K_ASSERT(!kCreateTask(&task3Handle, Task3, RK_NO_ARGS, "Task3", stack3, \
        STACKSIZE, 3, RK_PREEMPT));

    /* mutexes initialised with priority inheritance enabled */
    kMutexInit(&mutexA, RK_INHERIT);
    kMutexInit(&mutexB, RK_INHERIT);
}

VOID Task3(VOID *args)
{
    RK_UNUSEARGS
    while (1)
    {
        printf("@ %lums: [TL] Attempting to LOCK 'A' | Eff: %d | Nom: %d\r\n",
kTickGet(),
            runPtr->priority, runPtr->prioNominal);
}

```

```

    kMutexLock(&mutexA, RK_WAIT_FOREVER);

    printf("@%lums: [TL] LOCKED 'A' (in CS) | Eff: %d | Nom: %d\r\n",
           kTickGet(),
           runPtr->priority, runPtr->prioNominal);

    kDelay(60); /* <- important */

    printf("@%lums: [TL] About to UNLOCK 'A' | Eff: %d | Nom: %d\r\n",
           kTickGet(),
           runPtr->priority, runPtr->prioNominal);

    kMutexUnlock(&mutexA);

    printf("--->");
    printf("@%lums: [TL] Exit CS | Eff: %d | Nom: %d\r\n",
           kTickGet(),
           runPtr->priority, runPtr->prioNominal);

    kSleep(4);
}

}

VOID Task2(VOID *args)
{
    RK_UNUSEARGS
    while (1)
    {
        kSleep(5);

        printf("@%lums: [TM] Attempting to LOCK 'B' | Eff: %d | Nom: %d\r\n",
               kTickGet(),
               runPtr->priority, runPtr->prioNominal);
        kMutexLock(&mutexB, RK_WAIT_FOREVER);

        printf("@%lums: [TM] LOCKED 'B', now trying to LOCK 'A' | Eff: %d | Nom:
               %d\r\n",
               kTickGet(), runPtr->priority, runPtr->prioNominal);
        kMutexLock(&mutexA, RK_WAIT_FOREVER);

        printf("@%lums: [TM] LOCKED 'A' (in CS) | Eff: %d | Nom: %d\r\n",
               kTickGet(),
               runPtr->priority, runPtr->prioNominal);
        kMutexUnlock(&mutexA);

        printf("@%lums: [TM] UNLOCKING 'B' | Eff: %d | Nom: %d\r\n",
               kTickGet(),
               runPtr->priority, runPtr->prioNominal);

        kMutexUnlock(&mutexB);

        printf("--->");

        printf("@%lums: [TM] Exit CS | Eff: %d | Nom: %d\r\n",
               kTickGet(),
               runPtr->priority, runPtr->prioNominal);

```

```

    }

VOID Task1(VOID *args)
{
    RK_UNUSEARGS
    while (1)
    {
        kSleep(2);

        printf("@%lums: [TH] Attempting to LOCK 'B' | Eff: %d | Nom: %d\r\n",
kTickGet(),
            runPtr->priority, runPtr->prioNominal);

        kMutexLock(&mutexB, RK_WAIT_FOREVER);

        printf("@%lums: [TH] LOCKED 'B' (in CS) | Eff: %d | Nom: %d\r\n",
kTickGet(),
            runPtr->priority, runPtr->prioNominal);

        kMutexUnlock(&mutexB);

        printf("--->");

        printf("@%lums: [TH] Exit CS | Eff: %d | Nom: %d\r\n",
kTickGet(),
            runPtr->priority, runPtr->prioNominal);
    }
}

```

Result and comments:

```

>>> TL locks 'A'. Higher priority tasks are sleeping. <<<
@ 14720ms: [TL] Attempting to LOCK 'A' | Eff: 3 | Nom: 3
@ 14720ms: [TL] LOCKED 'A' (in CS) | Eff: 3 | Nom: 3

@14721ms: [TM] Attempting to LOCK 'B' | Eff: 2 | Nom: 2

>>> TM acquires 'B' and is blocked by TL on 'A'. TL inherits TM's priority. <<<
@14721ms: [TM] LOCKED 'B', now trying to LOCK 'A' | Eff: 2 | Nom: 2

>>> TH will blocked by TM on 'B': <<<
@14722ms: [TH] Attempting to LOCK 'B'| Eff: 1 | Nom: 1

>>> TM inherits TH's priority. TL inherits TH's priority via TM. <<<
@14780ms: [TL] About to UNLOCK 'A' | Eff: 1 | Nom: 3

>>> Upon unlocking 'A', TL is preempted by TM. It means TL's priority has been

```

```
restored, as it is no longer blocking a higher priority task. <<<
```

```
>>> Now TM acquires 'A' <<<
```

```
@14780ms: [TM] LOCKED 'A' (in CS) | Eff: 1 | Nom: 2
```

```
>>> After releasing 'A', but before releasing 'B', TM's priority is still '1', as it  
is blocking TH while holding 'B'. <<<
```

```
@14780ms: [TM] UNLOCKING 'B' | Eff: 1 | Nom: 2
```

```
>>> Upon unlocking 'B' TM is preempted by TH. (TM's priority has been restored.) <<<
```

```
@14780ms: [TH] LOCKED 'B' (in CS) | Eff: 1 | Nom: 1
```

```
>>> RESULT: even though priority inversion was enforced, tasks leave the nested lock  
ordered by their nominal priority. <<
```

```
-->@14780ms: [TH] Exit CS | Eff: 1 | Nom: 1
```

```
-->@14780ms: [TM] Exit CS | Eff: 2 | Nom: 2
```

```
-->@14780ms: [TL] Exit CS | Eff: 3 | Nom: 3
```

Importantly, the worst-case time is bounded by the time the lowest priority task holds a lock (60 ms in the example: 14720ms → 14780ms).

As for each priority update we check each waiting queue for each mutex a task owns, the time-complexity is linear $O(\text{owner} * \text{mutex})$. But, typically no task ever holds more than a few mutexes. Yet, one should not be encouraged to nest locks if not needed.

5.2.5. Mutexes vs Binary Semaphores

There is (or used to be) a lot of fuss about whether binary semaphores are appropriate to use as locks. As a practical guideline, if all tasks sharing the resource have the same priority, using a binary semaphore *can be appropriate*—because a binary semaphore is considerably faster. It all depends on the case.

The drawback is the lack of ownership: any task can accidentally release the resource. On a large codebase, this can become a real problem. Nonetheless, this is a problem for semaphores in general.

For tasks with different priorities, binary semaphores should never be considered for mutual exclusion unless priority inversion is not a problem (how?).

Counting semaphores initialised as 1 is too risky. Besides the priority inversion, if the count ever increases above 1, mutual exclusion is lost, and multiple tasks can enter the critical section at once.

5.3. Scheduler Lock

Often, we need a task to perform operations without being preempted. A mutex serialises access to a code region but does not prevent a task from being preempted while operating on data. Depending on the case, this can lead to inconsistent data state.

An aggressive way is to disable interrupts globally. For kernel services often it is the only way to keep data integrity. On the higher level it is feasible for very short operations and/or when you need to protect data from interrupts altogether.

A less aggressive approach is to make the task non-preemptible with `kSchLock()` before entering the critical region and `kSchUnlock()` when leaving. This way, interrupts are still being sensed, and even higher-priority tasks might switch to a ready state, but the running thread will not be preempted.

The priority inversion it potentially causes is bounded. If a higher-priority task is readied while the scheduler is locked, the context switch happens immediately after unlocking.

Note that for locking/unlocking the scheduler the global interrupts will be disabled for the time to increment/decrement a counter, therefore, if your atomic operation is as short as that (3 to 4 cycles), disabling/enabling global interrupts is a better alternative.

To add to the discussion, when two threads need to access the same data to 'read-modify-write', a lock-free mechanism is the LDREX/STREX operations of ARMv7M (or more generally C11 atomics). They do not avoid preemptions, and particularly in ARMv7m, if the data is touched by an ISR before the store-exclusive concludes, the ownership is lost. Typically used for multi-core spin-locking.

5.4. Sleep Queue

Sleep Queue Control Block

Task Waiting Queue

`RK_SLEEP_QUEUE` is a queue of tasks sleeping until explicitly woken.

This primitive was formerly named `RK_EVENT`. Queue names usually reflect the condition they represent (for example, `notFull`, `notEmpty`, `goWriters`, `goReaders`).

An `RK_SLEEP_QUEUE` is stateless: it does not record whether an event has happened. Therefore `kSleepQueueWait(&q, timeout)` always attempts to sleep the caller. `RK_NO_WAIT` is effectively a no-op on this primitive.

`kSleepQueueSignal(&q)` wakes one waiter (highest priority first). `kSleepQueueWake(&q, n, &u)` wakes at most `n` waiters. `kSleepQueueFlush(&q)` is the wake-all form (`n = 0`).

`kSleepQueueQuery(&q, &n)` returns the number of waiting tasks.

`kSleepQueueSuspend(&q, taskHandle)` can move a `READY` task to that sleep queue (task context or ISR).

`kSleepQueueReady(&q, taskHandle)` does the opposite, moving a specific task from that sleep queue to READY.

`kSleepQueueWake()` is synchronous and executes immediately in thread context. It cannot be called from ISR context.

Because Sleep Queues are stateless, they are prone to lost wake-ups when used alone. They are best used as building blocks for Conditional Critical Regions (monitor-like constructs).

5.5. Conditional Critical Regions

Task Events and semaphores work by atomically updating state and testing predicates that control execution flow (for example, `pend` on a semaphore with count `0` blocks the caller).

A critical region guarded by a lock is either free or taken. What if we need to wait on a richer condition?

We can combine a Mutex and a Sleep Queue, but there is an important detail: operations on both must be coordinated atomically.

When that combination is wrapped as data plus methods, it is a *Monitor*.

5.5.1. Monitor Invariants

1. a single task can be active within a monitor;
2. only the active task within a monitor can check or change its state.

Since tasks signal or sleep *within* a monitor based on its *internal* state, and only one task can be active, we need to establish how to keep both invariants true.

5.5.2. Signalling Discipline

At any moment a single task can be active within a monitor. When the sleeping task is signalled, there are 3 common disciplines to follow: signal-and-leave (Hansen), signal-and-wait (Hoare) or signal-and-continue (Mesa).

Arguably, the most common is *signal-and-continue* - rather than leaving or suspending itself the active task might continue within the monitor. That is possible if the active task holds a lock the waking task needs to acquire to enter. Upon leaving, the active task must release the lock.

The major implication is that by the time the woken task enters the monitor, the condition it was waiting for might no longer be true. It sounds odd because a Monitor is about encapsulating a conditional critical region so no outsiders change its state. But, either a flush, a bad design—or a preemption anomaly—can violate that somehow.


Mesa Monitor has a typical test-loop pattern:

```

--- snippet ---
while (condition.isFalse())
{
    /*unlock-wait sequence*/
    DISABLE_PREEMPTION
    unlock(mutex);
    sleep(condition)
    ENABLE_PREEMPTION
    lock(mutex);
    /* when waking, the while clause is tested again */
}
--- snippet ---

```

The figure below shows a producer-consumer problem if implemented on a Monitor-idiom.



The **UNLOCK-SLEEP** sequence within the testing loop has preemption disabled because after releasing the lock, the task cannot be allowed to resume within the monitor again, for any reason that is not the monitor predicate being satisfied.

5.5.3. Condition Variable Model for (Mesa) Monitors

The *Condition Variable Model* allows a task to wait within a monitor-construct and if active, operate using `signal()`, `wait()` and `broadcast()` respecting the monitor invariant.

Sleep Queues are actually the canonical *Condition Variable* introduced by Hoare:

Note that a condition "variable" is neither true nor false; indeed, it does not have any stored value accessible to the program. In practice, a condition variable will be represented by an (initially empty) queue of processes which are currently waiting on the condition; but this queue is invisible both to waiters and signallers. This design of the condition variable has been deliberately kept as primitive and rudimentary as possible (...)

(Monitors: An Operating System Structuring Concept, Hoare, 1974)

RKO follows *pthreads condition variable* semantics, aligned with Mesa monitors.

The difference is that there is no standalone *CondVar* primitive. The programmer combines Sleep Queues and Mutexes and uses these helpers:

- `kCondVarWait(&sleepq, &mutex, timeout)`
- `kCondVarSignal(&sleepq)`
- `kCondVarBroadcast(&sleepq)`

The `condWait` is the real helper. When using it, a Mesa testing-loop reduces to:

```
while(!condition)
{
    kCondVarWait(&condQueue, &monitorLock, timeout);
}
```

If you need a monitor policy different from Mesa, you can build it from the same primitives.



`kCondVarWait`, `kCondVarSignal`, and `kCondVarBroadcast` are task-context APIs and cannot be called from ISRs.

5.5.3.1. Usage Example: Synchronisation Barrier

A given number of tasks must reach a point in the program before *all can proceed*, so every task calls a `barrWait(&barrier)` to catch up with the set of tasks it must synchronise.

The last task entering the barrier will broadcast a signal to all tasks waiting for the wake condition.

At any moment within a Monitor a single task is RUNNING (what is an invariant of the kernel), all other tasks within the monitor are either SLEEPING (for some condition) or BLOCKED (on a mutex).

```
/* Synchronisation Barrier */

typedef struct
{
    RK_MUTEX lock;
    RK_SLEEP_QUEUE allSynch;
    UINT count; /* number of tasks in the barrier */
    UINT round; /* increased every time all tasks synch */
    UINT nRequired; /* number of tasks required */
} Barrier_t;

VOID BarrierInit(Barrier_t *const barPtr, UINT nRequired)
{
    kMutexInit(&barPtr->lock, RK_INHERIT);
    kSleepQueueInit(&barPtr->allSynch);
    barPtr->count = 0;
    barPtr->round = 0;
    barPtr->nRequired = nRequired;
```

```

}

VOID BarrierWait(Barrier_t *const barPtr)
{
    UINT myRound = 0;
    kMutexLock(&barPtr->lock, RK_WAIT_FOREVER);

    /* save round number */
    myRound = barPtr->round;
    /* increase count on this round */
    barPtr->count++;

    if (barPtr->count == barPtr->nRequired)
    {
        /* reset counter, inc round, broadcast to sleeping tasks */
        barPtr->round++;
        barPtr->count = 0;
        kCondVarBroadcast(&barPtr->allSynch);
    }
    else
    {
        /* sequence: a proper wake signal might happen after inc round */
        while ((UINT)(barPtr->round - myRound) == 0U)
        {
            RK_ERR err = kCondVarWait(&barPtr->allSynch, &barPtr->lock,
RK_WAIT_FOREVER);
            K_ASSERT(err==RK_ERR_SUCCESS);
        }
    }
    kMutexUnlock(&barPtr->lock);
}

#define N_REQUIRED 3

Barrier_t syncBarrier;

VOID kApplicationInit(VOID)
{
    K_ASSERT(!kCreateTask(&task1Handle, Task1, RK_NO_ARGS, "Task1", stack1, STACKSIZE,
2, RK_PREEMPT));
    K_ASSERT(!kCreateTask(&task2Handle, Task2, RK_NO_ARGS, "Task2", stack2, STACKSIZE,
3, RK_PREEMPT));
    K_ASSERT(!kCreateTask(&task3Handle, Task3, RK_NO_ARGS, "Task3", stack3, STACKSIZE,
1, RK_PREEMPT));
    BarrierInit(&syncBarrier, N_REQUIRED);
}

VOID Task1(VOID* args)
{
    RK_UNUSEARGS
}

```

```
while (1)
{
    kPuts("Task 1 is waiting at the barrier...\r\n");
    BarrierWait(&syncBarrier);
    kPuts("Task 1 passed the barrier!\r\n");
    kSleep(8);
}
}

VOID Task2(VOID* args)
{
    RK_UNUSEARGS
    while (1)
    {
        kPuts("Task 2 is waiting at the barrier...\r\n");
        BarrierWait(&syncBarrier);
        kPuts("Task 2 passed the barrier!\r\n");
        kSleep(5);
    }
}

VOID Task3(VOID* args)
{
    RK_UNUSEARGS
    while (1)
    {
        kPuts("Task 3 is waiting at the barrier...\r\n");
        BarrierWait(&syncBarrier);
        kPuts("Task 3 passed the barrier!\r\n");
        kSleep(3);
    }
}
```

```

Task 3 is waiting at the barrier...
Task 2 is waiting at the barrier...
Task 1 is waiting at the barrier...
Task 3 passed the barrier!
Task 1 passed the barrier!
Task 2 passed the barrier!
Task 3 is waiting at the barrier...
Task 2 is waiting at the barrier...
Task 1 is waiting at the barrier...
Task 3 passed the barrier!
Task 1 passed the barrier!
Task 2 passed the barrier!
Task 3 is waiting at the barrier...
Task 2 is waiting at the barrier...
Task 1 is waiting at the barrier...
Task 3 passed the barrier!
Task 1 passed the barrier!
Task 2 passed the barrier!
Task 3 is waiting at the barrier...
Task 2 is waiting at the barrier...
Task 1 is waiting at the barrier...
Task 3 passed the barrier!
Task 1 passed the barrier!
Task 2 passed the barrier!

```



Monitor/CondVars are not to be signalled by ISRs. Because as they do not accumulate and an ISR is not a normal task, it can't be within a monitor and a signal might be lost.

5.5.3.2. Usage Example: Readers Writers Lock

Several readers and writers share a piece of memory. Readers can concurrently access the memory to read; a single writer is allowed (otherwise, data would be corrupted).

When a writer finishes, it checks for any readers waiting. If there is, the writer flushes the readers waiting queue. If not, it wakes a single writer, if any. When the last reader finishes, it signals a writer.

Every read or write operation begins with an acquire and finishes with a release.

*PS: This RWLock implementation has a **reader-preference** policy, as when a writer finishes, it flushes sleeping readers. When the last reader finishes, it will signal writer waiting queue.*

```

/* RW-Lock */

/* a single writer is allowed if there are no readers */
/* several readers are allowed if there is no writer*/
typedef struct
{
    RK_MUTEX      lock;
    RK_SLEEP_QUEUE  writersGo;
    RK_SLEEP_QUEUE  readersGo;
    INT           rwCount; /* number of active readers if > 0 */
                           /* active writer if -1 */

```

```

}RwLock_t;

VOID RwLockInit(RwLock_t *const rwLockPtr)
{
    kMutexInit(&rwLockPtr->lock, RK_INHERIT);
    kSleepQueueInit(&rwLockPtr->writersGo);
    kSleepQueueInit(&rwLockPtr->readersGo);
    rwLockPtr->rwCount = 0;
}

/* A writer can acquire if rwCount = 0 */
/* An active writer is indicated by rwCount = -1; */
VOID RwLockAcquireWrite(RwLock_t *const rwLockPtr)
{
    kMutexLock(&rwLockPtr->lock, RK_WAIT_FOREVER);
    /* if different than 0, there are either writers or readers */
    /* sleep to be signalled */
    while (rwLockPtr->rwCount != 0)
    {
        kCondVarWait(&rwLockPtr->writersGo, &rwLockPtr->lock, RK_WAIT_FOREVER);
        /* mutex is locked when waking up*/
    }
    /* woke here, set an active writer */
    rwLockPtr->rwCount = -1;
    kMutexUnlock(&rwLockPtr->lock);
}

/* a writer releases, waking up all waiting readers, if any */
/* if there are no readers, a writer can get in */
VOID RwLockReleaseWrite(RwLock_t *const rwLockPtr)
{
    kMutexLock(&rwLockPtr->lock, RK_WAIT_FOREVER);

    rwLockPtr->rwCount = 0; /* indicate no writers*/

    /* if there are waiting readers, flush */
    ULONG nWaitingReaders=0;
    kSleepQueueQuery(&rwLockPtr->readersGo, &nWaitingReaders);
    if (nWaitingReaders > 0)
    {
        /* condVarBroadcast is just an alias for an event flush */
        kCondVarBroadcast(&rwLockPtr->readersGo);
    }
    else
    {
        /* wake up a single writer if any */
        kCondVarSignal(&rwLockPtr->writersGo);
    }
    kMutexUnlock(&rwLockPtr->lock);
}

```

```

/* a reader can acquire if there are no writers */
VOID RwLockAcquireRead(RwLock_t *const rwLockPtr)
{
    kMutexLock(&rwLockPtr->lock, RK_WAIT_FOREVER);
    /* if there is an active writer, sleep */
    while (rwLockPtr->rwCount < 0)
    {
        kCondVarWait(&rwLockPtr->readersGo, &rwLockPtr->lock, RK_WAIT_FOREVER);
        /* mutex is locked when waking up*/
    }
    /* increase rwCount, so its > 0, indicating readers */
    rwLockPtr->rwCount++;
    kMutexUnlock(&rwLockPtr->lock);
}

/* a reader releases and wakes a single writer */
/* if it is the last reader */
VOID RwLockReleaseRead(RwLock_t *const rwLockPtr)
{
    kMutexLock(&rwLockPtr->lock, RK_WAIT_FOREVER);
    rwLockPtr->rwCount--;
    if (rwLockPtr->rwCount == 0)
    {
        kCondVarSignal(&rwLockPtr->writersGo);
    }
    kMutexUnlock(&rwLockPtr->lock);
}

```

In the image below, 4 tasks—a fast writer (Task 1), a slow writer (Task 4) and two readers (Task3 is faster than Task2)—reading from and writing to a shared UINT variable:

[readerwriter 4] | <https://kernel0org.wordpress.com/wp-content/uploads/2025/06/readerwriter-4.png>

5.6. Message Queue

The key mechanism for Message Passing in *RK0* is a *Message Queue*. Messages are passed *by copy* and have a fixed-size of 4-,8-,16- or 32-byte, i.e., 1, 2, 4 or 8 words.

Two other abstractions are constructed on top of Queues:

1. *Port*, which is a server endpoint that enqueues *Remote Invocations*, or *Procedure Calls* from clients.
2. *Mailbox*, a 1-word message queue. It can be used as is, and is used by the kernel as reply route from servers to clients.

Message Queue Control Block

Buffer Address

Message Size

Number of Mesages

Write Position

Read Position

Server Status

Owner Task

Notify callback

Waiting queue

5.6.1. Size of a Message

Each declared queue has a *fixed message-size* at initialisation, and can assume, 1, 2, 4 or 8 WORDS (4, 8, 16, 32 BYTES). This constraint is intentional. Word-aligned copies are faster, predictable and safer for type casting.

(A word-aligned *single copy* will take ~5 cycles in Cortex-M3/4/7, and ~6 cycles on Cortex-M0/M0+.)

Ports have messages whose types are labeled as:

- `RK_PORT_MESG_<2,4,8>WORD`, which the first two are reserved meta-data. For variable payload one uses the `RK_PORT_MESG_COOKIE`, passing an opaque spointer as payload.

5.6.2. Blocking and non-blocking behaviour

A producer task can optionally block on a full queue, switching its state to `SENDING`. A consumer (optionally) blocks on an empty queue switching its state to `RECEIVING`.

Note that when using non-blocking calls (`RK_NO_WAIT`) they immediately return if unsuccessful.

One can use a `peek()` to read from the head of a queue without extracting the message. A `jam()` method is used to put a message on the queue head, but it does not overwrite. `jam()` is meaningless for Mailboxes.

A `postovw()` overwrites the oldest message and is allowed solely for 1-message queues. If N-message queues were allowed to be overwritten, all unread messages would leak as read and write pointers are adjusted on the operation. In case one needs overwriting full queues continuously a classic ring buffer will do.

`peek()` and `postovw()` are normally used on Mailboxes for last-message semantics—the Mailbox never goes empty; a new message is placed by overwriting the current one.

5.6.3. Ownership and Priority Inversion

Message Queues (and by extension, Mailboxes) *can* have owners, and thus, only the owners can

receive at these channels. One might think it makes it a *Port*, but not yet. The fact is, if a receiver is blocked on full queue, the sender can boost its priority.

Ports are for _ synchronous procedure calls_. The server needs to run at the client's priority, which means either boosting or demoting priority.

5.6.4. Notify callback

A callback can be registered for when a queue sends a message successfully, as a means of event notification.

5.6.5. Usage Examples

Below there is a rich set of usage examples. They demonstrate how to use message-passing in RK0, its API and helper macros to declare storage and messages at the appropriate size.

5.6.5.1. Averaging Sensor Values

A task receives measured sensor values from an ISR on a periodic rate. (A Soft Timer emulates the ISR).

Then it enqueues this data to a consumer - that will process the average value for each of 4 sensors.

The inter-task communication is designed as follows:

1. The producer pends on a Mailbox that an ISR posts to. An application timer emulates this ISR.
2. The data extracted from the Mailbox is placed in a queue with the processing task as the consumer.
3. As the producer's priority must be higher than that of the consumer, eventually, the queue will get full.
4. The producer drops the last message when the queue is full and signals the consumer.
5. Now the consumer has a batch of data to work until the next sensor update. It will block (pend on a signal) whenever the queue is empty.

```
/* helpers for pend and signal */

#define kPend(timeout)
    do
        \
    \
    \
    kTaskEventGet(0x1, RK_FLAGS_ANY, NULL, timeout); \
} while (0)

#define kSignal(taskhandle)
    do
        \
    \
    \
    kTaskEventSet(taskhandle, 0x01); \
} while (0)
```

```

/* sensor type */
typedef enum
{
    TEMPERATURE=1, HUMIDITY, CO2, FLOW
}SensorType_t;

struct sensorMsg
{
    SensorType_t sensorType;
    ULONG sensorValue;
};

typedef struct sensorMsg Mesg_t;

#define N_SENSOR      4
#define AVG_WINDOW_SIZE   10 /* 10 samples */

/* the queue */
RK_MESG_QUEUE sensorStream;
/* convenience macro to declare the queue storage */
#define N_MESSAGE 8
RK_DECLARE_MESGQ_BUF(mesgBuf, Mesg_t, N_MESSAGE)

/* timer to mimic isr */
RK_TIMER timerT1;

/* the mailbox the sensor task pends */
RK_MAILBOX sensorBox;
static Mesg_t sample = {0};
static UINT sampleErr;
VOID callBackISR(VOID* args);

VOID kApplicationInit( VOID)
{
    K_ASSERT(!kMesgQueueInit(&sensorStream,
                           mesgBuf,
                           RK_MESGQ_MESG_SIZE(Mesg_t), /* set mesg size */
                           N_MESSAGE));
    /* timer @ every 10 ms */
    K_ASSERT(!kTimerInit(&timerT1, 0, 10, callBackISR, NULL, RK_TIMER_RELOAD));
    K_ASSERT(!kMailboxInit(&sensorBox));
}

VOID callBackISR(VOID *args)
{
    RK_UNUSEARGS
    sample.sensorType = (rand() % 4) + 1;
    switch (sample.sensorType)
    {
        case TEMPERATURE:

```

```

        sample.sensorValue = ( ULONG) rand() % 50;
        break;
    case HUMIDITY:
        sample.sensorValue = ( ULONG) rand() % 100;
        break;
    case CO2:
        sample.sensorValue = ( ULONG) rand() % 1000;
        break;
    case FLOW:
        sample.sensorValue = ( ULONG) rand() % 10;
        break;
    default:
        break;
    }
/* Mailbox carries one word: post a pointer to sample */
Mesg_t *samplePtr = &sample;
RK_ERR err = kMailboxPost(&sensorBox, &samplePtr, RK_NO_WAIT);
if (err != RK_ERR_SUCCESS)
    sampleErr++;
}

/* Producer - higher priority, blocks on mailbox */
VOID Task1(VOID *args)
{
    RK_UNUSEARGS
    Mesg_t *recvSample = NULL;
    while (1)
    {
        /* receive a pointer */
        RK_ERR errmbox = kMailboxPend(&sensorBox, &recvSample, RK_WAIT_FOREVER);
        K_ASSERT( errmbox==RK_ERR_SUCCESS);

        /* enqueue by copy into the stream */
        RK_ERR err = kMesgQueueSend(&sensorStream, recvSample, RK_NO_WAIT);

        K_ASSERT(err >= 0); /* either succesful or unsuccesful */
        if (err == RK_ERR_SUCCESS)
        {
            CHAR const *sensorTypeStr = NULL;
            if (recvSample->sensorType == 1)
                sensorTypeStr = "TEMP";
            if (recvSample->sensorType == 2)
                sensorTypeStr = "HUM";
            if (recvSample->sensorType == 3)
                sensorTypeStr = "CO2";
            if (recvSample->sensorType == 4)
                sensorTypeStr = "FLOW";
            printf( "ENQ: [%@lums, %s, %lu] \r\n", kTickGet(), sensorTypeStr,
                    recvSample->sensorValue);
        }
    }
}

```

```

/* full, drop this sample and signal task2 */
else if (err == RK_ERR_MESGQ_FULL)
{
    kSignal(task2Handle);
}
}

/* for each sensor:
   . a ring buffer of AVG_WINDOW_SIZE values
   . sum of values
   . an index table (=enum - 1 eg., HUMIDITY IDX=2-1=1)
*/
static ULONG ringBuf[N_SENSOR][AVG_WINDOW_SIZE];
static ULONG ringSum[N_SENSOR] = {0};
static UINT ringIndex[N_SENSOR] = {0};

void Task2( void *args)
{

    RK_UNUSEARGS
    Mesg_t readSample;
    while (1)
    {
        RK_ERR err = kMesgQueueRecv(&sensorStream, &readSample, RK_NO_WAIT);
        if (err == RK_ERR_SUCCESS)
        {
            UINT sensorIdx = readSample.sensorType - 1;

/* remove oldest sample */
            ULONG oldest = ringBuf[sensorIdx][ringIndex[sensorIdx]];
            ringSum[sensorIdx] -= oldest;

/* push new sample */
            ringBuf[sensorIdx][ringIndex[sensorIdx]] = readSample.sensorValue;
            ringSum[sensorIdx] += readSample.sensorValue;

/* index incr-wrap */
            ringIndex[sensorIdx]++;
            ringIndex[sensorIdx] %= AVG_WINDOW_SIZE;

/* simple average */
            ULONG avg = ringSum[sensorIdx] / AVG_WINDOW_SIZE;

            CHAR const *sensorTypeStr = NULL;
            if (readSample.sensorType == 1)
                sensorTypeStr = "TEMP";
            if (readSample.sensorType == 2)
                sensorTypeStr = "HUM";
            if (readSample.sensorType == 3)

```

```

        sensorTypeStr = "CO2";
        if (readSample.sensorType == 4)
            sensorTypeStr = "FLOW";

        printf( "DEQ: [%@%lums, %s, %lu] | AVG: %lu \r\n", kTickGet(),
                sensorTypeStr, readSample.sensorValue, avg);

    }
    else
    {
        kPend(RK_WAIT_FOREVER);
    }

}

```

OUTPUT:

```

ENQ: [@550ms, CO2, 571]
ENQ: [@560ms, FLOW, 4]
ENQ: [@570ms, FLOW, 4]
ENQ: [@580ms, HUM, 25]
ENQ: [@590ms, CO2, 931]
ENQ: [@600ms, CO2, 487]
ENQ: [@610ms, FLOW, 7]
ENQ: [@620ms, HUM, 79]

```

>>> Queue is full. Now offload and process. Note the order remains <<<

```

DEQ: [@630ms, CO2, 571] | AVG: 460
DEQ: [@631ms, FLOW, 4] | AVG: 5
DEQ: [@632ms, FLOW, 4] | AVG: 5
DEQ: [@633ms, HUM, 25] | AVG: 52
DEQ: [@634ms, CO2, 931] | AVG: 553
DEQ: [@635ms, CO2, 487] | AVG: 549
DEQ: [@636ms, FLOW, 7] | AVG: 5
DEQ: [@637ms, HUM, 79] | AVG: 55

```

>>> Consumer is preempted <<<

```

ENQ: [@640ms, CO2, 913]
ENQ: [@650ms, CO2, 134]
ENQ: [@660ms, HUM, 47]
ENQ: [@670ms, HUM, 30]
ENQ: [@680ms, TEMP, 7]
ENQ: [@690ms, CO2, 726]
ENQ: [@700ms, FLOW, 7]
ENQ: [@710ms, TEMP, 43]

```

```

DEQ: [@720ms, CO2, 913] | AVG: 578
DEQ: [@721ms, CO2, 134] | AVG: 543

```

```

DEQ: [@722ms, HUM, 47] | AVG: 51
DEQ: [@723ms, HUM, 30] | AVG: 44
DEQ: [@724ms, TEMP, 7] | AVG: 20
DEQ: [@725ms, CO2, 726] | AVG: 592
DEQ: [@726ms, FLOW, 7] | AVG: 5
DEQ: [@727ms, TEMP, 43] | AVG: 23

```

5.6.5.2. Queue Select using Notify Callback

A task is receiving from many queues and need to know which one has been able to complete.

The `notifyCbk(queue*)` is executed every time a send is successful. In this case it is using an Event Signal to a task. The Signals Flag indicate the queue number - as a contract - of which queue has completed sends. Note that as sends may coalesce, while a flag caps at 1, the consumer will drain each queue until it is empty, or it is preempted

```

/* Many-to-1 queue channels */

/* Consumer Select queue based on its event
flags, that a send callback for each queue runs */

#define LOG_PRIORITY 4 /* keep logger as lowest-priority user task */
#define STACKSIZE 256

#define NQUEUES 3 /* number of queues */
#define QSIZ 8 /* depth of each queue */

#define Q0_FLAG    (1UL << 0)
#define Q1_FLAG    (1UL << 1)
#define Q2_FLAG    (1UL << 2)
#define QFLAGS     (ULONG)(Q0_FLAG | Q1_FLAG | Q2_FLAG)

typedef struct
{
    RK_TASK_HANDLE producer;
    UINT payload;
} MESG_t;

/* Succesful Send callbacks */
/* each callback follows this pattern */
static inline
VOID sendNotify0(RK_MESG_QUEUE *qPtr)
{
    (VOID)qPtr;
    kTaskEventSet(consumerHandle, Q0_FLAG);
    /* Q1 flag for queue1 and so forth */
}

/* each callback is installed using kMesgQueueInstallSendCbk
on kApplicationInit() */

```

```

/* helper to send */
static inline
VOID enqueueSample(RK_MESG_QUEUE *qPtr UINT payload)
{
    MESG_t mesg = {
        .payload = payload,
        .producer = RK_RUNNING_HANDLE,
    };
    RK_ERR err = kMesgQueueSend(qPtr, &mesg, RK_WAIT_FOREVER);
    K_ASSERT(err == RK_ERR_SUCCESS);
}

VOID Prod0Task(VOID *args)
{
    RK_UNUSEARGS
    UINT seq = 0U;

    while (1)
    {
        enqueueSample(&queues[0], seq++);
        kSleepRelease(25); /* fast producer */
    }
}

VOID Prod1Task(VOID *args)
{
    UINT seq = 0U;
    RK_UNUSEARGS

    while (1)
    {
        enqueueSample(&queues[1], seq++);

        /* every fourth sample, also tickle the third queue */
        if ((seq & 0x3U) == 0U)
        {
            enqueueSample(&queues[2], seq);
        }

        kSleepRelease(60);
    }
}

/* Consumer listens on all queues, selecting those on its signal flags. */
VOID ConsumerTask(VOID *args)
{
    RK_UNUSEARGS

    MESG_t recv = {0};

```

```

ULONG flags = 0UL;

while (1)
{
    flags = 0UL;
    kTaskEventGet(QFLAGS, RK_EVENT_FLAGS_ANY, &flags,
                  RK_WAIT_FOREVER);

    for (UINT i = 0; i < NQUEUES; ++i)
    {
        if (flags & (1UL << i))
        {
            while (kMesgQueueRecv(&queues[i], (VOID*)&recv, RK_NO_WAIT) ==
                   RK_ERR_SUCCESS)
            {
                logPost("Q%u <- sender=%s payload=%u", i,
                        RK_TASK_NAME(recv.producer), recv.payload);
            }
        }
    }
}

```

```

0 ms :: Q1 <- sender=Prod1 payload=0
250 ms :: Q0 <- sender=Prod0 payload=1
500 ms :: Q0 <- sender=Prod0 payload=2
600 ms :: Q1 <- sender=Prod1 payload=1
750 ms :: Q0 <- sender=Prod0 payload=3
1000 ms :: Q0 <- sender=Prod0 payload=4
1200 ms :: Q1 <- sender=Prod1 payload=2
1250 ms :: Q0 <- sender=Prod0 payload=5
1500 ms :: Q0 <- sender=Prod0 payload=6
1750 ms :: Q0 <- sender=Prod0 payload=7
1800 ms :: Q1 <- sender=Prod1 payload=3
1800 ms :: Q2 <- sender=Prod1 payload=4
2000 ms :: Q0 <- sender=Prod0 payload=8
2250 ms :: Q0 <- sender=Prod0 payload=9
2400 ms :: Q1 <- sender=Prod1 payload=4
2500 ms :: Q0 <- sender=Prod0 payload=10
2750 ms :: Q0 <- sender=Prod0 payload=11
3000 ms :: Q0 <- sender=Prod0 payload=12
3000 ms :: Q1 <- sender=Prod1 payload=5
3250 ms :: Q0 <- sender=Prod0 payload=13
3500 ms :: Q0 <- sender=Prod0 payload=14
3600 ms :: Q1 <- sender=Prod1 payload=6
3750 ms :: Q0 <- sender=Prod0 payload=15
4000 ms :: Q0 <- sender=Prod0 payload=16
4200 ms :: Q1 <- sender=Prod1 payload=7
4200 ms :: Q2 <- sender=Prod1 payload=8
4250 ms :: Q0 <- sender=Prod0 payload=17

```

```

4500 ms :: Q0 <- sender=Prod0 payload=18
4750 ms :: Q0 <- sender=Prod0 payload=19
4800 ms :: Q1 <- sender=Prod1 payload=8
5000 ms :: Q0 <- sender=Prod0 payload=20
5250 ms :: Q0 <- sender=Prod0 payload=21

```

5.6.5.3. Synchronous Client-Server Procedure Call

The example below computes a CRC on the server for the client's payload and returns it as the reply code. Logically, this is unbuffered: the client blocks until it receives the reply.

For `kPortSendRecv()`, each client task must register its reply mailbox once using `kRegisterMailbox()`.

Note that the server may temporarily adopt the client priority while serving the request. After `kPortReplyDone()`, it returns to nominal priority.

```

#include <application.h>
#include <logger.h>
#define STACKSIZE 256
#define PORT_MSG_WORDS 4U /* 2 words meta + 2 words payload */
#define PORT_CAPACITY 16

/* tasks */
RK_DECLARE_TASK(serverHandle, ServerTask,      stack1, STACKSIZE)
RK_DECLARE_TASK(clientHandle, ClientTask,      stack2, STACKSIZE)

/* port */
static RK_PORT serverPort;
RK_DECLARE_PORT_BUF(portBuf, PORT_MSG_WORDS, PORT_CAPACITY)

/* 4-word message format; first two words are metadata */
typedef RK_PORT_MESG_4WORD RpcMsg;

static inline UINT crc32(const VOID *data, ULONG size);
static inline BYTE xorshift8(void);

VOID kApplicationInit(void)
{
    K_ASSERT(!kCreateTask(&serverHandle, ServerTask, RK_NO_ARGS, "Server", stack1,
STACKSIZE, 1, RK_PREEMPT));
    K_ASSERT(!kCreateTask(&clientHandle, ClientTask, RK_NO_ARGS, "Client", stack2,
STACKSIZE, 2, RK_PREEMPT));

    /* init port */
    K_ASSERT(!kPortInit(&serverPort, portBuf, PORT_MSG_WORDS, PORT_CAPACITY,
serverHandle));

    logInit();
}

```

```

VOID ServerTask(VOID *args)
{
    RK_UNUSEARGS
    RpcMsg msg;
    while(1)
    {
        /* receive next request; server may adopt client priority here */
        K_ASSERT(!kPortRecv(&serverPort, &msg, RK_WAIT_FOREVER));

        BYTE *vector = (BYTE*) msg.payload[0];
        ULONG size = msg.payload[1];
        UINT crc = crc32(vector, size);

        logPost("[SERVER] Will Reply CRC=0x%04X | Eff Prio=%d | Nom Prio=%d",
                crc, runPtr->priority, runPtr->prioNominal);

        /* must end with kPortReplyDone */
        K_ASSERT(!kPortReplyDone(&serverPort, (ULONG const*)&msg, crc));

        logPost("[SERVER] Finished. | Eff Prio: %d | Nom Prio: %d", runPtr->priority,
                runPtr->prioNominal);
    }
}

VOID ClientTask(VOID *args)
{
    RK_UNUSEARGS
    static BYTE vec[8];
    for (UINT i = 0; i < 8; ++i)
        vec[i] = xorshift8();

    RK_MAILBOX replyBox;
    K_ASSERT(!kMailboxInit(&replyBox));
    K_ASSERT(!kRegisterMailbox(clientHandle, &replyBox));
    RpcMsg msg = {0};
    msg.payload[0] = (ULONG) vec; /* pointer as one word */
    msg.payload[1] = 8;           /* number of bytes */

    UINT reply = 0;
    while(1)
    {
        /* Send-Receive: a call */
        UINT want = crc32(vec, 8);
        K_ASSERT(!kPortSendRecv(&serverPort, (ULONG*)&msg, &reply,
                               RK_WAIT_FOREVER));
        logPost("[CLIENT] Need=0x%04X | Recvd=0x%04X", want, reply);
        /* if reply is correct, generate a new payload */
        if (want == reply)
            for (UINT i = 0; i < 8; ++i) vec[i] = xorshift8();
        kSleepRelease(1000);
    }
}

```

```
}
```

```
28000 ms :: [CLIENT] Need=0xC2A6C337 | Recvd=0xC2A6C337
29000 ms :: [SERVER] Will Reply CRC=0x93F4110A | Eff Prio=2 | Nom Prio=1
29000 ms :: [SERVER] Finished. | Eff Prio: 1 | Nom Prio: 1
29000 ms :: [CLIENT] Need=0x93F4110A | Recvd=0x93F4110A
30000 ms :: [SERVER] Will Reply CRC=0x7A8FA006 | Eff Prio=2 | Nom Prio=1
30000 ms :: [SERVER] Finished. | Eff Prio: 1 | Nom Prio: 1
30000 ms :: [CLIENT] Need=0x7A8FA006 | Recvd=0x7A8FA006
31000 ms :: [SERVER] Will Reply CRC=0x9051C8B1 | Eff Prio=2 | Nom Prio=1
31000 ms :: [SERVER] Finished. | Eff Prio: 1 | Nom Prio: 1
31000 ms :: [CLIENT] Need=0x9051C8B1 | Recvd=0x9051C8B1
32000 ms :: [SERVER] Will Reply CRC=0x11F29117 | Eff Prio=2 | Nom Prio=1
32000 ms :: [SERVER] Finished. | Eff Prio: 1 | Nom Prio: 1
```

5.7. Most-Recent Message Protocol (MRM)

MRM Control Block

MRM Buffer Allocator

Data Buffer Allocator

Current MRM Buffer Address

Data Size (Message Size)

MRM Buffer

Data Buffer Address

Readers Count

Data Buffer

Application-dependent

There is little practical difference between a message that does not arrive and one with no valid (stale) data. But when wrong (or stale) data is processed - e.g., to define a set point on a loop - a system can fail badly.



Design Choice: provide a broadcast asynchronous message-passing scheme that guarantees data freshness and integrity for all readers.

Benefits: The system has a mechanism to meet strict deadlines that cannot be predicted on design time.

Control loops reacting to unpredictable time events—like a robot scanning an environment or a drive-by-wire system—require a different message-passing approach. Readers cannot "look at the past" and cannot block. The most recent data must be delivered lock-free and have guaranteed

integrity.

As *Ports*, the *MRM* is a high-level mechanism. It was chosen to be provided as a kernel service, given its distinctive nature and suitability for *RKO* applications.

5.7.1. Functional Description

An *MRM* works as a *1-to-many asynchronous Mailbox* - with a lock-free specialisation that enables several readers to get the most recent deposited message with no integrity issues. Whenever a reader reads an *MRM* buffer, it will find the most recent data transmitted. It can also be seen as an extension of the Double Buffer pattern for a 1:N communication.

The core idea of the *MRM* protocol is that readers can only access the buffer that is classified as the '*most recent buffer*'. After a writer *publish()* a message, that will be the only message readers can *get()*—any former message being processed by a reader was grabbed *before* a new *publish()* - and, from now on, can only be *unget()*, eventually returning to the pool.

To clarify further, the communication steps are listed:

1. A producer first reserves an *MRM* Buffer - the reserved *MRM* Buffer is not available for reading until it is published.
2. A message buffer is allocated and filled, and its address is within an *MRM* Buffer. The producer *publishes* the message. From now on, it is *the most recent message*. Any former published buffer is no longer visible to new readers
3. A reader starts by *getting* an *MRM* Buffer. A *get()* operation delivers a copy of the message to the reader's scope. Importantly, this operation increases the number of readers associated to that *MRM* Buffer.

Before ending its cycle, the task releases (*unget()*) the buffer; on releasing, the kernel checks if the caller task is the last reader and if the buffer being released is not the current *MRM* Buffer.

If the above conditions are met, the *unget()* operation will return the *MRM* buffer to the pool. If there are more readers, OR if it is the current buffer, it remains available.

When the *reserve* operation detects that the most recent buffer still has readers, a new buffer is allocated to be written and published. If it has no readers, it is reused.

This way, the worst case is a sequence of *publish()* with no *unget()* at all—this would lead to the writer finding no buffer to reserve. This is prevented by making: **N Buffers = N tasks + 1**.

5.7.2. MRM Control Block Configuration

What might lead to some confusion when initialising an *MRM* Control Block is the need for two different pools:

- One pool will be the storage for the *MRM* Buffers, which is the data structure for the mechanism.
- Another pool is for the actual payload. The messages.

Both pools must have the same number of elements: the number of tasks communicating + 1.

- The size of the data buffers is application-dependent - and is passed as a number of *words*. The minimal message size is 32-bit.
- If using data structures, keep it aligned to 4 to take advantage of the performance of aligned memory.

5.7.3. Usage Example

Consider a modern car - speed variations are of interest in many modules. With a somehow "naive" approach, let us consider three modules and how they should react when speed varies:

1. **Cruiser Control:** For the Cruiser Control, a speed increase might signify the driver wants manual control back, and it will likely turn off.
2. **Windshield Wipers:** If they are on, a speed change can affect the electric motor's adjustments to the air resistance.
3. **Radio:** Speed changes reflect the aerodynamic noise - the radio volume might need adjustment.

As the variations are unpredictable, we need a mechanism to deliver the last speed in order of importance for all these modules. From highest to lowest priority, Cruise, Wipers, and Radio are the three modules that range from safety to comfort.

To emulate this scenario, we can write an application with a higher priority task that sleeps and wakes up at pseudo-random times to produce random values that represent the (unpredictable) speed changes.

The snippet below has 4 periodic tasks. Tasks are periodic using the `kSleepRelease()` primitive.

There is a sequential counter. When a task wakes and see it was not incremented it just runs.

The producer publishes new data at a random interval, preempting whatever task is running at the moment. Given the random updates,

```
typedef struct
{
    UINT speed;
    ULONG timeStamp;
} Mesg_t;

#define STACKSIZE 256
#define N_MRM (5) /* Number of MRMNs N Tasks + 1 */
#define MRM_MESG_SIZE (sizeof(Mesg_t) / 4) /* In WORDS */
RK_MRM MRMCtl; /* MRM control block */
RK_MRM_BUF buf[N_MRM]; /* MRM pool */
Mesg_t data[N_MRM]; /* message data pool */

RK_DECLARE_TASK(speedSensorHandle, SpeedSensorTask, stack1, STACKSIZE)
RK_DECLARE_TASK(cruiserHandle, CruiserTask, stack2, STACKSIZE)
```

```

RK_DECLARE_TASK(wiperHandle, WiperTask, stack3, STACKSIZE)
RK_DECLARE_TASK(radioHandle, RadioTask, stack4, STACKSIZE)
volatile UINT seq = 0;
VOID kApplicationInit(VOID)
{
    kCreateTask(&speedSensorHandle, SpeedSensorTask, RK_NO_ARGS, "SpeedTsk",
               stack1, STACKSIZE, 1, RK_PREEMPT);

    kCreateTask(&cruiserHandle, CruiserTask, RK_NO_ARGS, "CruiserTsk", stack2,
               STACKSIZE, 2, RK_PREEMPT);

    kCreateTask(&wiperHandle, WiperTask, RK_NO_ARGS, "WiperTsk", stack3,
               STACKSIZE, 3, RK_PREEMPT);

    kCreateTask(&radioHandle, RadioTask, RK_NO_ARGS, "RadioTsk", stack4,
               STACKSIZE, 4, RK_PREEMPT);

    kMRMInit(&MRMCtl, buf, data, N_MRM, MRM_MESG_SIZE);

    logInit(5);
}

VOID SpeedSensorTask(VOID *args)
{
    RK_UNUSEARGS

    Mesg_t sendMesg = {0};
    while (1)
    {
        RK_TICK sleepTicks = ((RK_TICK)rand() % 18) + 1;
        kSleep(sleepTicks);
        RK_TICK currTick = kTickGetMs();
        UINT speedValue = (UINT)(rand() % 170) + 1;
        sendMesg.speed = speedValue;
        sendMesg.timeStamp = currTick;
        /* grab a buffer */
        RK_MRM_BUF *bufPtr = kMRMReserve(&MRMCtl);
        if (bufPtr != NULL)
        {
            K_ASSERT(!kMRMPublish(&MRMCtl, bufPtr, &sendMesg));
            printf("!!!!! @%lums SPEED UPDATE: %u mph\r\n", kTickGetMs(),
speedValue);
            seq += 1;
        }
        else
        { /* cannot fail */
            logError("MRM protocol could not find a free buffer\r\n");
        }
        /* publish */
    }
}

```

```

}

VOID CruiserTask(VOID *args)
{
    RK_UNUSEARGS
    Mesg_t recvMesg = {0};
    while (1)
    {
        RK_MRM_BUF *readBufPtr = kMRMGet(&MRMCtl, &recvMesg);
        if (readBufPtr)
        {
            logPost("CRUISER: (%u mph, %lu ms)", recvMesg.speed,
recvMesg.timeStamp);

            kMRMUnget(&MRMCtl, readBufPtr);
        }
        kSleepRelease(4);
    }
}

VOID WiperTask(VOID *args)
{
    RK_UNUSEARGS
    Mesg_t recvMesg = {0};

    while (1)
    {

        RK_MRM_BUF *readBufPtr = kMRMGet(&MRMCtl, &recvMesg);
        if (readBufPtr)
        {
            logPost("WIPERS: (%u mph, %lu ms)", recvMesg.speed,
recvMesg.timeStamp);

            kMRMUnget(&MRMCtl, readBufPtr);
        }
        kSleepRelease(7);
    }
}

VOID RadioTask(VOID *args)
{
    RK_UNUSEARGS
    Mesg_t recvMesg = {0};
    while (1)

    {

        RK_MRM_BUF *readBufPtr = kMRMGet(&MRMCtl, &recvMesg);

        if (readBufPtr)
        {

```

```

        logPost("RADIO: (%u mph, %lu ms) ", recvMesg.speed,
recvMesg.timeStamp);
        kMRMUnget(&MRMCtl, readBufPtr);
    }
    kSleepRelease(11);
}
}

```

Thus, different situations can happen:

- All tasks read the updated pair (speed, time)
- Not all tasks receive the updated pair because another update happens in between.
- No tasks receive an update - because another happens too soon.
- No update happens between in the period of a given task. It receives the same value. No problems.

All these cases are on the log:

Logs show: (last speed record, record time)

```

!!!@120ms SPEED UPDATE: 164 mph
120 ms :: CRUISER: (164 mph, 120 ms)
140 ms :: WIPERS: (164 mph, 120 ms)
!!! @150ms SPEED UPDATE: 80 mph
160 ms :: CRUISER: (80 mph, 150 ms)
200 ms :: CRUISER: (80 mph, 150 ms)
210 ms :: WIPERS: (80 mph, 150 ms)
220 ms :: RADIO: (80 mph, 150 ms)
240 ms :: CRUISER: (80 mph, 150 ms)
!!!@280ms SPEED UPDATE: 49 mph
280 ms :: CRUISER: (49 mph, 280 ms)
280 ms :: WIPERS: (49 mph, 280 ms)
320 ms :: CRUISER: (49 mph, 280 ms)
330 ms :: RADIO: (49 mph, 280 ms)
350 ms :: WIPERS: (49 mph, 280 ms)
360 ms :: CRUISER: (49 mph, 280 ms)
400 ms :: CRUISER: (49 mph, 280 ms)
420 ms :: WIPERS: (49 mph, 280 ms)
440 ms :: CRUISER: (49 mph, 280 ms)
440 ms :: RADIO: (49 mph, 280 ms)
!!!@450ms SPEED UPDATE: 87 mph
480 ms :: CRUISER: (87 mph, 450 ms)
490 ms :: WIPERS: (87 mph, 450 ms)
520 ms :: CRUISER: (87 mph, 450 ms)
!!! @540ms SPEED UPDATE: 110 mph
550 ms :: RADIO: (110 mph, 540 ms)
560 ms :: CRUISER: (110 mph, 540 ms)
560 ms :: WIPERS: (110 mph, 540 ms)

```

```
600 ms :: CRUISER: (110 mph, 540 ms)
630 ms :: WIPERS: (110 mph, 540 ms)
!!! @640ms SPEED UPDATE: 22 mph
640 ms :: CRUISER: (22 mph, 640 ms)
660 ms :: RADIO: (22 mph, 640 ms)
```

The highlight is that controllers can keep their pace, while receiving fresh data - you can see it on the timestamp on the image.

Again, they might receive the same data more than once or miss samples; what is important is that *they are not lagging and consuming stale data*.

Chapter 6. Error Handling

6.1. Fail fast

While tracing and error handling are yet to be largely improved (and that is when the 1.0.0 version will be released), currently *RKO* employs a policy of *failing fast* in **debug mode**.

When Error Checking is enabled, every kernel call will be 'defensive', checking for correctness of parameters and invariants, null dereferences, etc.

In these cases is more useful to allow the first error to halt the execution by calling an Error Handler function to observe the program state.

A trace structure records the address of the running TCB, its current stack pointer, the link register (that is, the PC at kErrorHandler was called), and a time stamp.

This record is on a `.noinit` RAM section, so it is visible if CPU resets. A fault code is stored in a global `faultID` and on the trace structure. Developers can hook in custom behaviour.

If the kernel is configured to not halt on a fault, but Error Checking is enabled, functions will return negative values in case of an error.

On the other hand, when Error Checking is disabled or `NDEBUG` is defined nothing is checked, reducing code size and improving performance.

(*Some deeper internal calls have assertion. For those, only `NDEBUG` defined ensures they are disabled.*)

6.2. Stack Overflow

Stack overflow is detected (not prevented) using a "stack painting" with a sentinel word. Stack Overflow detection is enabled by defining the assembler preprocessor `_KDEF_STACKOVFLW` when compiling.

As a matter of fact, sizing your stack is something you must do diligently when programming a system. I would say a mechanism for stack overflow detection is on the bottom of the list of 'must-have' features.

One can take advantage of the static task model - *it is possible to predict offline* the deepest call within any task. The compiler flag `-fstack-usage` creates `.su` files indicating the depth of every function within a module. This is an example of compilation-unit output:

```
core/src/ksema.c:34:8:kSemaphoreInit    88  static
core/src/ksema.c:74:8:kSemaphorePend    96  static
core/src/ksema.c:189:8:kSemaphorePost   88  static
core/src/ksema.c:265:8:kSemaphoreFlush  72  static
core/src/ksema.c:306:5:kSemaphoreQuery  40  static
core/src/kmutex.c:128:8:kMutexInit     16  static
core/src/kmutex.c:165:8:kMutexLock    120 static
```

```
core/src/kmutex.c:325:8:kMutexUnlock    96  static  
core/src/kmutex.c:425:6:kMutexQuery  56  static
```

These are the worst cases. Now, you identify the depth of the longest *chain of calls* for a task using these services and add a generous safety margin — 30%. The cap depends on your budget.

Importantly, you also have to size the System Stack. This initial size is defined in `linker.ld` by the symbol `Min_Stack_Size`. In this case, account for the depth of `main()`, `kApplicationInit()`, and all interrupt handlers; again, inspect the longest call chain depth. Assume interrupts always add to the worst static depth, and account for *nested interrupts*.

6.3. Deadlocks

There are deadlock-recovery techniques in the literature, but they are generally unfeasible here. The kernel provides bounded waiting, enforces priority-ordered waiting queues, applies mutex priority inheritance, and offers lock-free primitives and period-drift compensation. None of these techniques prevents deadlocks by itself (and with bounded blocking plus lock-free primitives, one can still get *livelocks*).

- Ordered Locking:

For those programming the application, despite following the RMS rule of higher priority for higher request rate tasks, the golden rule for locking is acquiring resources in an unidirectional order *throughout the entire application*:

```
acquire(A);  
acquire(B);  
acquire(C);  
  
.  
  
.  
  
release(C);  
release(B);  
release(A);
```

This breaks circular waiting.

For instance:

```
TaskA:  
wait(R1);  
wait(R2);  
/* critical section */  
signal(R2);  
signal(R1);
```

```
TaskB:  
wait(R1);
```

```
wait(R2);
/* critical section */
signal(R2);
signal(R1);
```

But, if:

```
TaskA:
wait(R1);
wait(R2);
.
```

```
TaskB:
wait(R2);
wait(R1);
.
```

There are some possible outcomes:

1. Deadlock:

- TaskA runs: acquires R1
- TaskB runs: acquires R2
- TaskA runs: tries to acquire R2 — blocked
- TaskB runs: tries to acquire R1 — blocked

2. No deadlock:

- TaskA runs: acquires R1
- TaskA runs: acquires R2 (nobody is holding R2)
- TaskA releases both; TaskB runs and acquires both (in either order)

Overall, there is no deadlock if tasks do not overlap in critical sections. That is why systems run for years without deadlocks and eventually: *plotf*.

- Use a 'master-lock' with 'try' semantics:

Another technique that can be employed is if one needs to acquire multiple locks—acquire them all or none using a try-lock ([RK_NO_WAIT](#)). If any of the tries fail, the task gives up on acquiring the resources and backs off, releasing all successful locks to retry later (most simply, using a sleep queue). That is easier said than done, though, and, as mentioned, if not well done, instead of deadlocks, one gets livelocks.

(*Livelocks* are when a couple of tasks keep running, but the program does not advance.)

Chapter 7. RKO Services API

7.1. Convention

- A kernel call starts with a lowercase **k**. Typically it is followed by a kernel object identifier and an action.

```
kSemaphorePend(&sema, 800); /* pend on a semaphore; 800 ticks time-out */
```

- When **k** is followed by an action, it is acting on the caller task.

```
kSleep(150); /* sleep the caller task for 150 ticks */
```

- Some calls can act either on the caller or on another task:

```
/* stores the signal flags of the task identified by task1Handle on queryValue */
kTaskEventQuery(task1Handle, &queryValue);

/* retrieves its own signal flags */
kTaskEventQuery(NULL, &queryValue);
```

7.2. Return Values

With a few exceptions, kernel calls return a **RK_ERR** error code. **0** is a successful operation (**RK_ERR_SUCCESS**) and any negative value is an error that indicates failure. A positive value is an unsuccessful operation, but will not lead the system to failure (e.g., any unsuccessful **try** operation).

[kapi.h: RKO API](#)

